

Conceptualizing the Processing Model for Apache Spark Structured Streaming

GETTING STARTED WITH STRUCTURED STREAMING



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

**Streaming and its place in the
Apache Spark stack**

RDDs and DataFrames

Structured Streaming in Spark 2.x

Batch as a prefix of stream

**ClusterManager, SparkSession, and
Executors**

Prerequisites and Course Outline

Prerequisites

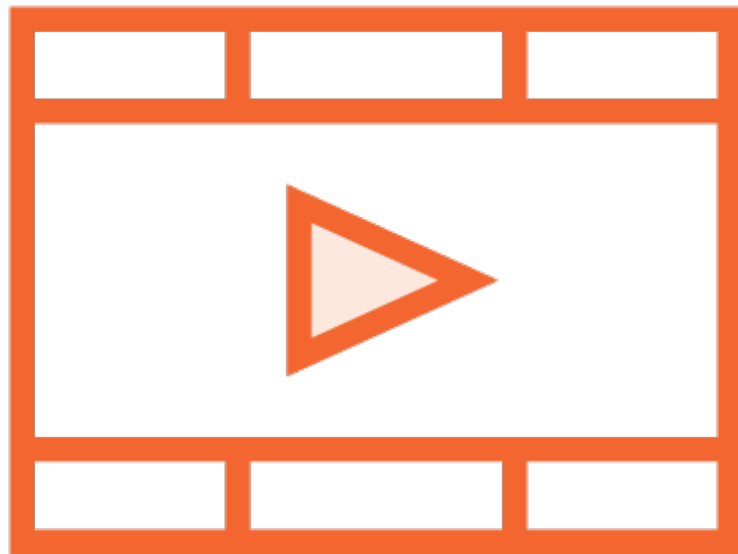


Comfortable programming in Python

**Some exposure to Apache Spark
and DataFrames**

**Exposure to streaming data would
be helpful, but is not required**

Prerequisite Courses



Apache Spark Fundamentals
Getting Started with Spark 2

Course Outline



Getting Started with Structured Streaming

Executing Streaming Queries

Understanding Scheduling and Checkpointing

Configuring Processing Models

Understanding Query Planning

Introducing Apache Spark

Hadoop

A diagram showing the three main components of Hadoop: HDFS, MapReduce, and YARN. These components are represented as three gray rectangular boxes arranged horizontally within a larger blue-bordered rectangle. Below each box is a descriptive text.

HDFS

**A file system to
manage the storage
of data**

MapReduce

**A framework to
define a data
processing task**

YARN

**A framework to run
the data processing
task**

Co-ordination Between Hadoop Blocks

MapReduce



**User defines map and
reduce tasks using the
MapReduce API**

YARN

HDFS

Co-ordination Between Hadoop Blocks

MapReduce



YARN

HDFS

**A job is triggered on the
cluster**

Co-ordination Between Hadoop Blocks

MapReduce

YARN



HDFS

YARN figures out where and how to run the job, and stores the result in HDFS

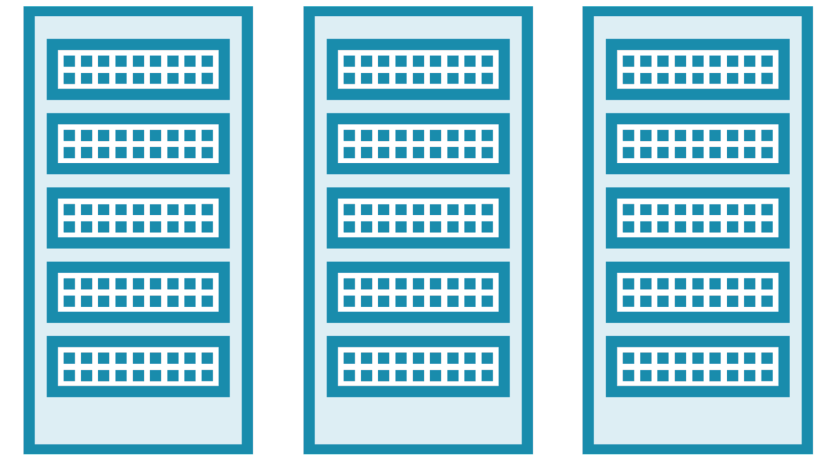
Apache Spark



General Purpose



Interactive



**Distributed
Computing**

An engine for data processing and analysis

Apache Spark



Analytics and ML on Big Data

Extremely powerful and popular Big Data technology

Distributed computing framework for general-purpose computing

Open-source from Apache

Written in Scala

Apache Spark



Real-time as well as batch

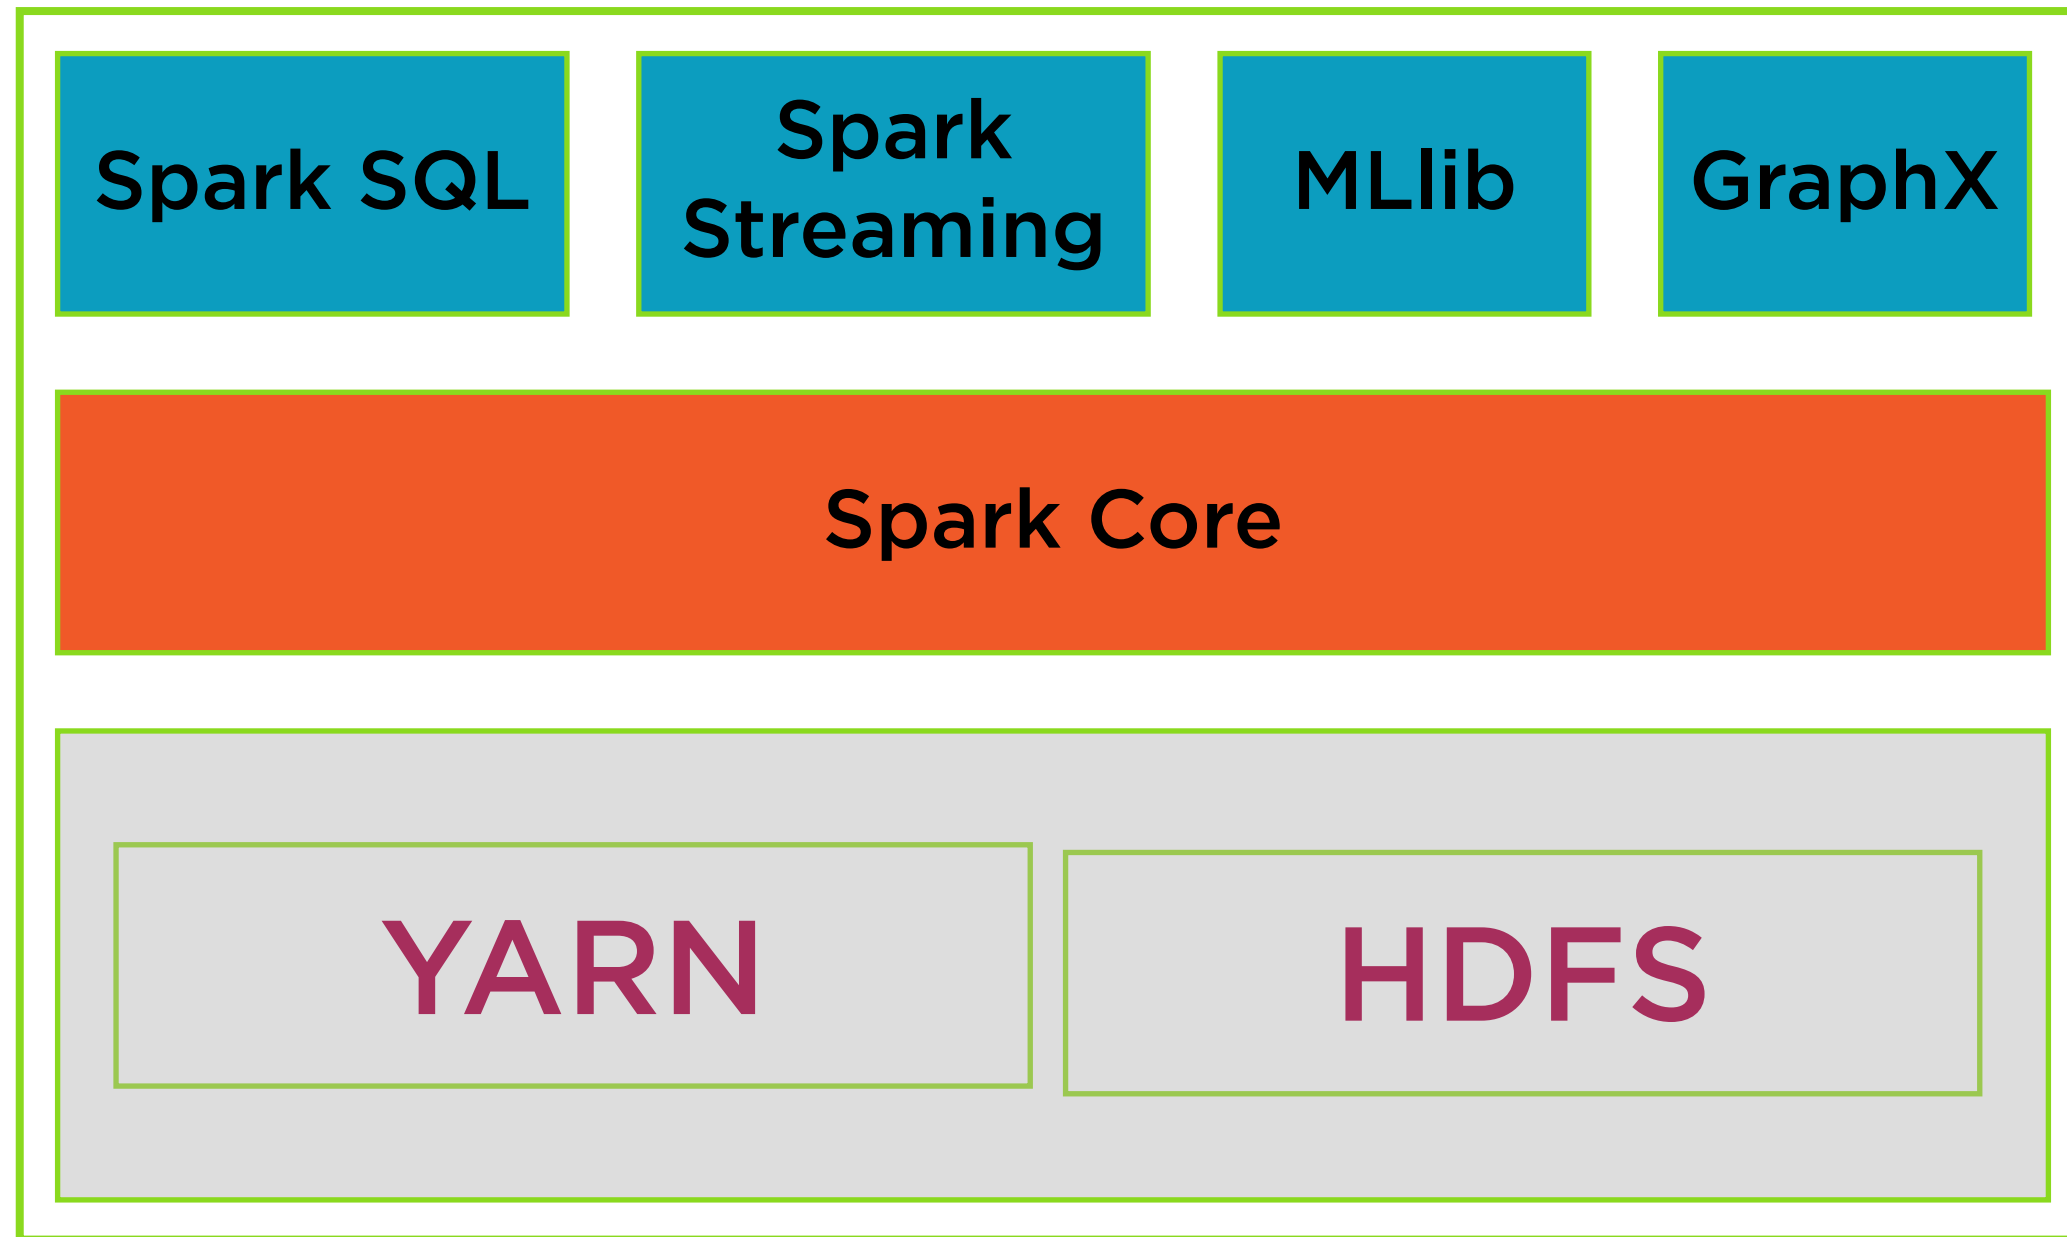
Interactive REPL environment

Read-Evaluate-Print-Loop

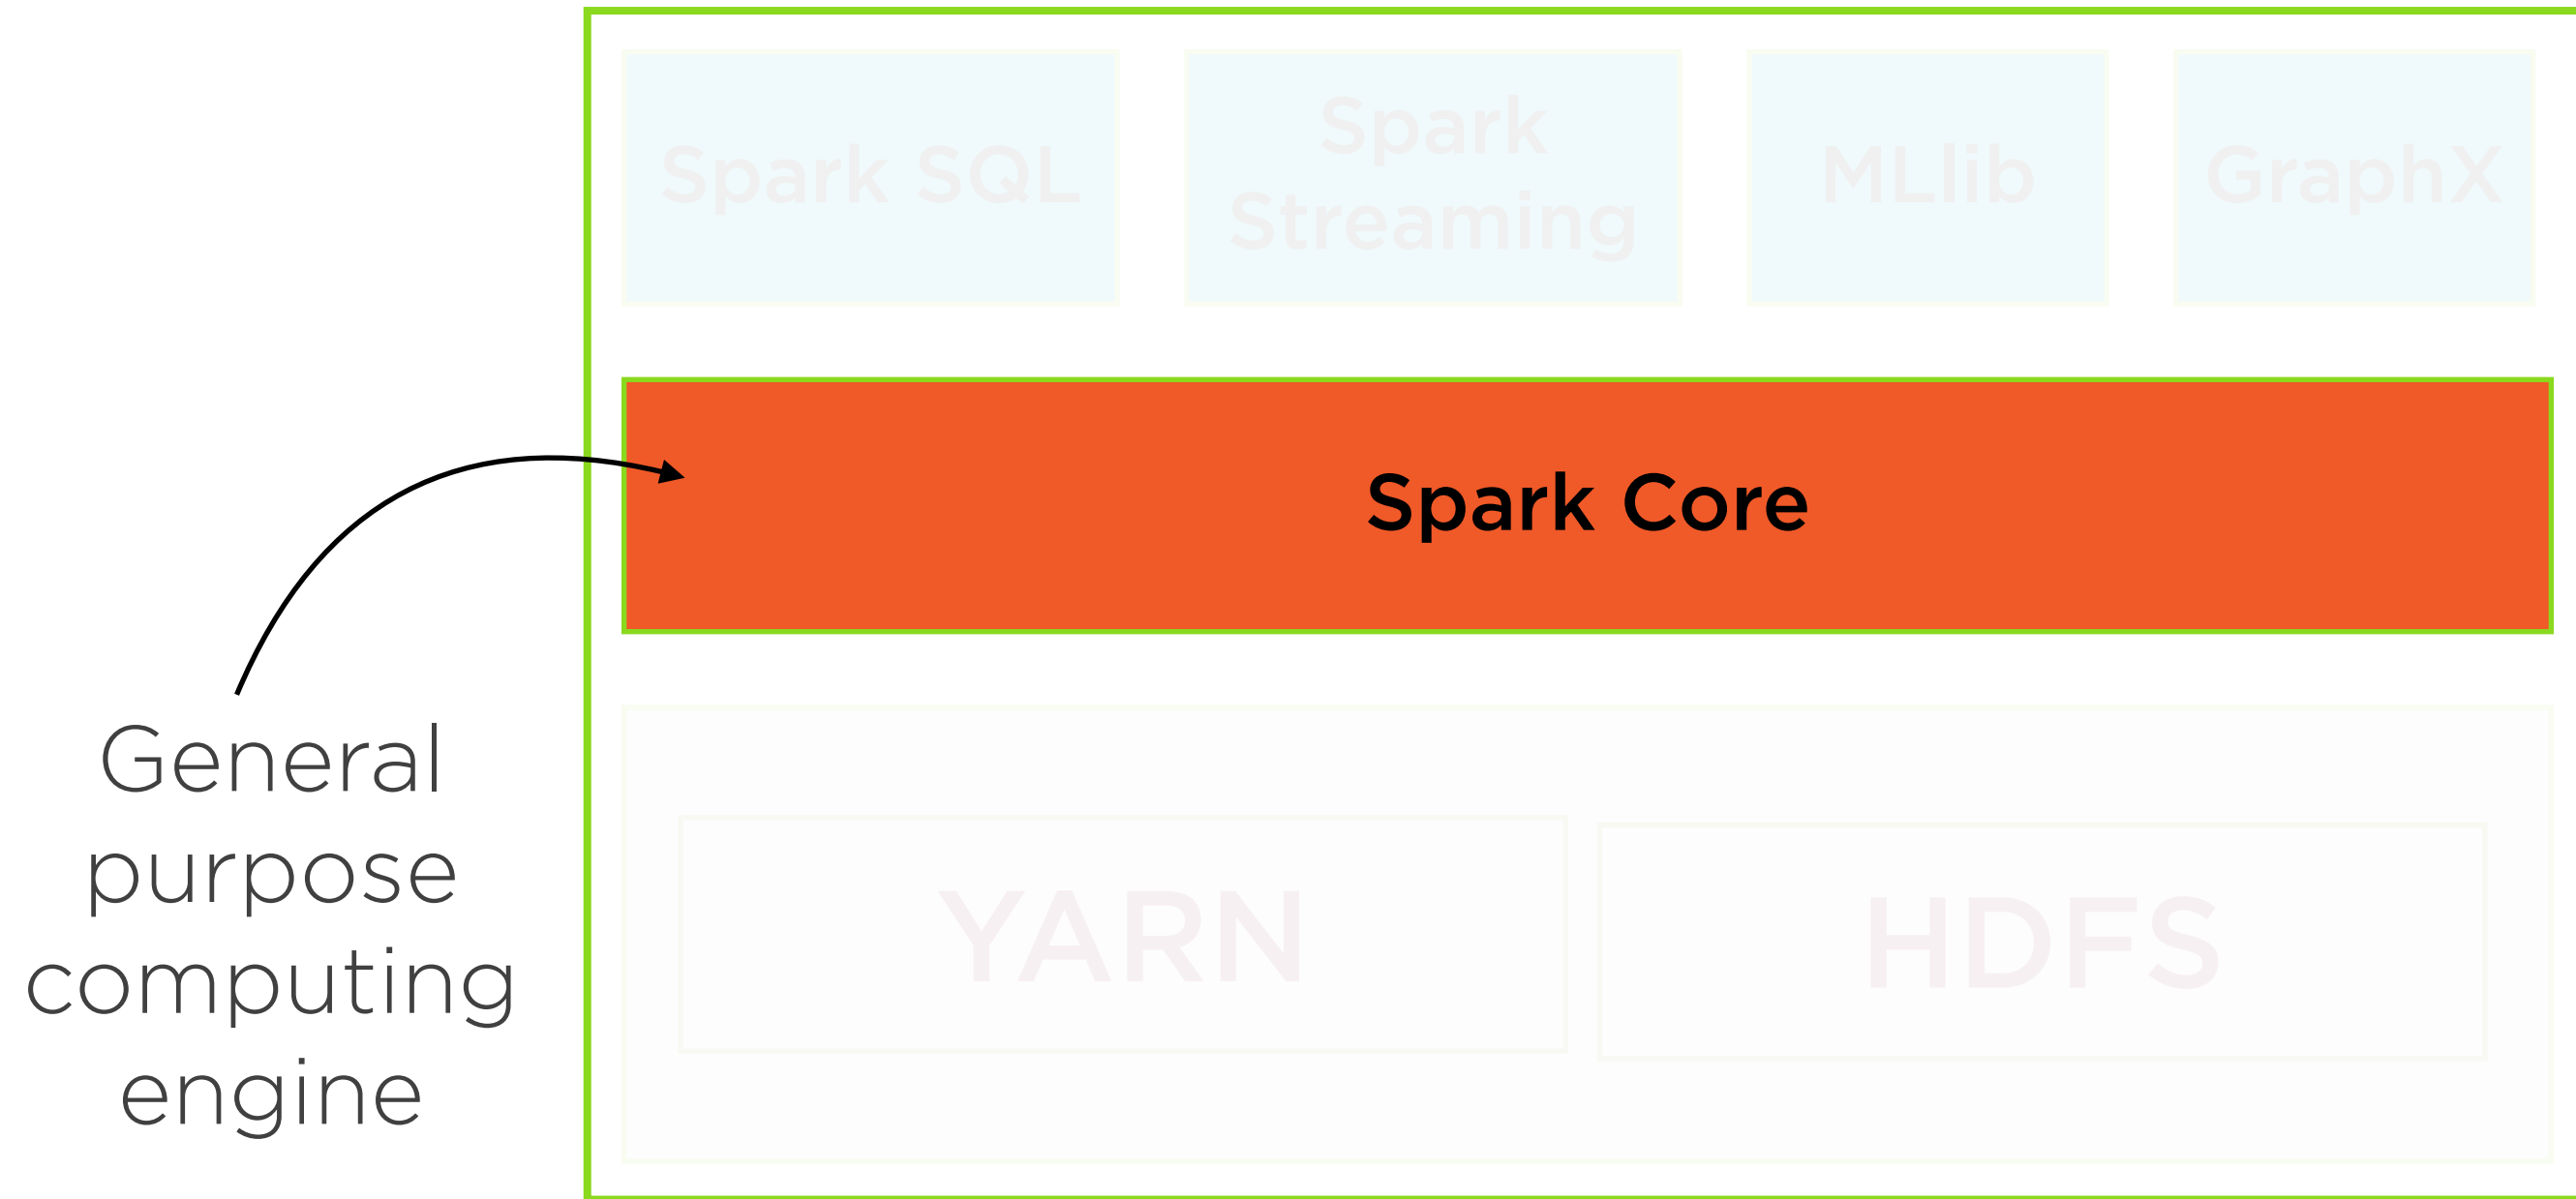
Support for multiple programming languages

- Python, Scala, R, ...

Apache Spark

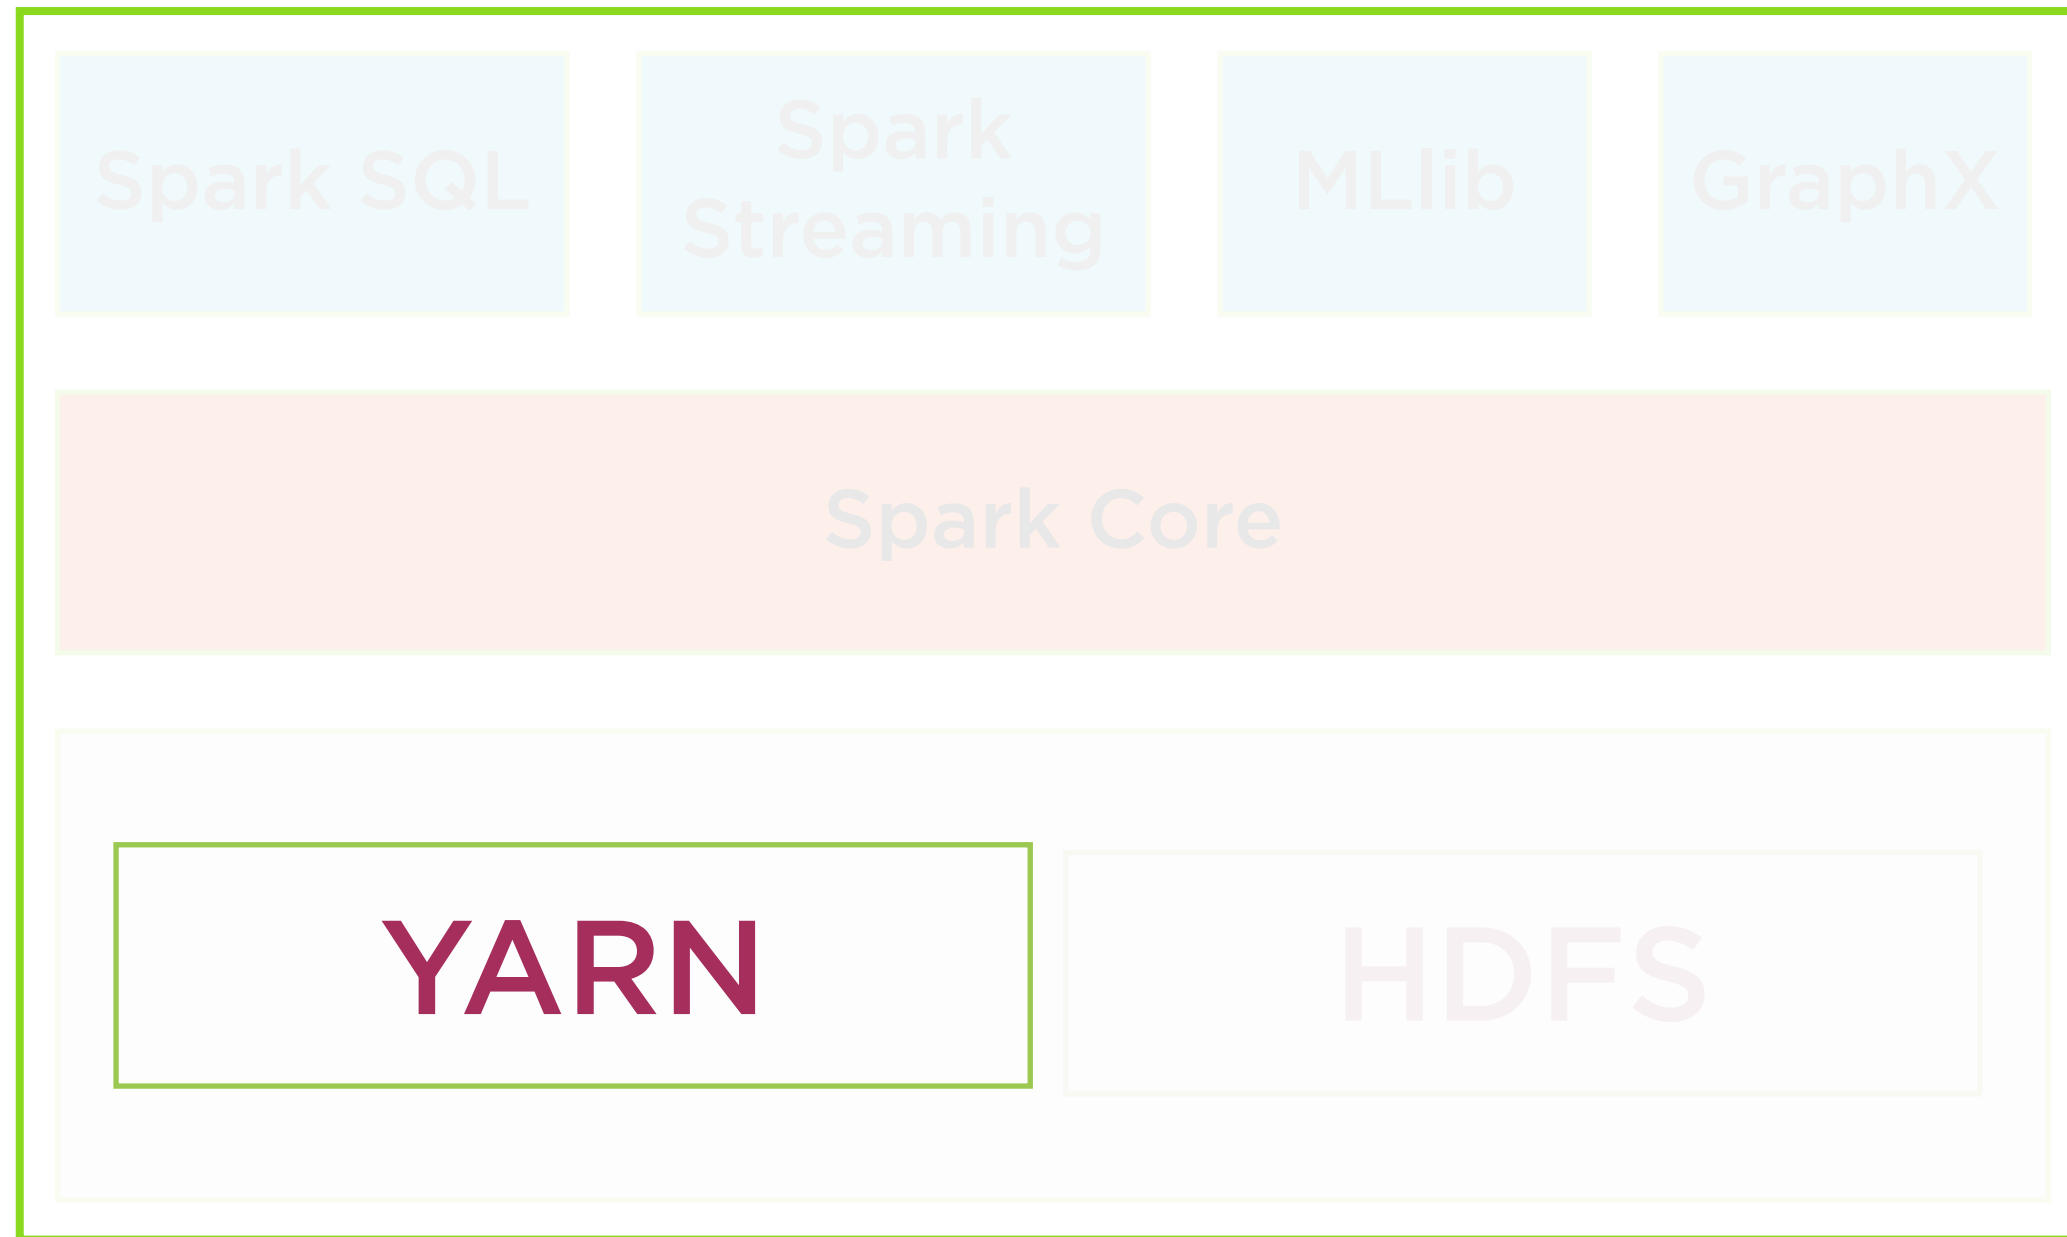


Apache Spark

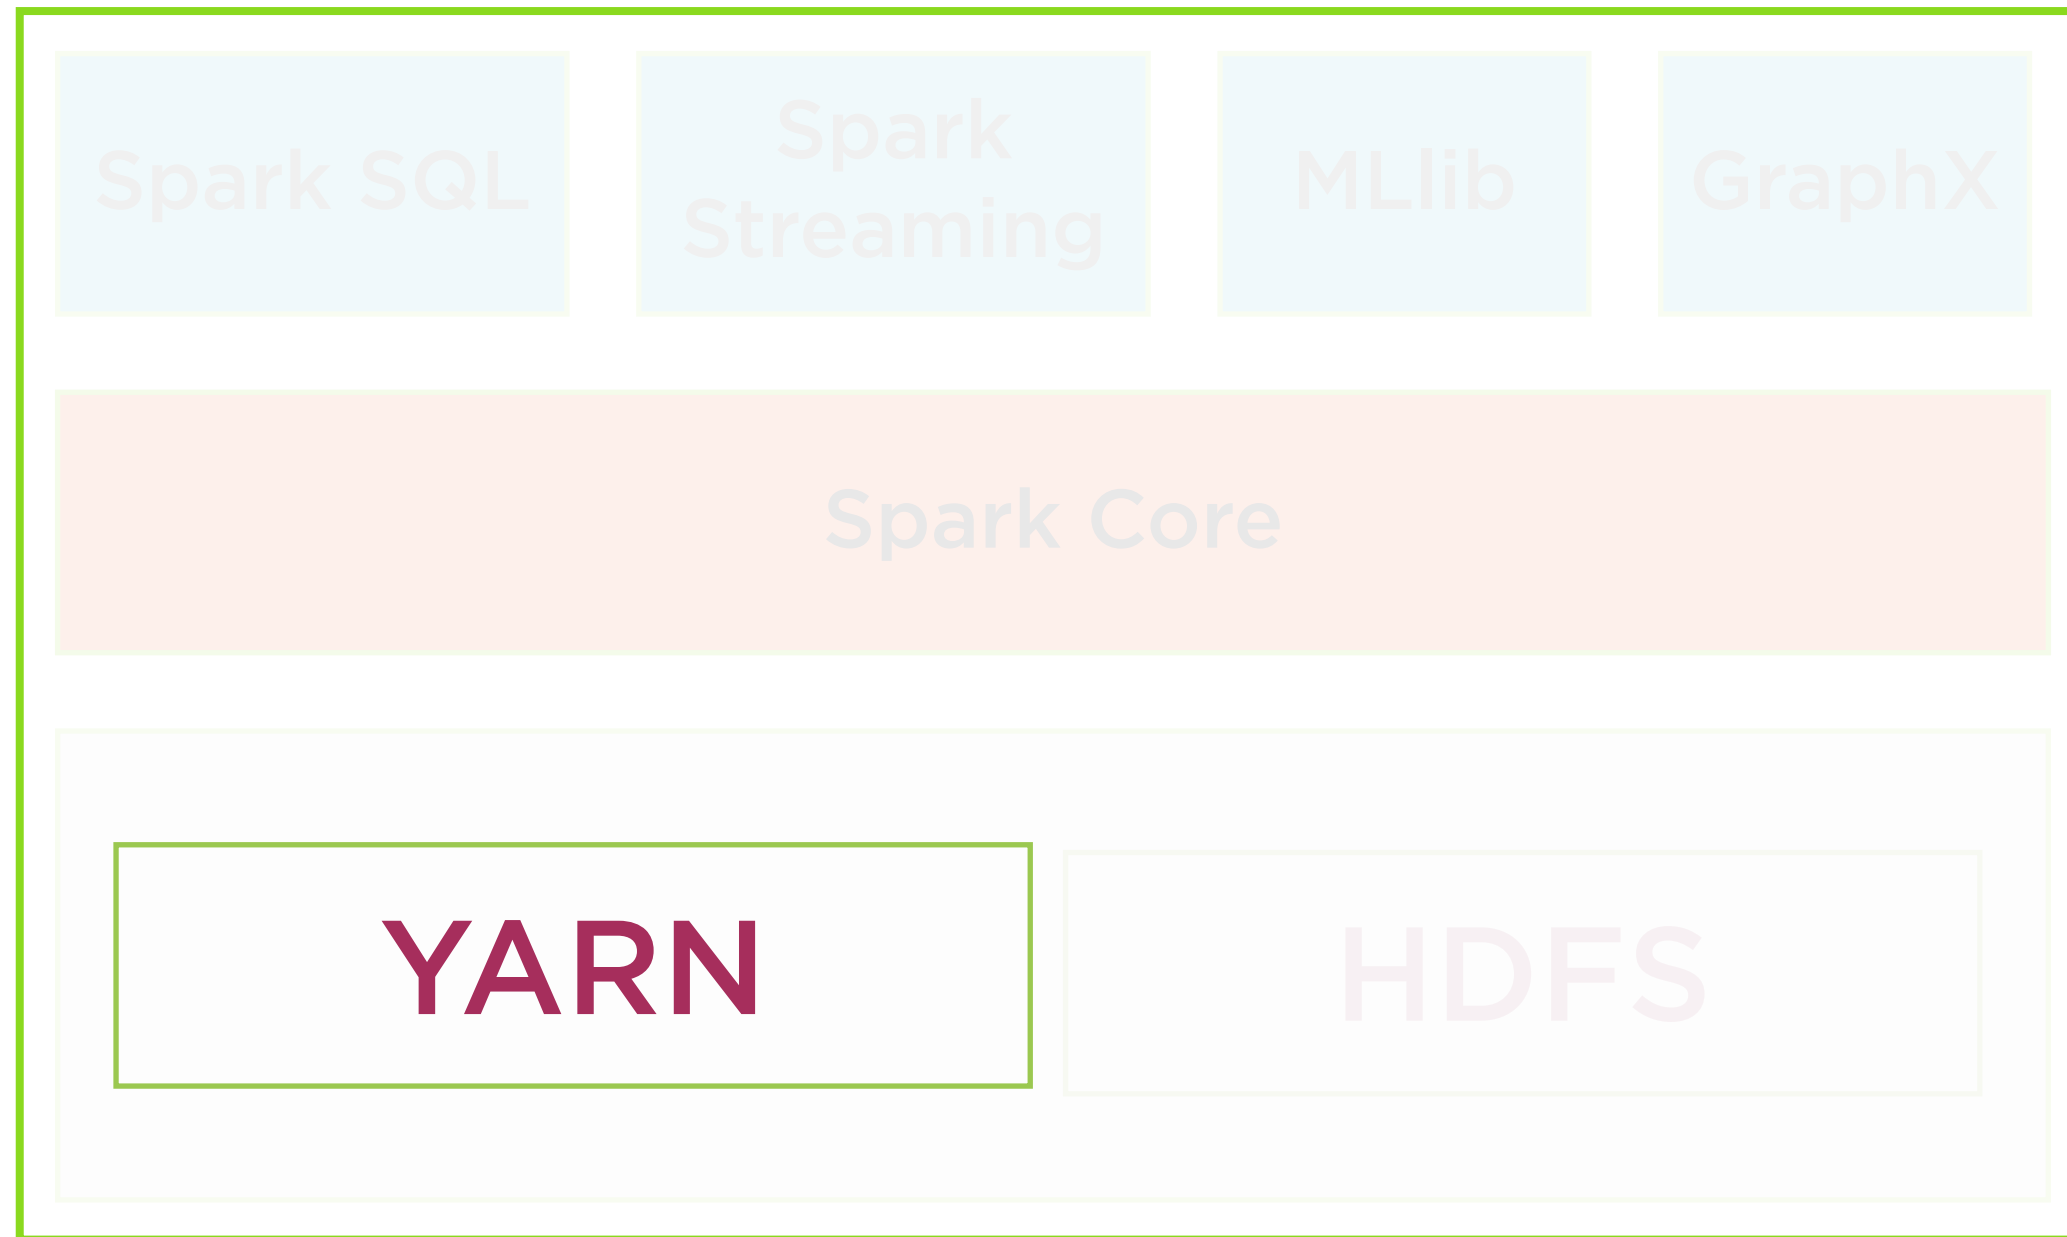


Apache Spark

Cluster
manager

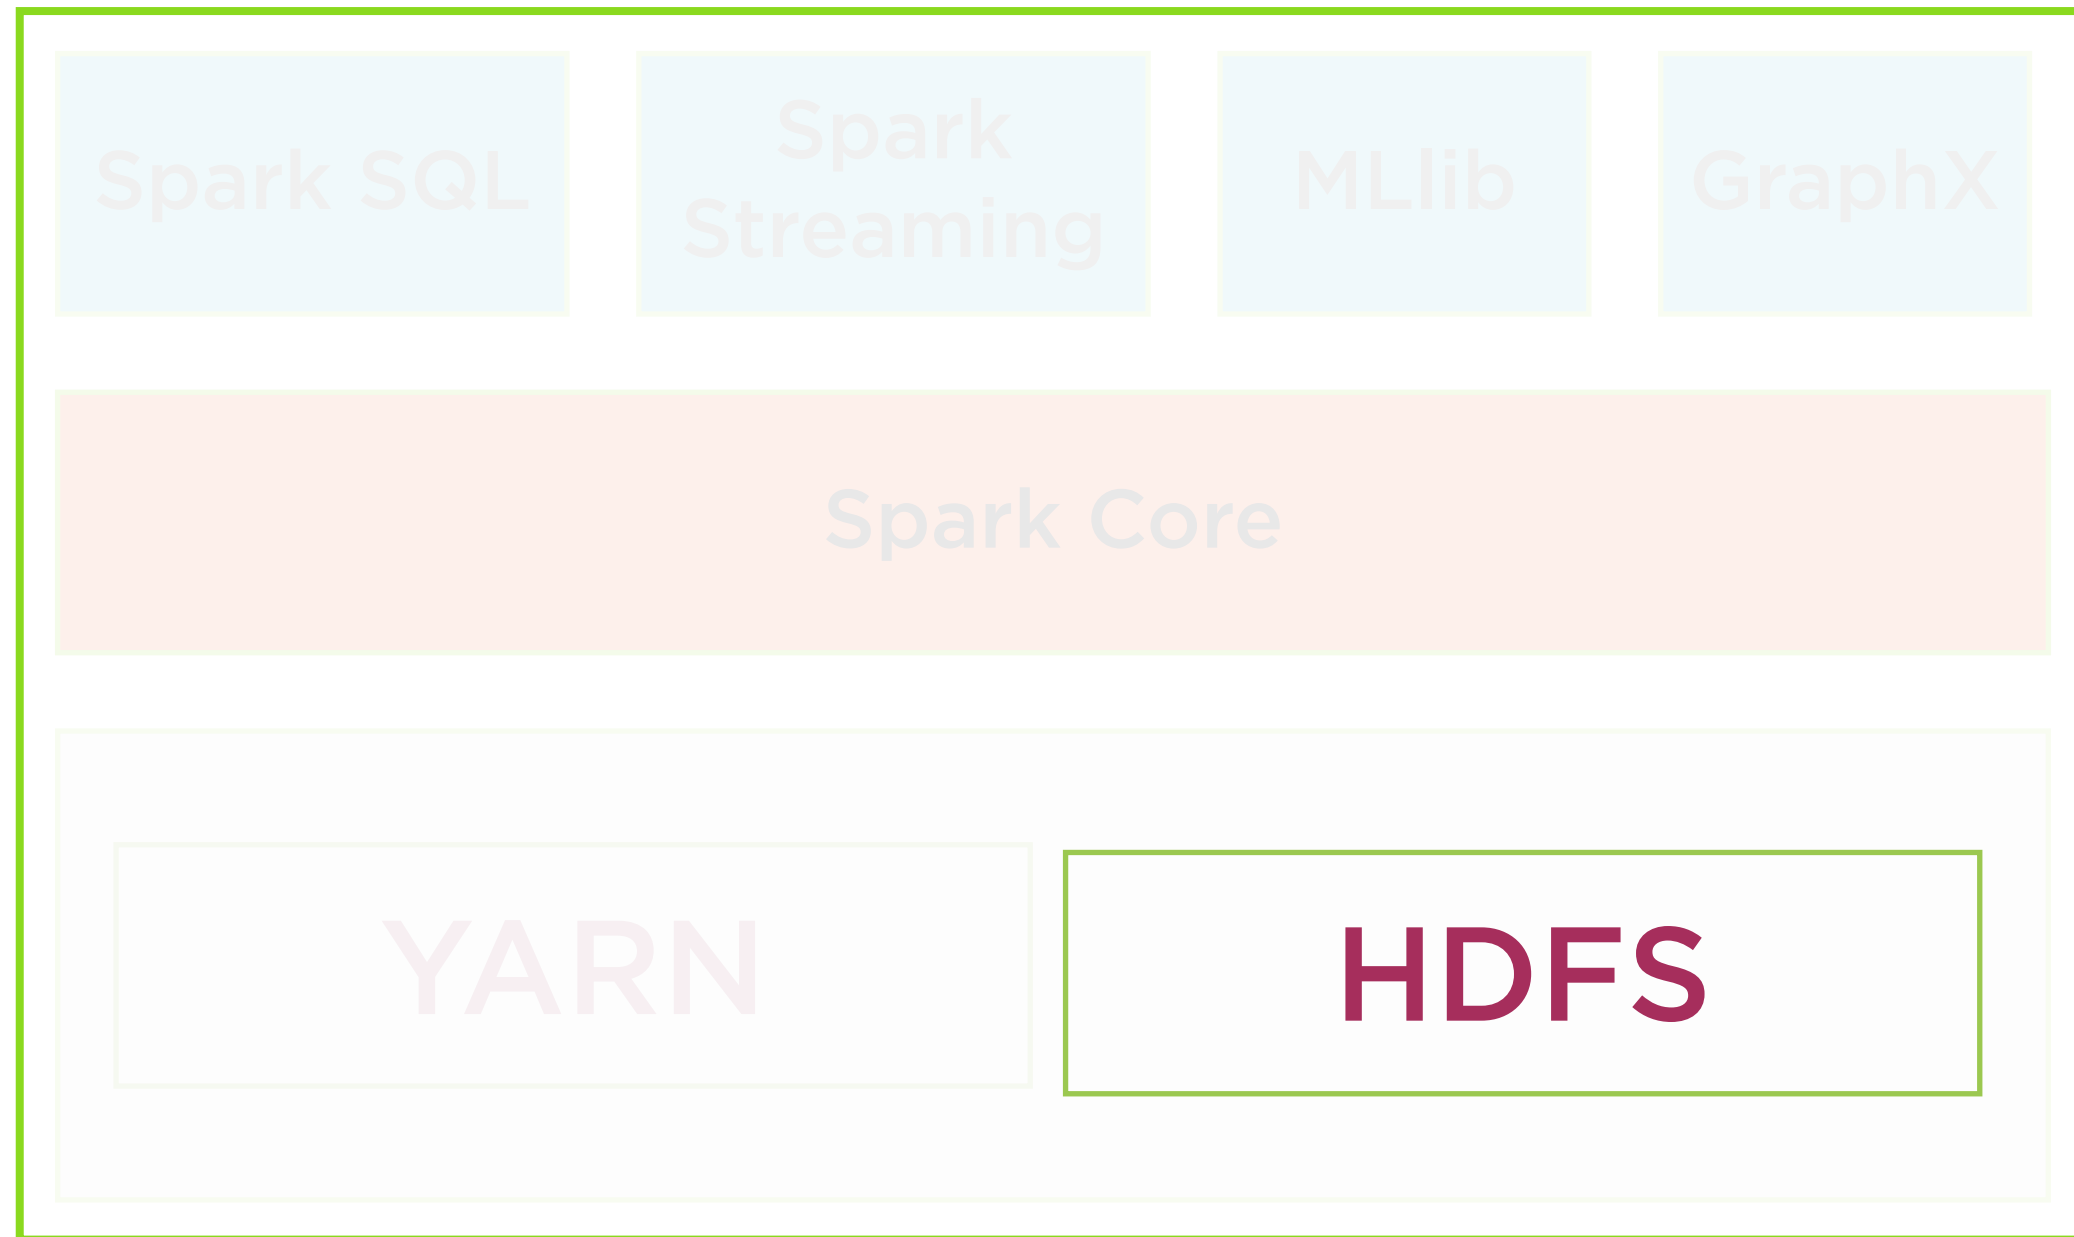


Apache Spark



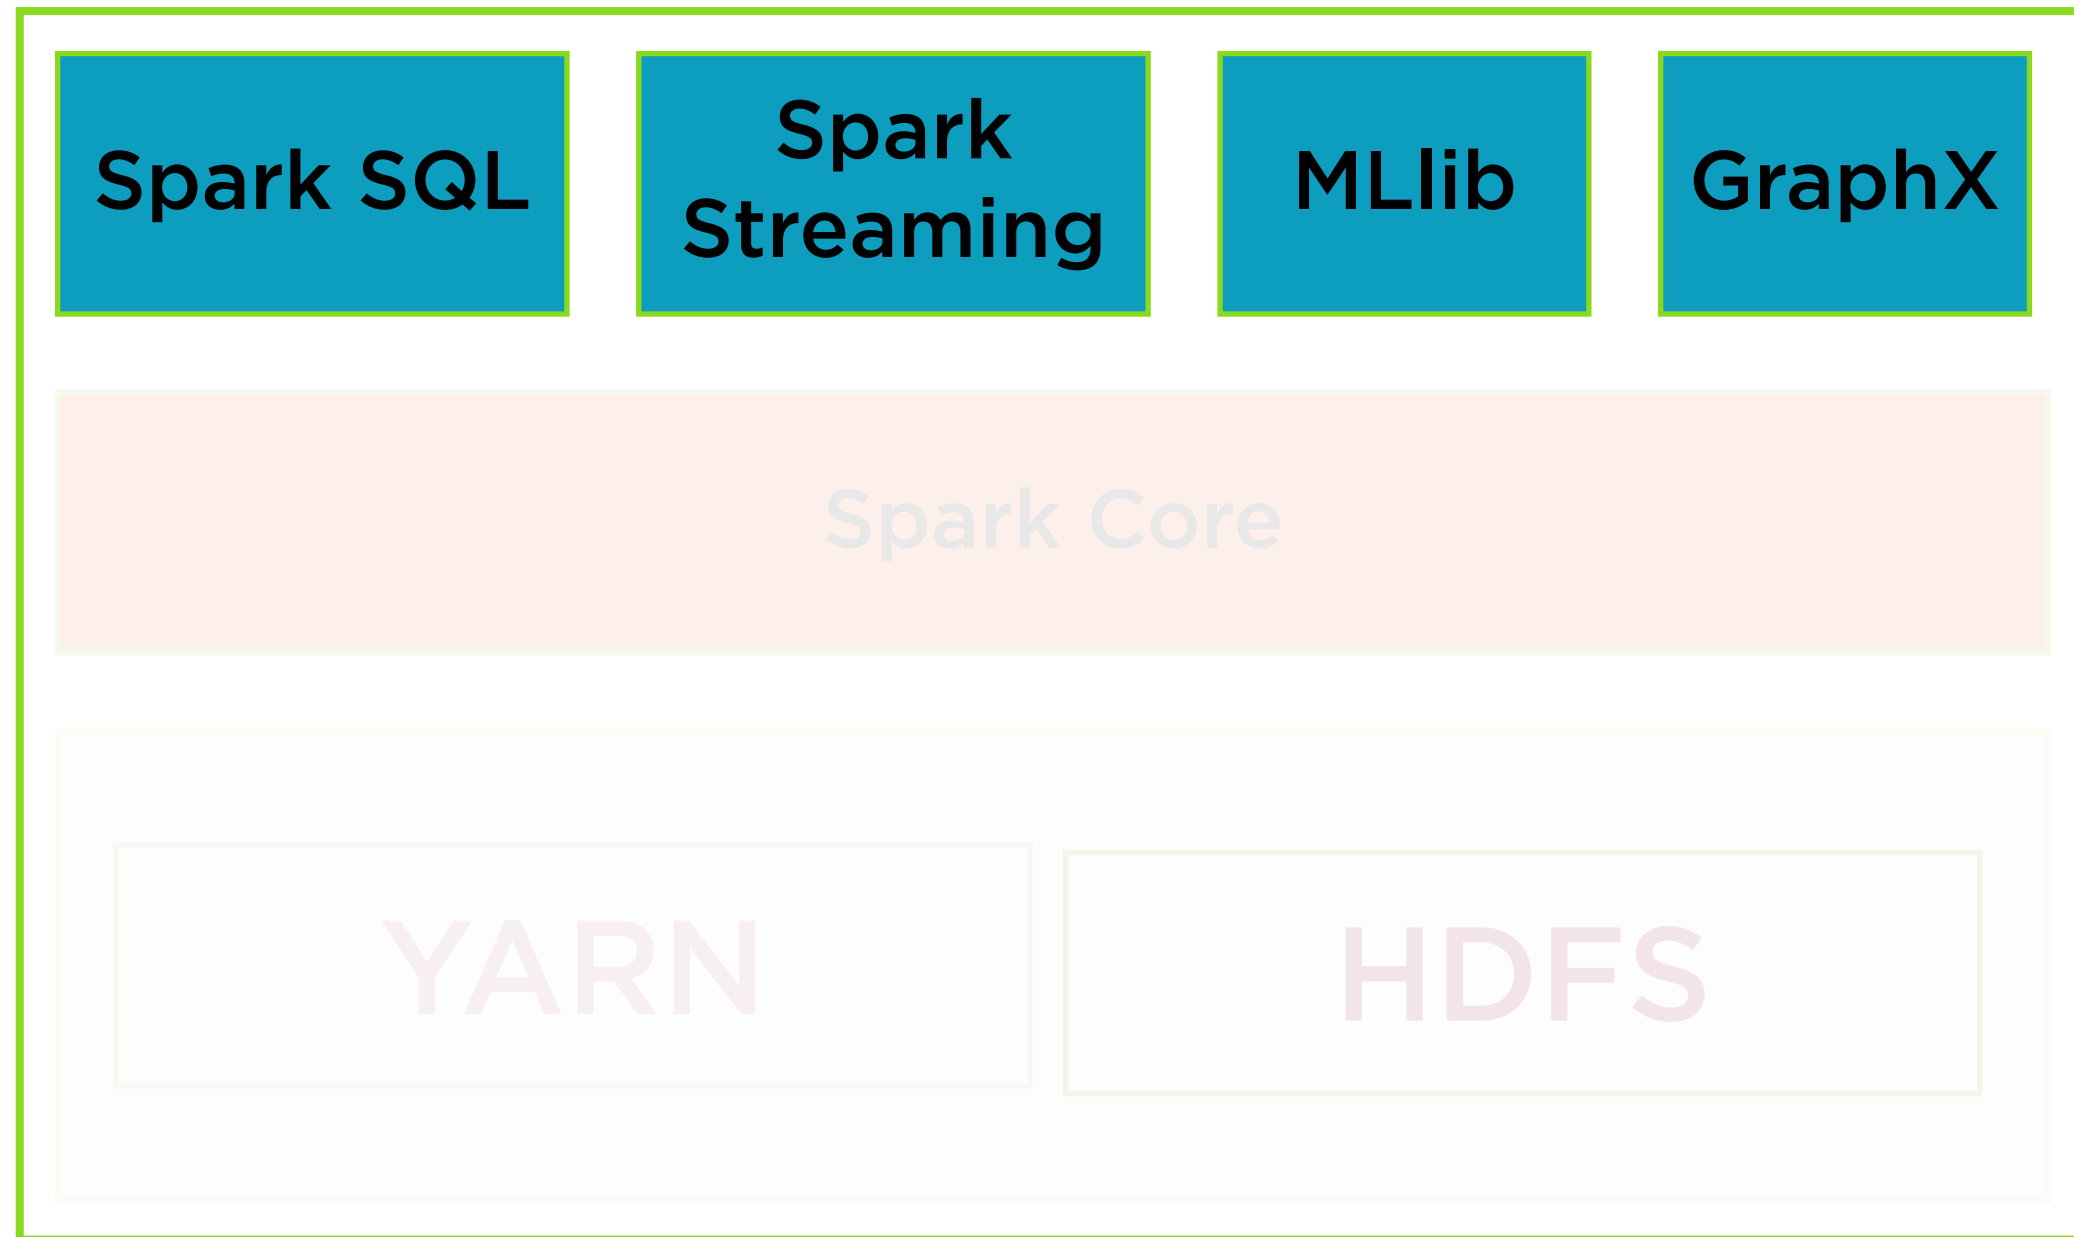
Alternatives:
Mesos, Spark
Standalone,
Kubernetes

Apache Spark



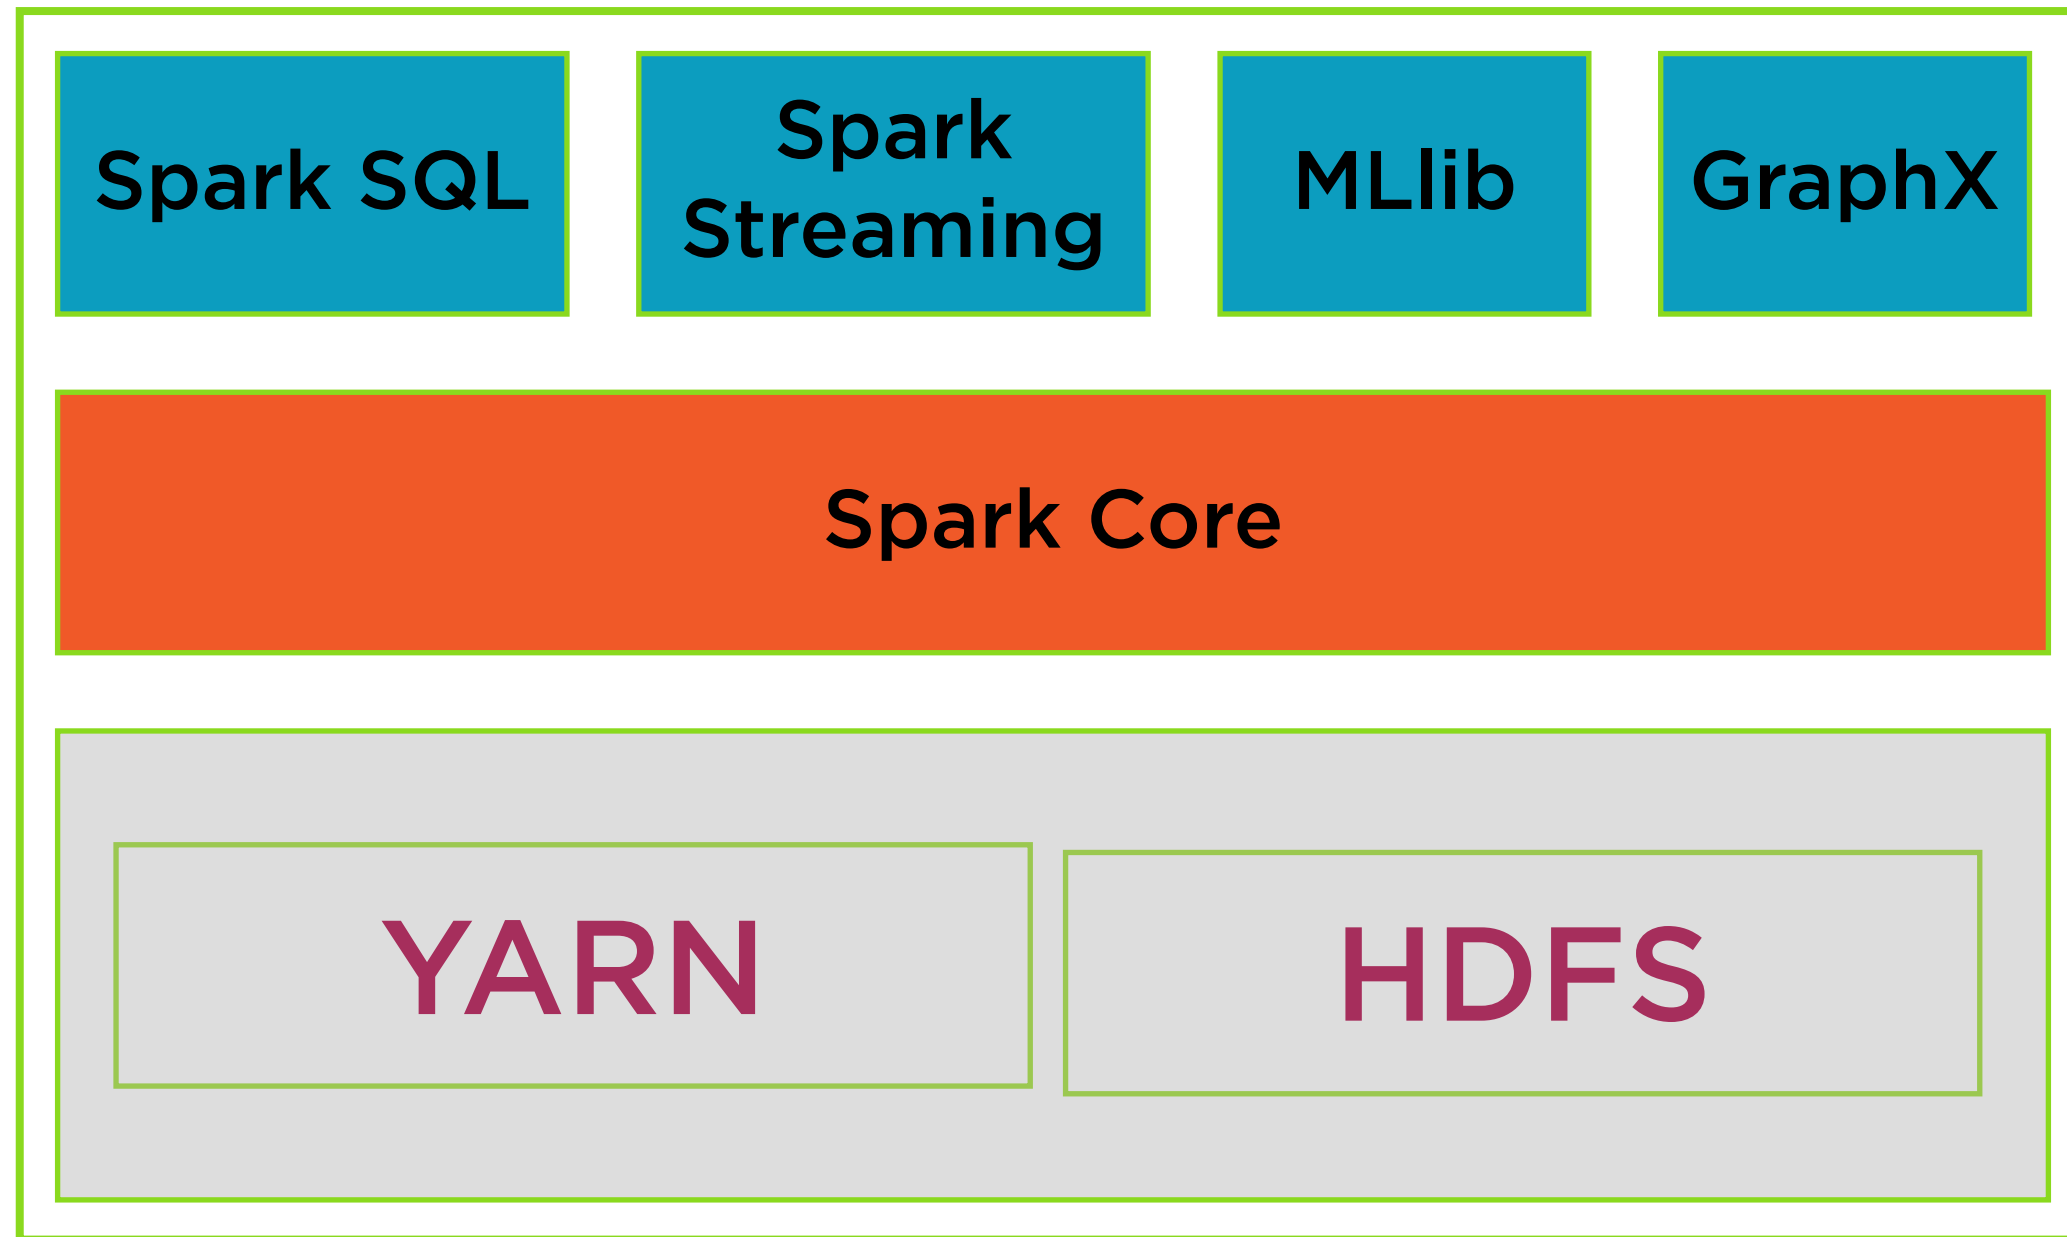
Distributed
Storage
system

Apache Spark



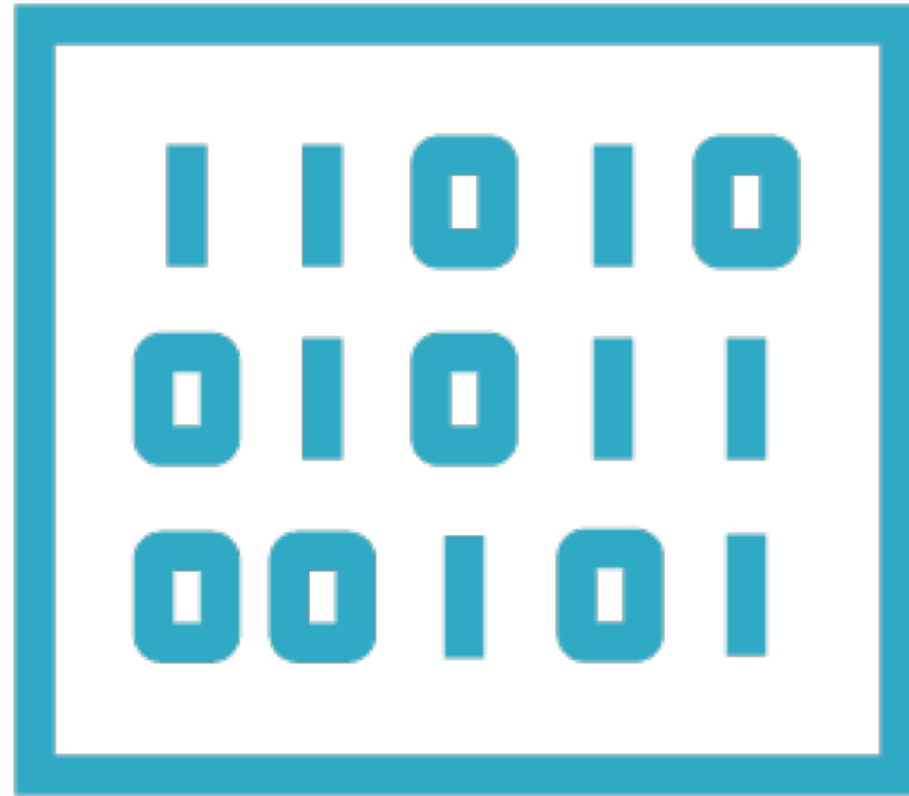
Spark
libraries

Apache Spark



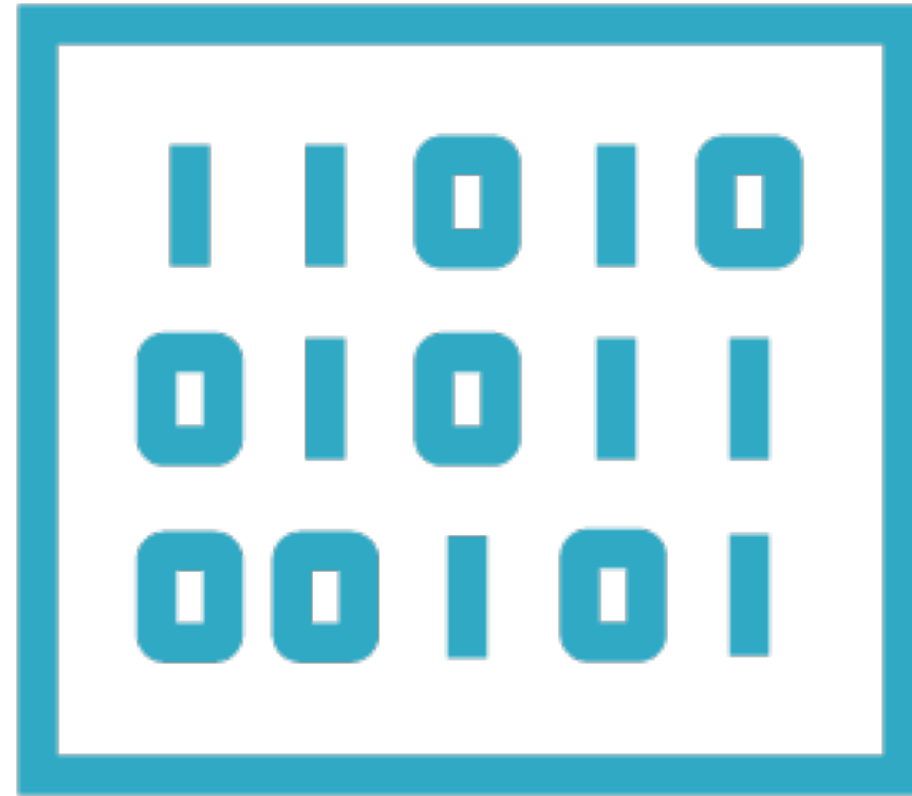
RDDs and DataFrames

Resilient Distributed Datasets



All operations in Spark are performed
on **in-memory objects**

Resilient Distributed Datasets



An RDD is a **collection** of entities - rows, records, an RDD is the basic data structure used in Spark 1.x

Why is this relevant in Spark 2?

RDDs are still the **fundamental building blocks** of Spark

Characteristics of RDDs

Partitioned

RDDs are split
across nodes in a
cluster

Immutable

RDDs, once
created, cannot be
changed

Resilient

Can be
reconstructed on
node crashes

RDDs Support Two Operations

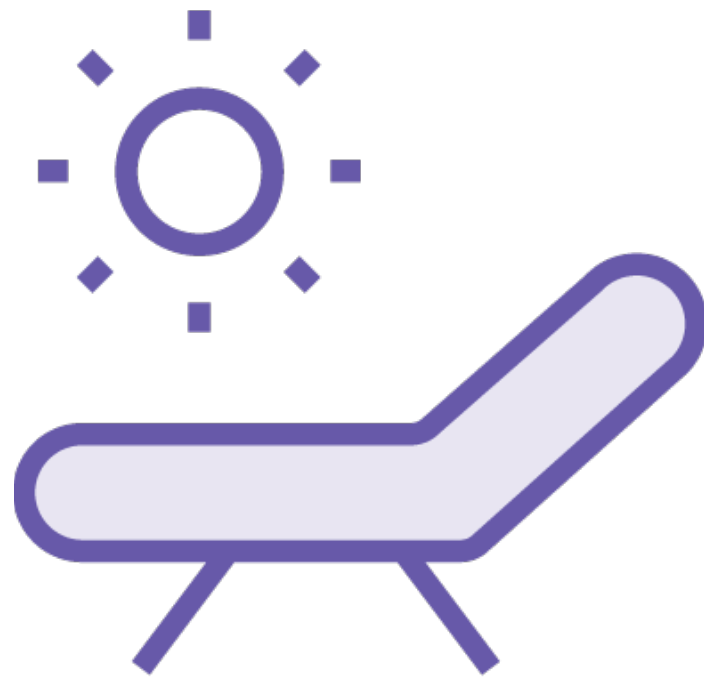
Transformation

Transform input RDDs into
another RDD

Action

Request a result, to a file, to
console window

Lazy Evaluation

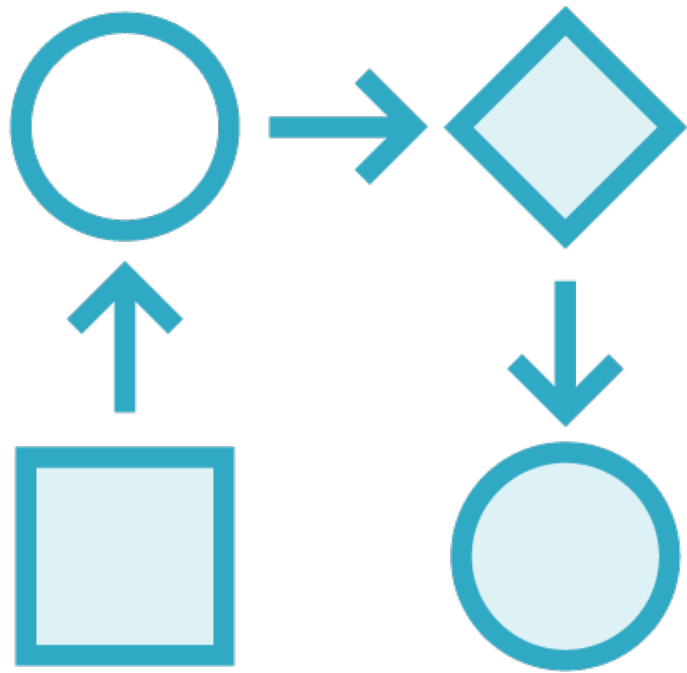


Spark keeps a record of the series of transformations

Transformations are not performed when defined

Transformations are materialized only when the user **requests a result**

Lineage



The record of transformations is called lineage

Allows RDDs to be reconstructed in case of node crashes

The basic data structure for
records in Spark 2.x is the
DataFrame

DataFrame: Data in Rows and Columns

| DATE | OPEN | ... | PRICE |
|------------|------|-----|-------|
| 2016-12-01 | 772 | ... | 779 |
| 2016-11-01 | 758 | ... | 747 |
| | | | |
| | | | |
| | | | |
| 2006-01-01 | 302 | ... | 309 |

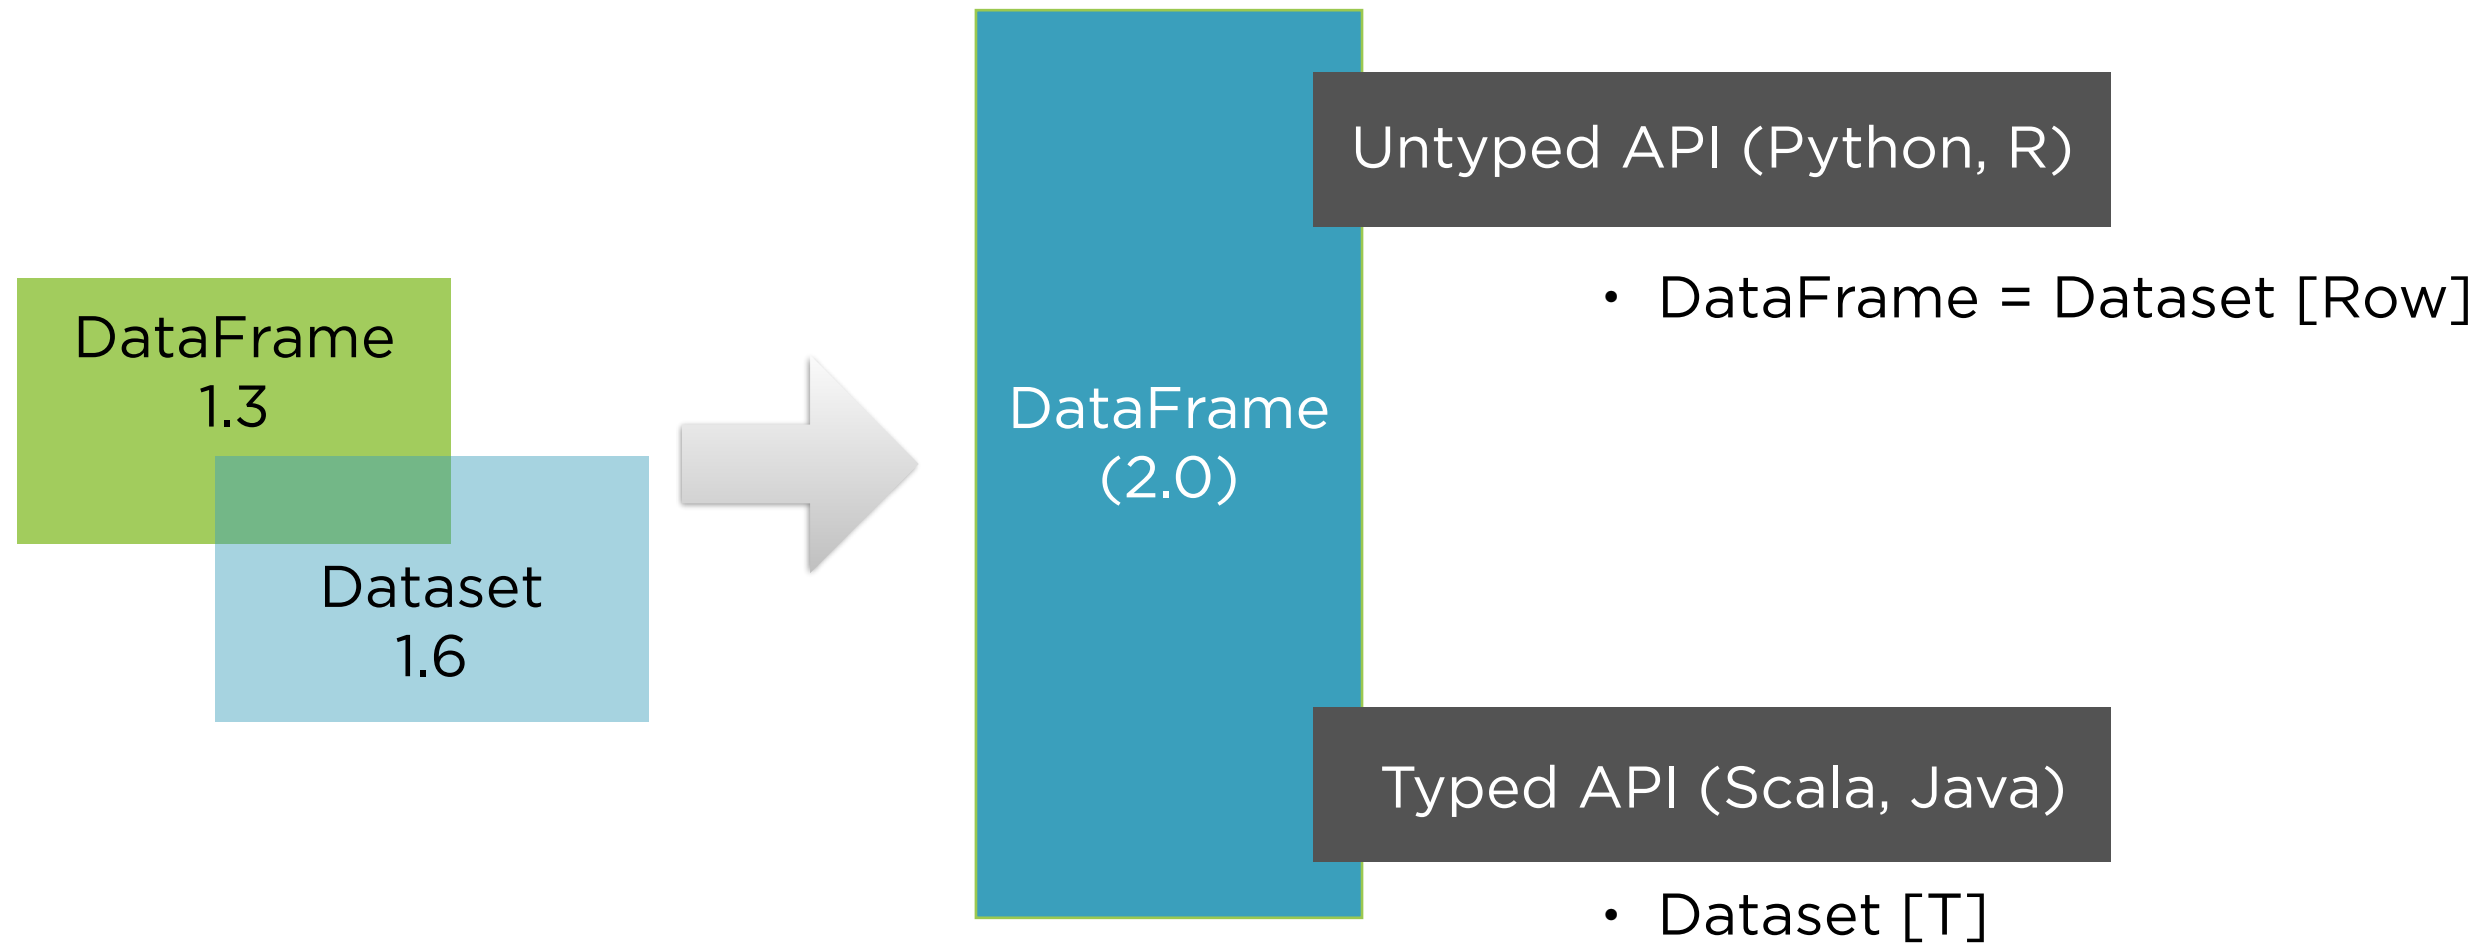
Each row represents
1 observation

DataFrame: Data in Rows and Columns

Each column
represents 1 variable
(a list or vector)

| DATE | OPEN | ... | PRICE |
|------------|------|-----|-------|
| 2016-12-01 | 772 | ... | 779 |
| 2016-11-01 | 758 | ... | 747 |
| | | | |
| | | | |
| | | | |
| 2006-01-01 | 302 | ... | 309 |

Unified API for DataFrames



DataFrames Built on Top of RDDs

Partitioned

RDDs are split
across nodes in a
cluster

Immutable

RDDs, once
created, cannot be
changed

Resilient

Can be
reconstructed on
node crashes

Streaming in Spark 1.x and Spark 2.x

Streaming Data Spark 1.x

```
2016-12-30 09:09:57,862 INFO
org.apache.hadoop.http.HttpServer2: Jetty bound to port
56745
2016-12-30 09:09:57,862 INFO org.mortbay.log: jetty-6.1.26
2016-12-30 09:09:58,037 INFO org.mortbay.log: Started
HttpServer2$SelectChannelConnectorWithSafeStartup@localhost:
56745
2016-12-30 09:09:58,124 INFO
org.apache.hadoop.hdfs.server.datanode.web.DatanodeHttpServe
r: Listening HTTP traffic on /0.0.0.0:50075
2016-12-30 09:09:58,239 INFO
```

Streaming Data Spark 1.x

```
2016-12-30 09:09:58,239 INFO
```

```
org.apache.hadoop.hdfs.server.datanode.web.DatanodeHttpServer: Listening HTTP traffic on /0.0.0.0:50075
```

```
HttpServer2$SelectChannelConnectorWithSafeStartup@localhost:56745
```

```
2016-12-30 09:09:58,037 INFO org.mortbay.log: Started
```

```
2016-12-30 09:09:57,862 INFO org.mortbay.log: jetty-6.1.26
```

```
2016-12-30 09:09:57,862 INFO org.apache.hadoop.p.http.HttpServer2: Jetty bound to port 56745
```

Each log message is **one entity** in this stream

Streaming Data Spark 1.x

```
2016-12-30 09:09:58,239 INFO
```

```
org.apache.hadoop.hdfs.server.datanode.web.DatanodeHttpServer: Listening HTTP traffic on /0.0.0.0:50075
```

```
HttpServer2$SelectChannelConnectorWithSafeStartup@localhost:56745
```

```
2016-12-30 09:09:58,037 INFO org.mortbay.log: Started
```

```
2016-12-30 09:09:57,862 INFO org.mortbay.log: jetty-6.1.26
```

```
2016-12-30 09:09:57,862 INFO org.apache.hadoop.p.http.HttpServer2: Jetty bound to port 56745
```

Spark works with stream data using the same **batch RDD abstraction**

Streaming Data Spark 1.x

2016-12-30 09:09:58,239 INFO

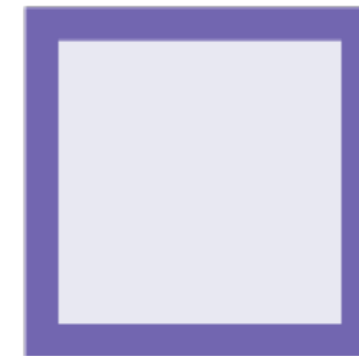
org.apache.hadoop.hdfs.server.datanode.web.DatanodeHttpServer: Listening HTTP traffic on /0.0.0.0:50075

HttpServer2\$SelectChannelConnectorWithSafeStartup@localhost:56745

2016-12-30 09:09:58,037 INFO org.mortbay.log: Started

2016-12-30 09:09:57,862 INFO org.mortbay.log: jetty-6.1.26

2016-12-30 09:09:57,862 INFO org.apache.hadoop.p.http.HttpServer2: Jetty bound to port 56745



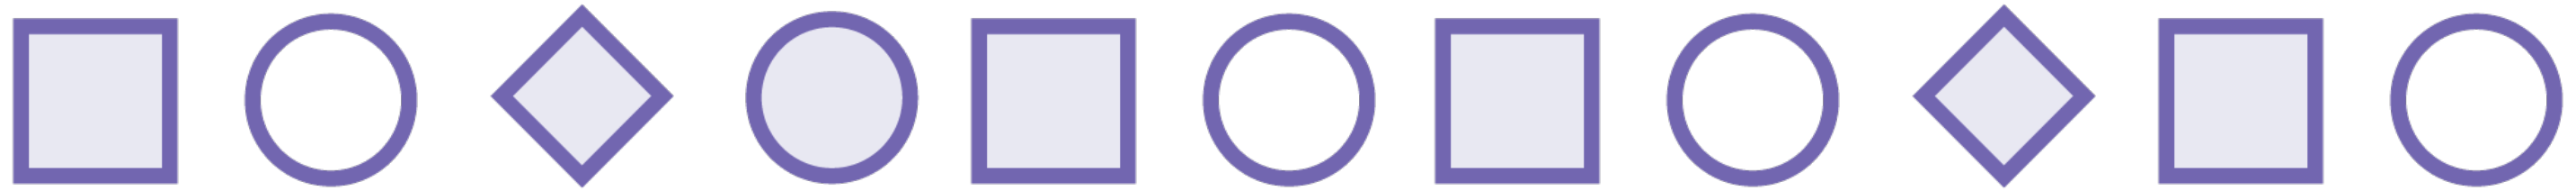
Streaming Data Spark 1.x



This stream of entities is represented as a
discretized stream or DStream

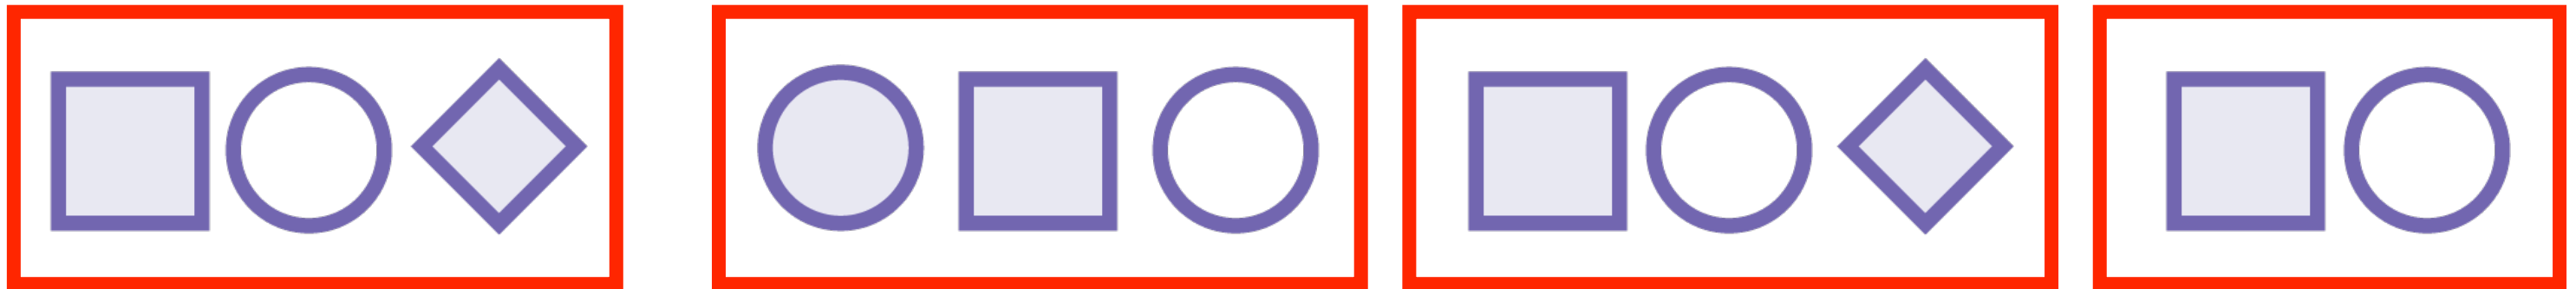
DStream = Sequence of RDDs

Streaming Data Spark 1.x



Entities => RDDs => DStream ?

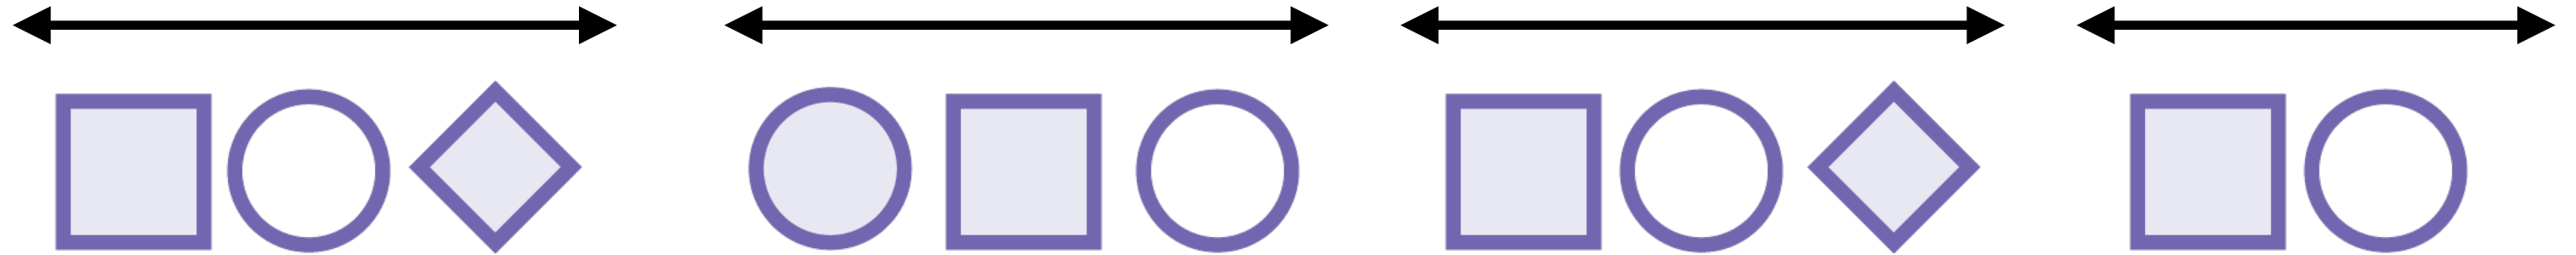
Streaming Data Spark 1.x



Entities in a stream are grouped into **batches**

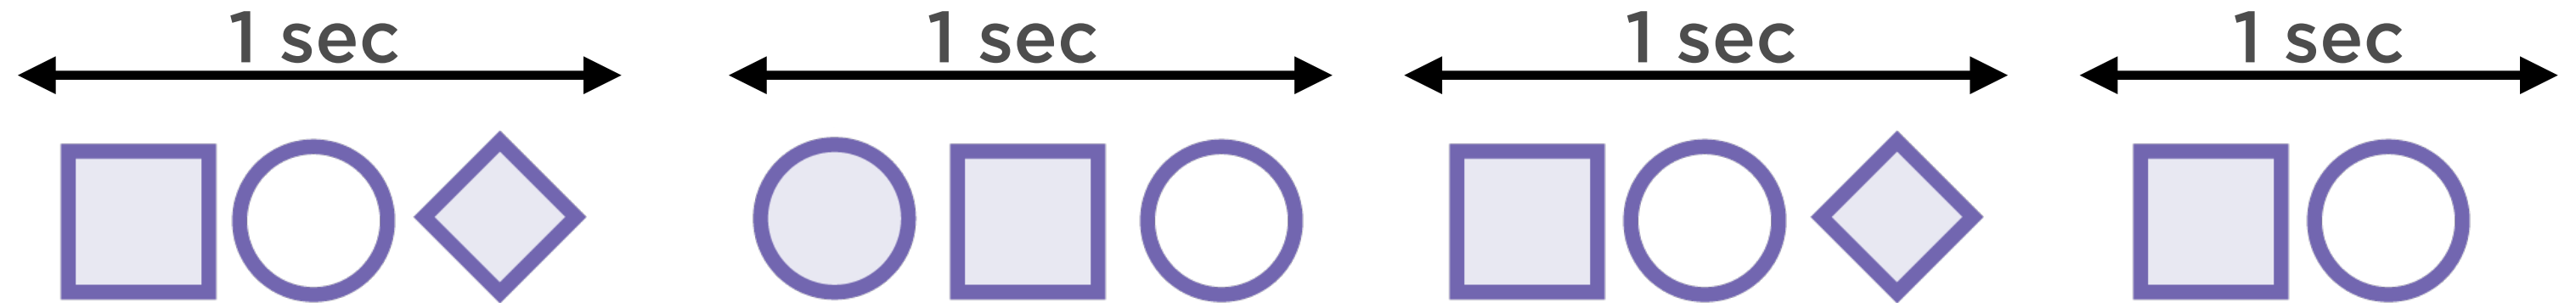
Each batch = 1 RDD

Streaming Data Spark 1.x



Batches are formed based on a **batch interval**

Streaming Data Spark 1.x



All entities received within the batch interval make one RDD

Streaming Data Spark 1.x



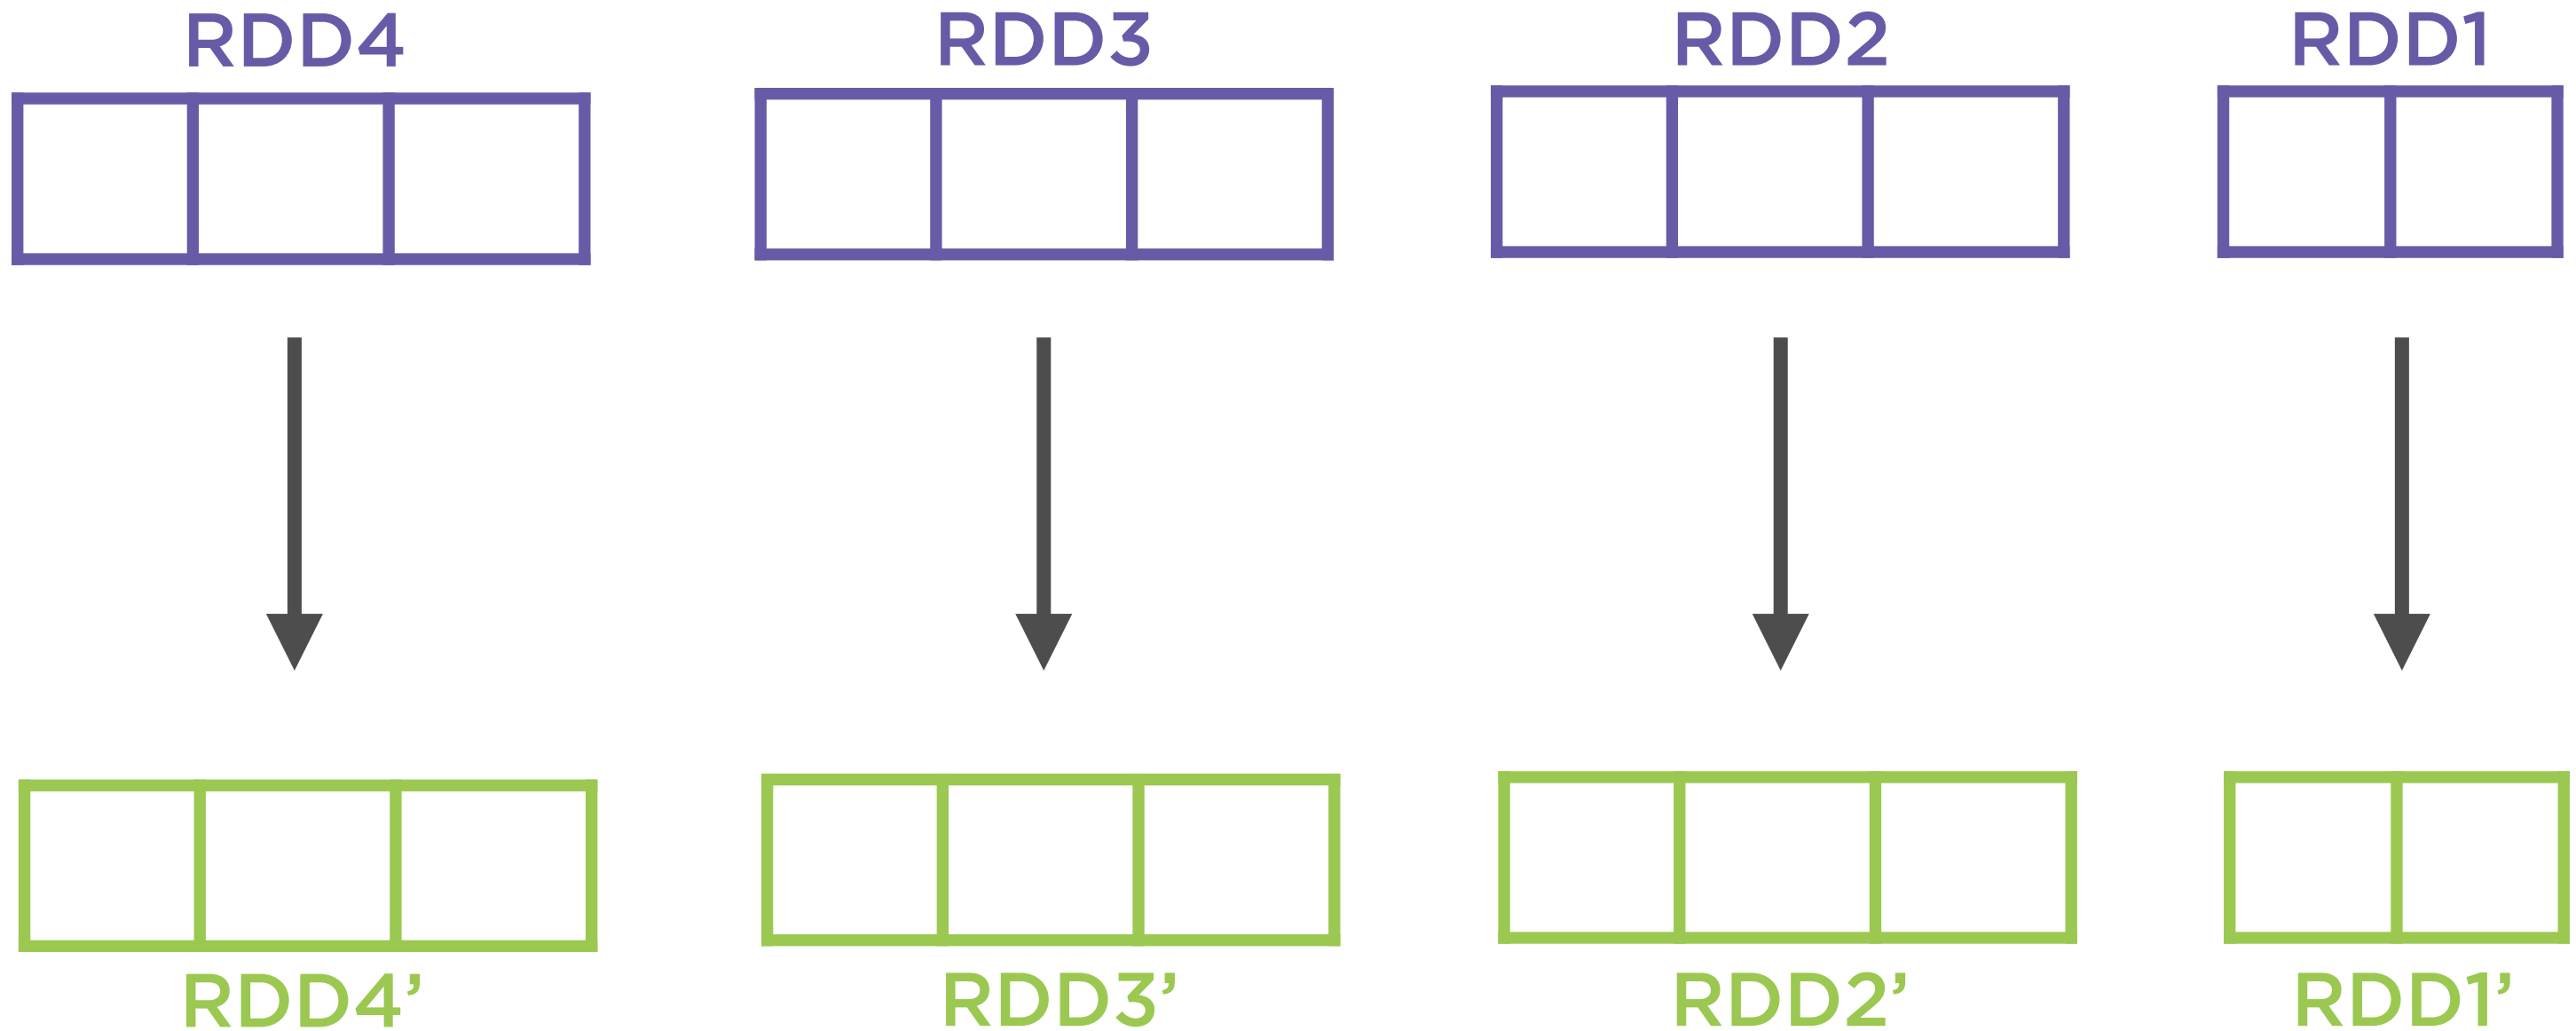
All entities received within the batch interval make one RDD

Streaming Data Spark 1.x

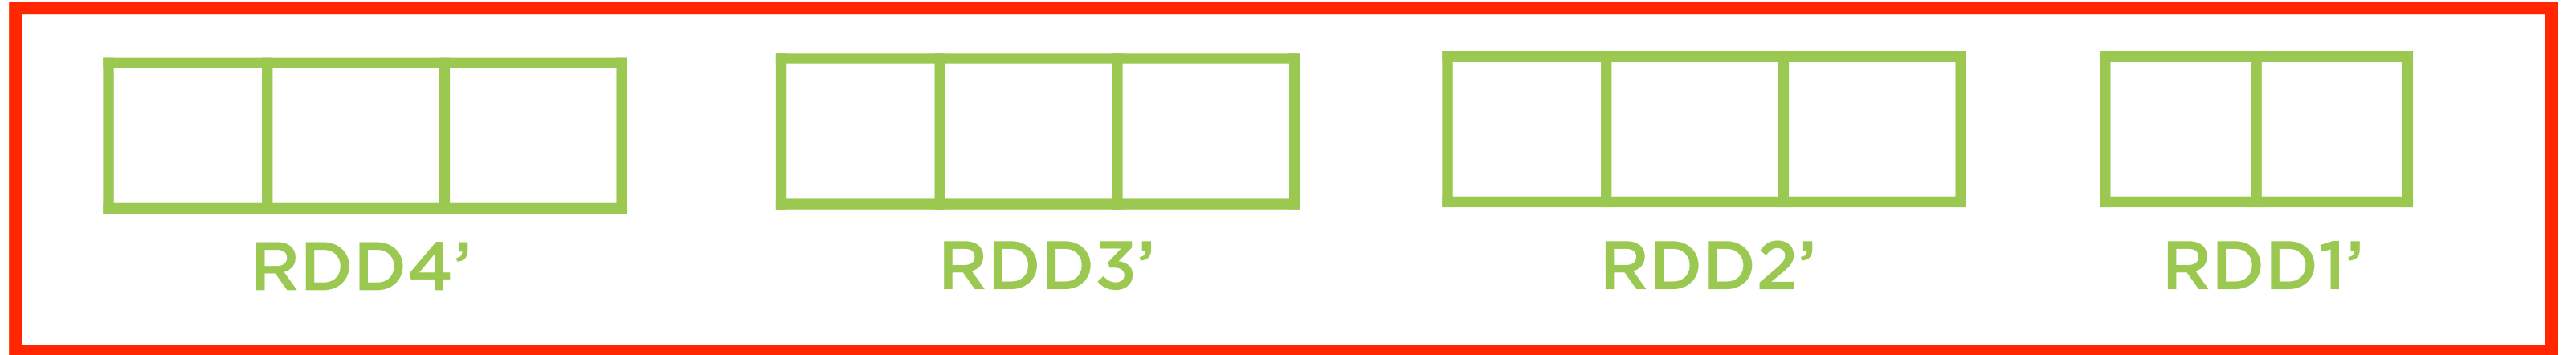


Within a DStream, Spark still performs operations on individual RDDs

Streaming Data Spark 1.x



Streaming Data Spark 1.x



Transformed DStream

Within DStreams, Spark 1.x
does **batch processing on
individual RDDs**

The basic data structure for
records in Spark 2.x is the
DataFrame

DataFrame: Data in Rows and Columns

| DATE | OPEN | ... | PRICE |
|------------|------|-----|-------|
| 2016-12-01 | 772 | ... | 779 |
| 2016-11-01 | 758 | ... | 747 |
| | | | |
| | | | |
| | | | |
| 2006-01-01 | 302 | ... | 309 |

Streaming Data Spark 2.x

Data stream



| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |

**Every data item
that is arriving on
the stream is like a
new row being
appended to the
input table**

Streaming Data Spark 2.x

Data stream



| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |

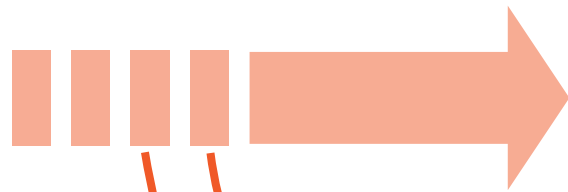
| | | |
|--|--|--|
| | | |
|--|--|--|

Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



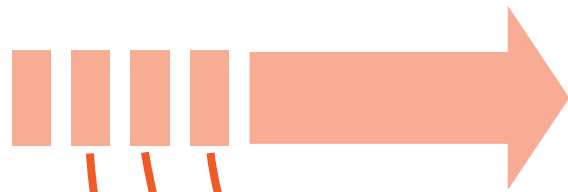
| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |

Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |

| | | |
|--|--|--|
| | | |
|--|--|--|

| | | |
|--|--|--|
| | | |
|--|--|--|

| | | |
|--|--|--|
| | | |
|--|--|--|

Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |

| | | |
|--|--|--|
| | | |
|--|--|--|

| | | |
|--|--|--|
| | | |
|--|--|--|

| | | |
|--|--|--|
| | | |
|--|--|--|

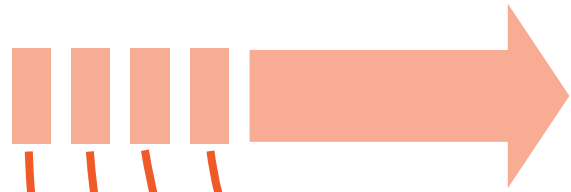
| | | |
|--|--|--|
| | | |
|--|--|--|

Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Batch is Simply A Prefix of Stream

Data stream

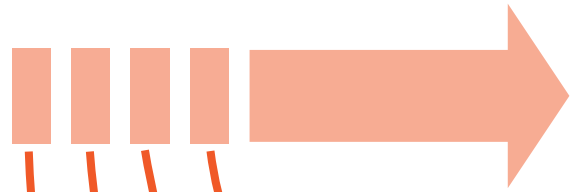


| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |

In other words, the input table (batch) is simply a prefix of the stream

Batch is Simply A Prefix of Stream

Data stream

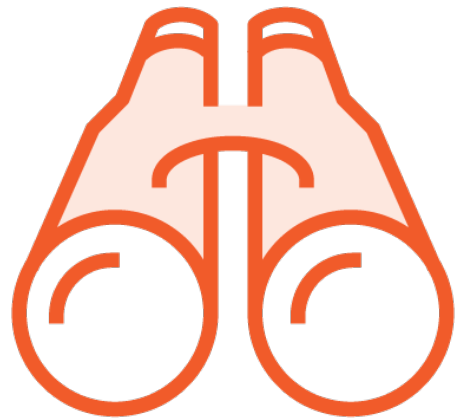


| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |

**All operations that
can be performed on
data frames can be
performed on the
stream**

Structured Streaming treats a live data stream as a table that is being **continuously** appended

Structured Streaming



What

A high-level API that
takes burden off user



How

Micro-batch
processing with
exactly-once fault-
tolerance



Why

Code virtually
identical for batch
and streaming

Structured Streaming

High-level API in Apache Spark 2.0 that supports continuous applications and replaces Spark Streaming

<https://databricks.com/blog/2016/07/28/continuous-applications-evolving-streaming-in-apache-spark-2-0.html>

Streaming and Structured Streaming

Streaming

Older

RDDs

No optimizations

Batch and streaming support not unified

Structured Streaming

Newer

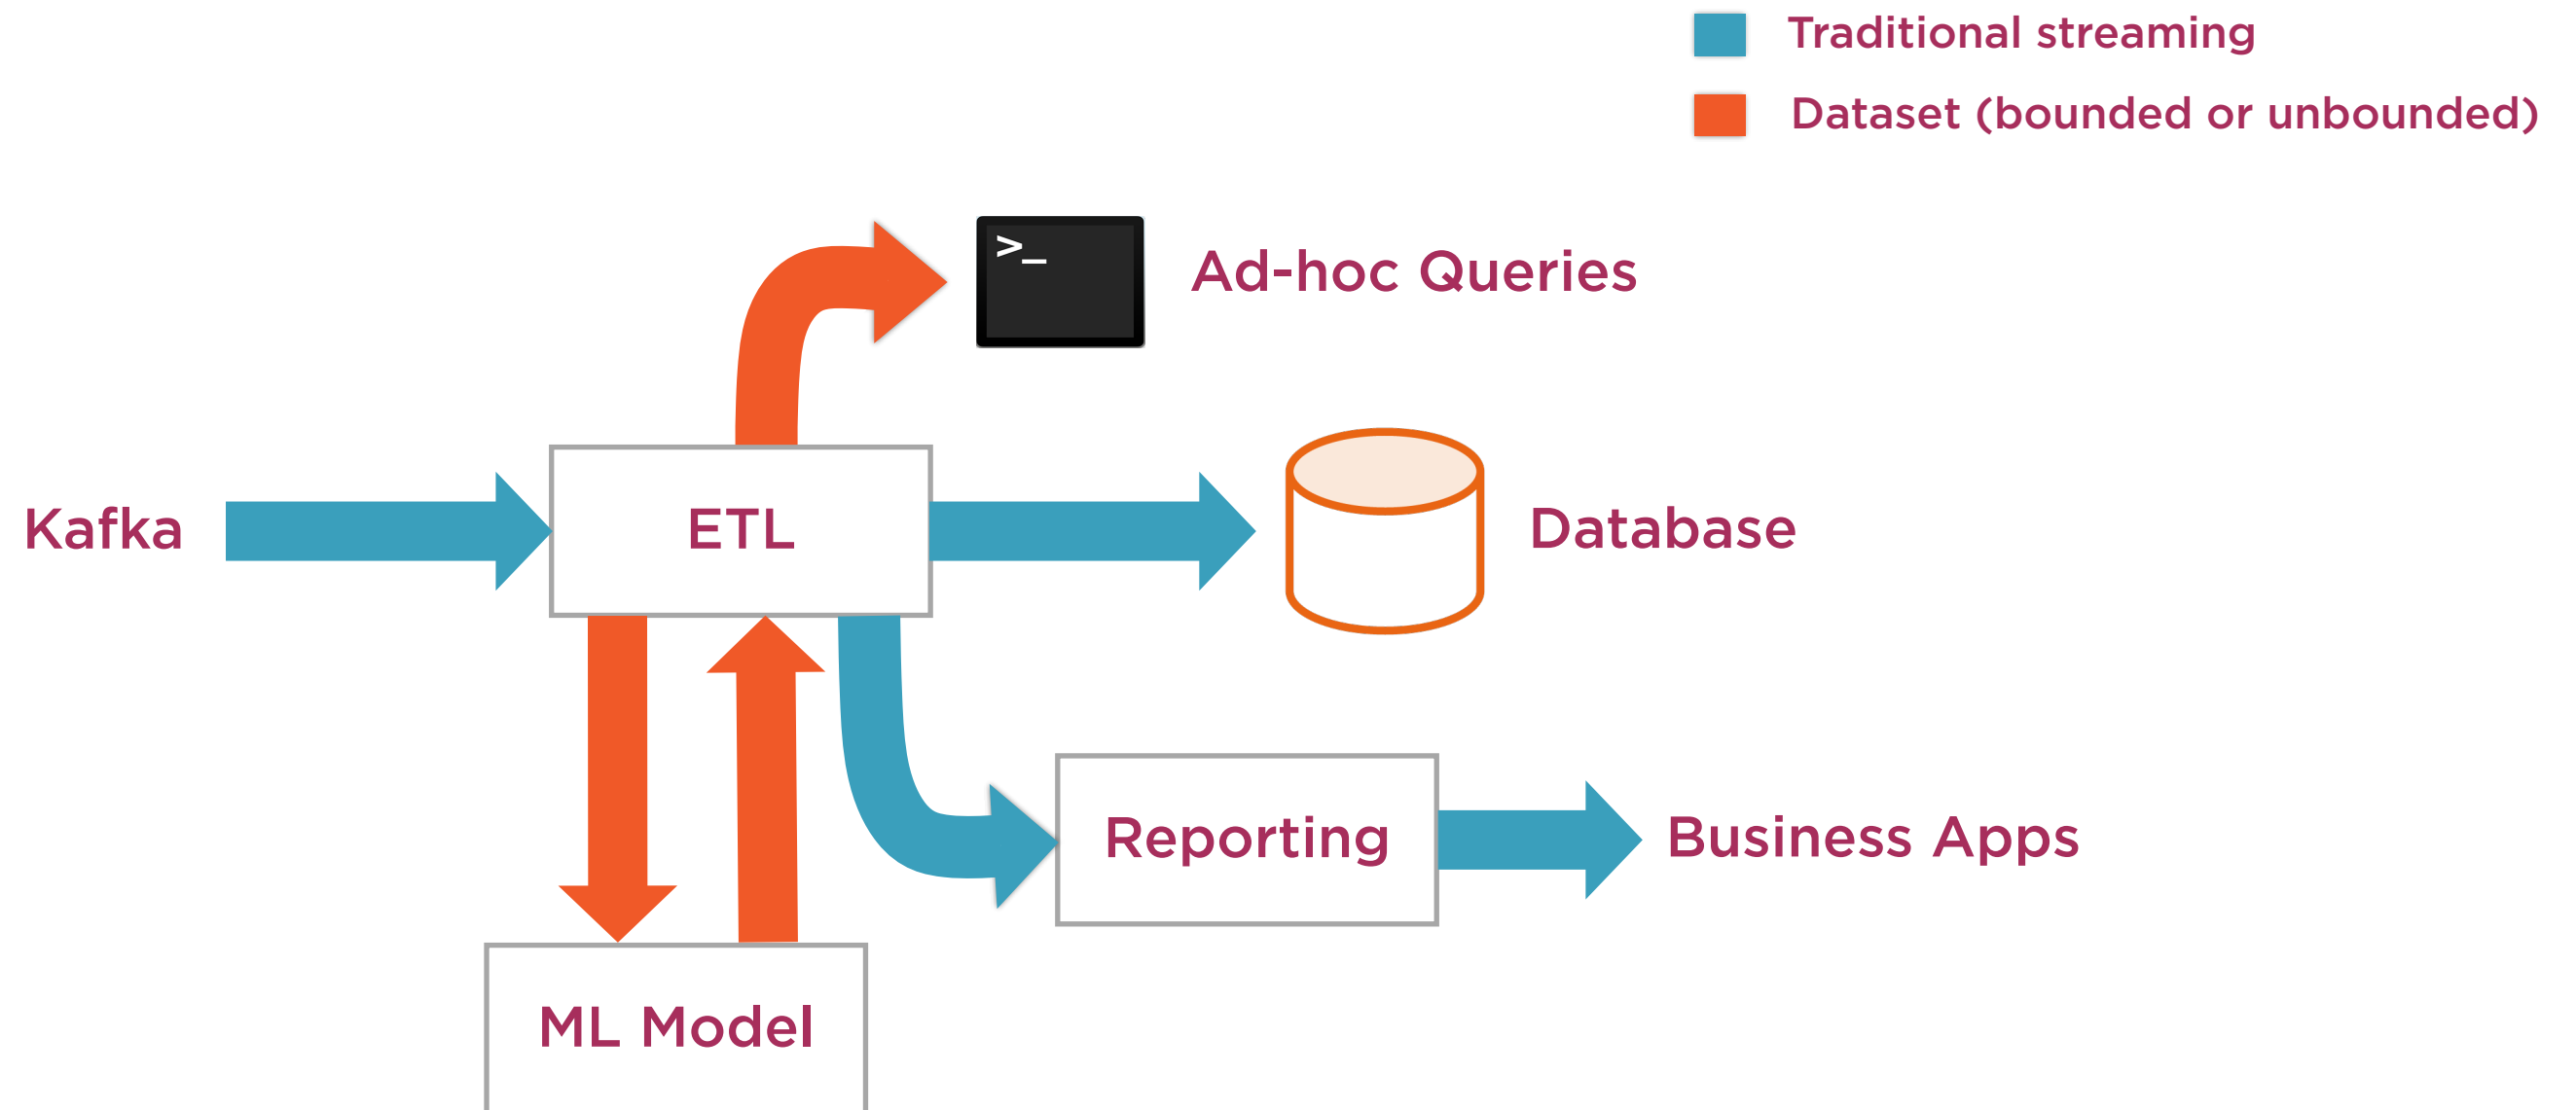
DataFrames

Optimizations on DataFrames

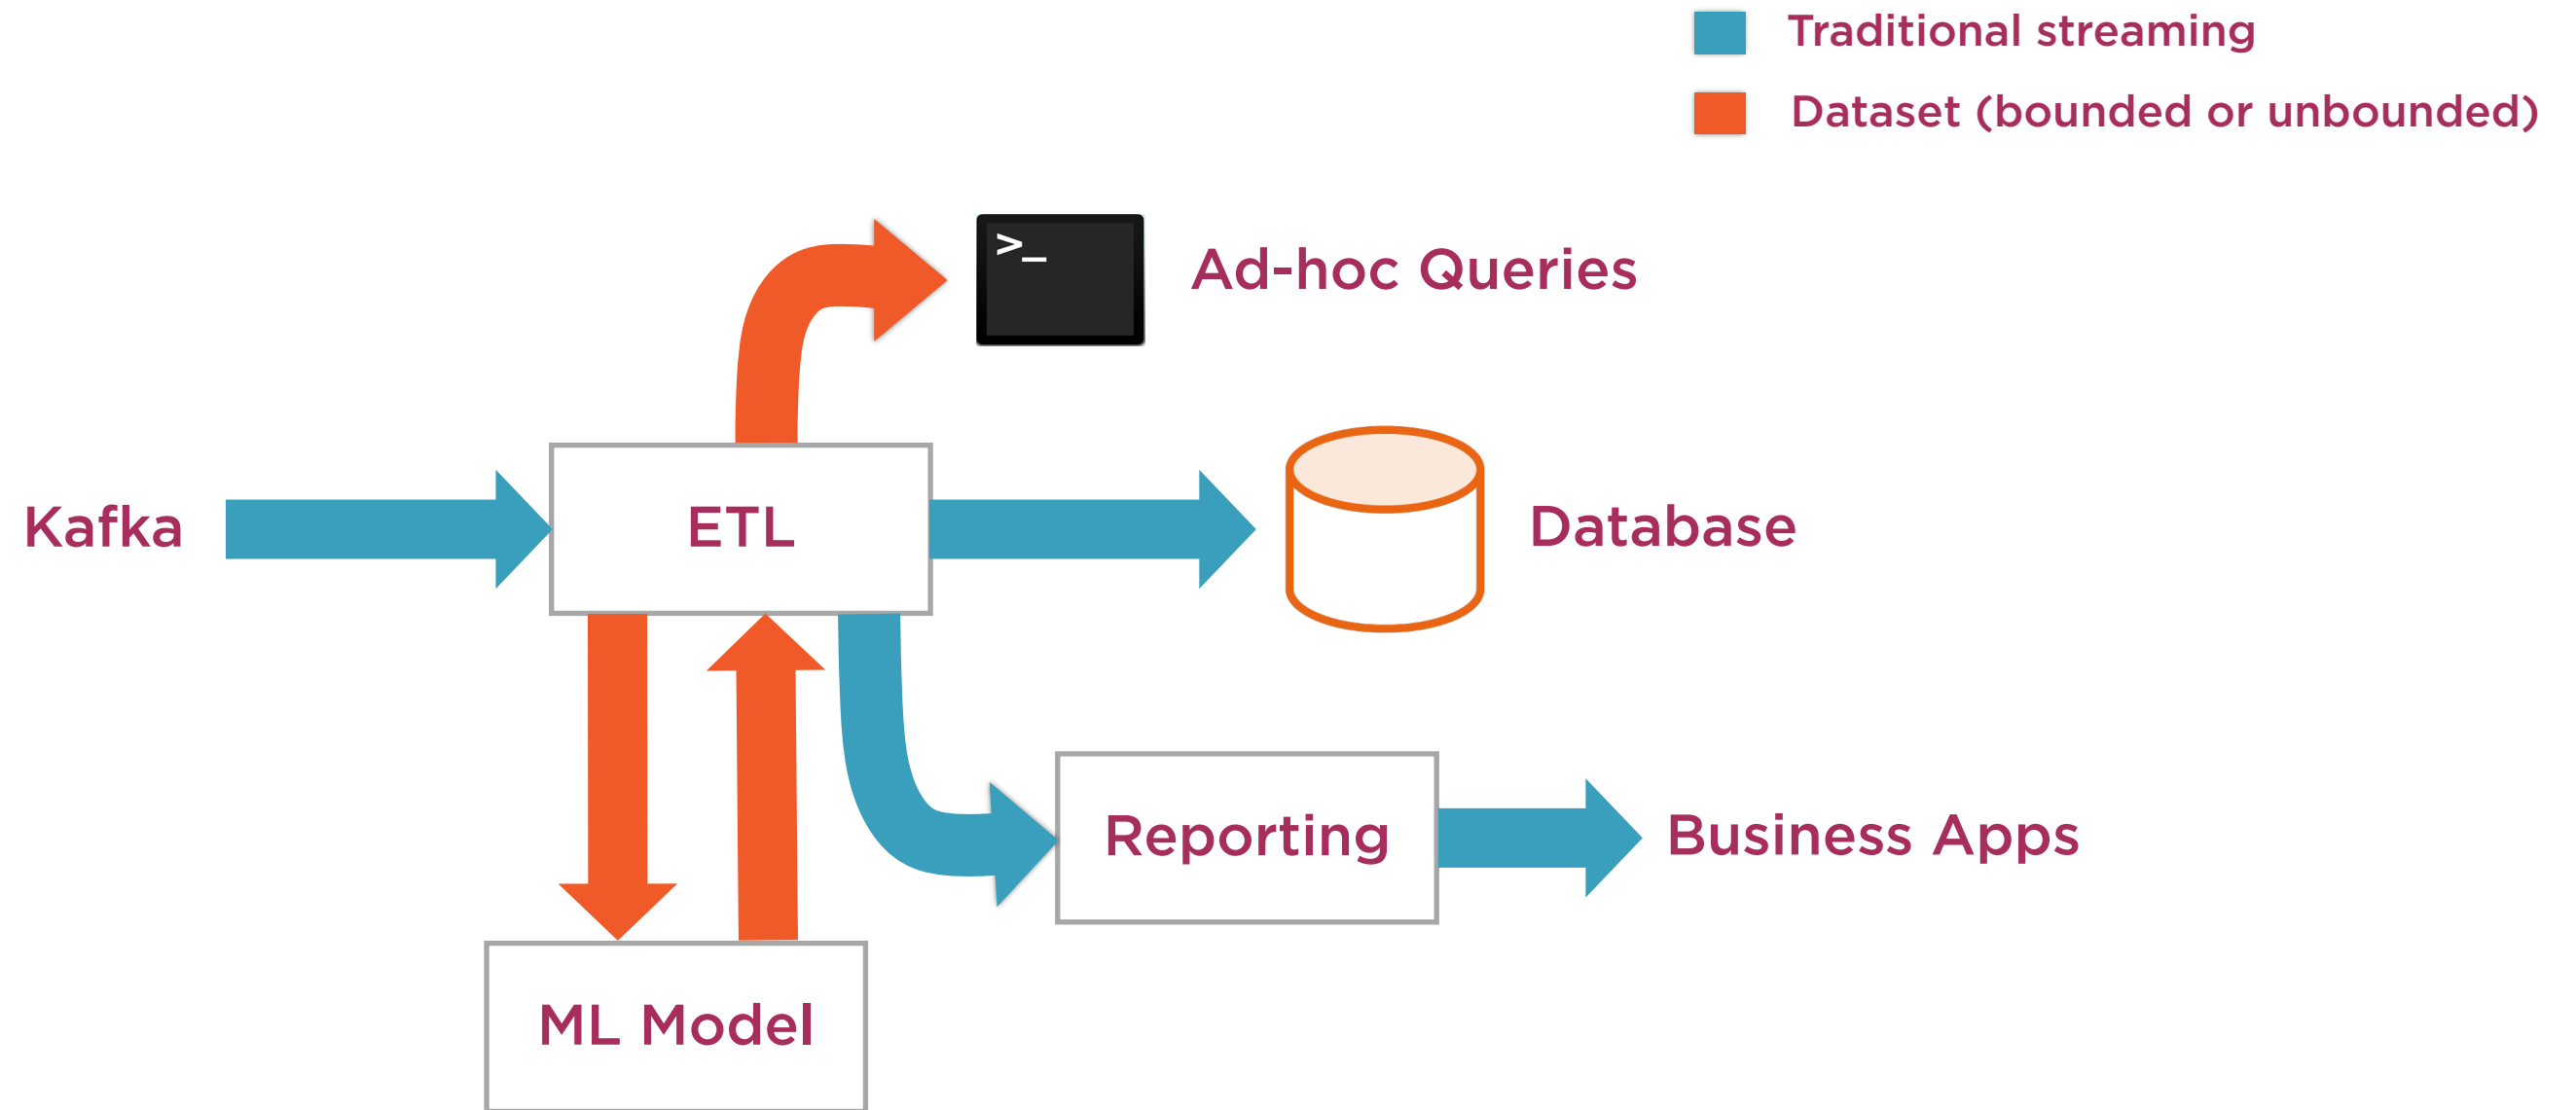
Unified support for batch and streaming

Structured Streaming

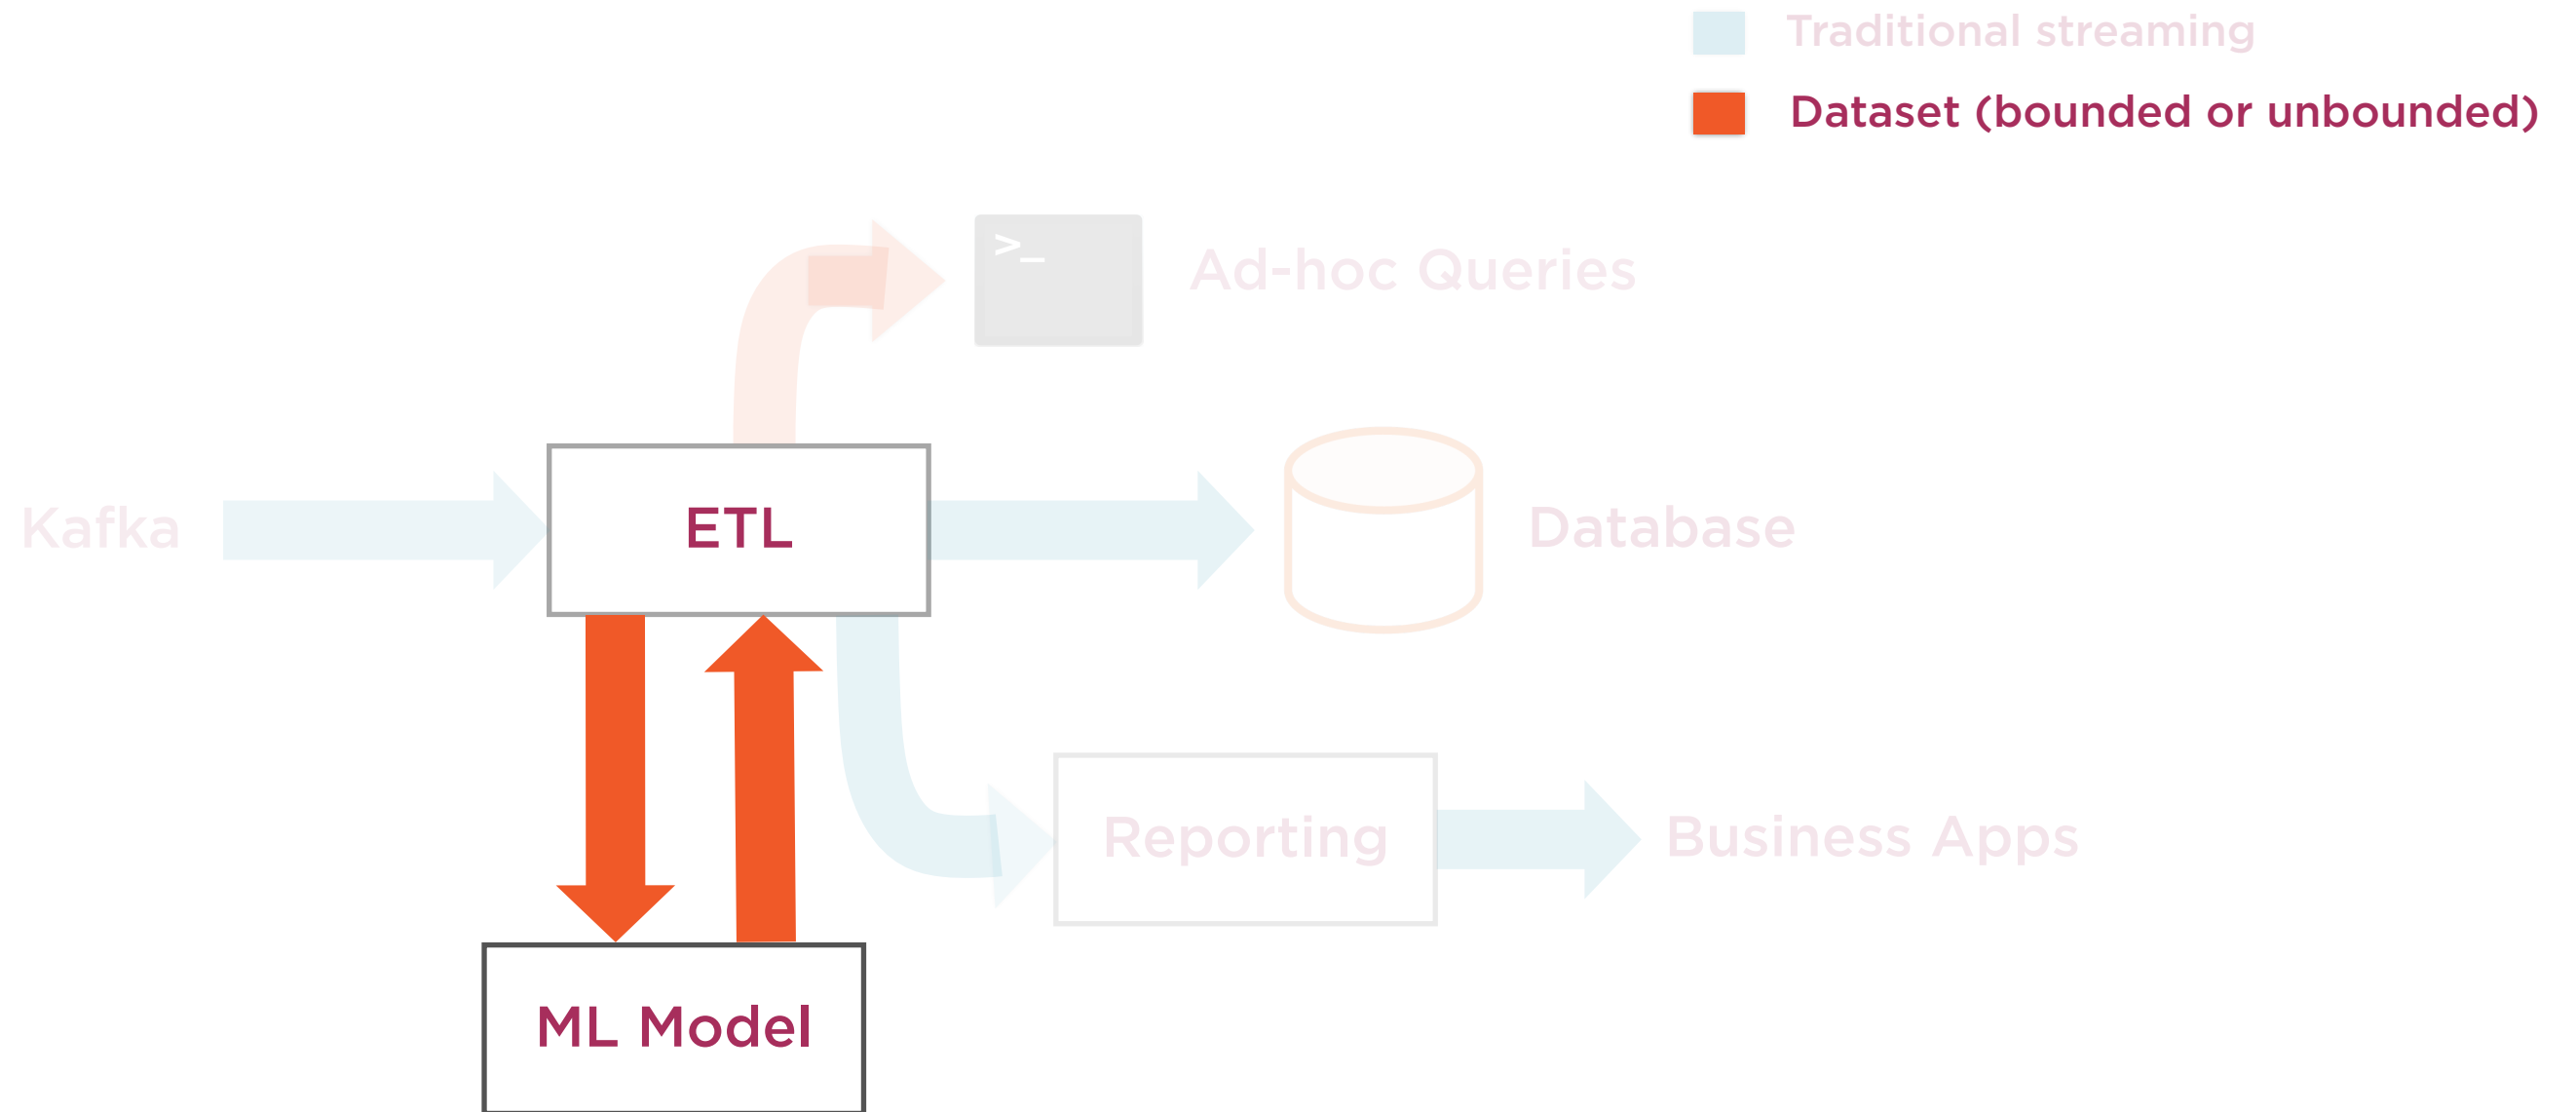
Spark 2.x: Continuous Applications



Structured Streaming

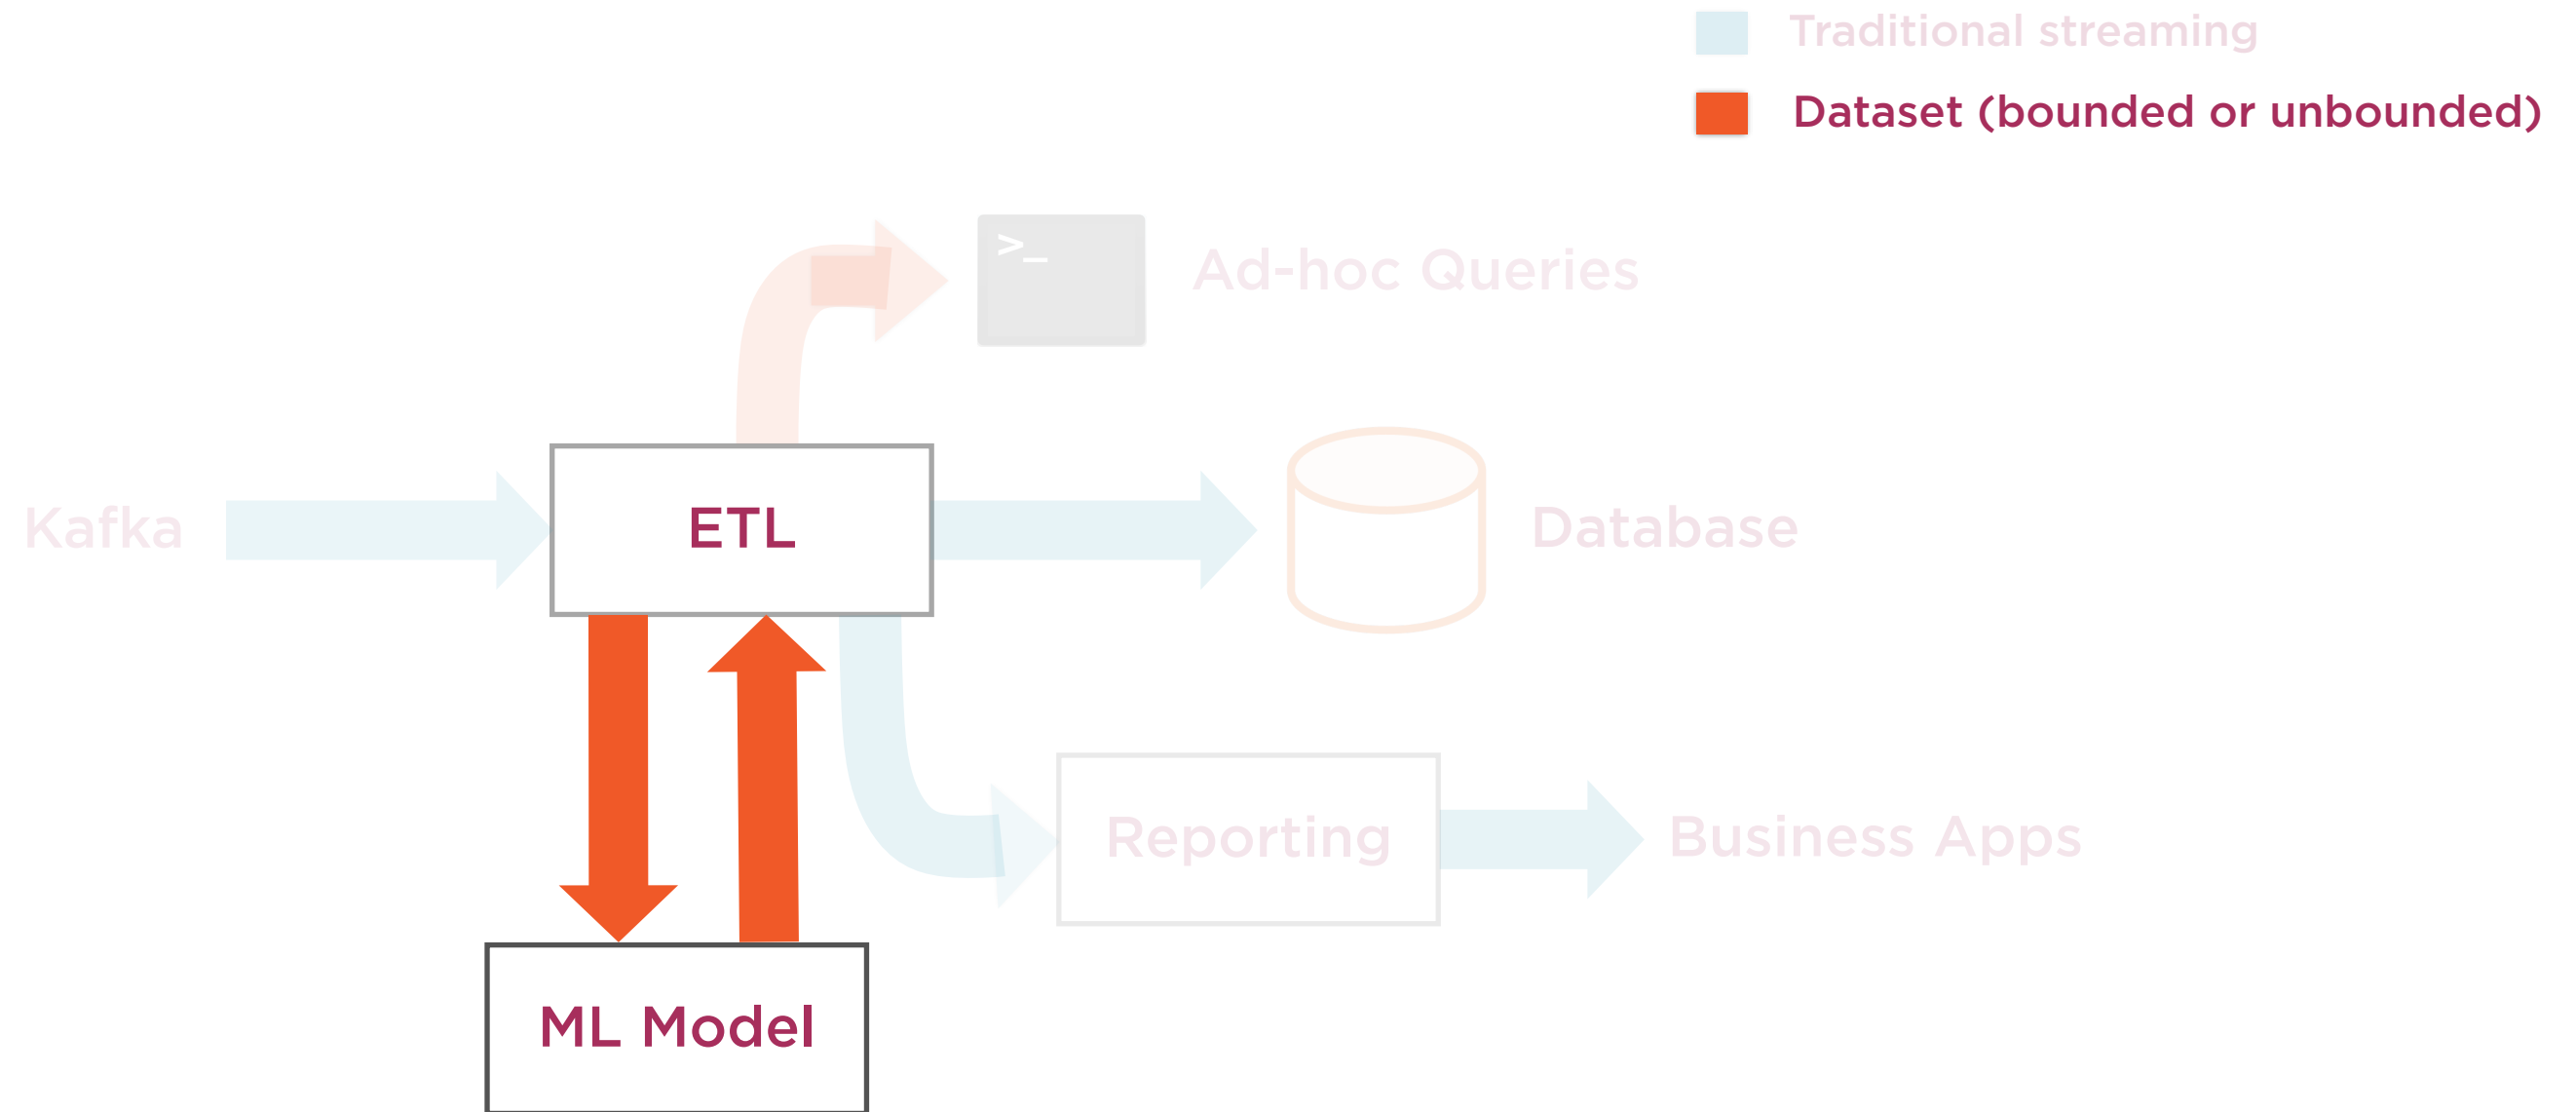


High-level User API



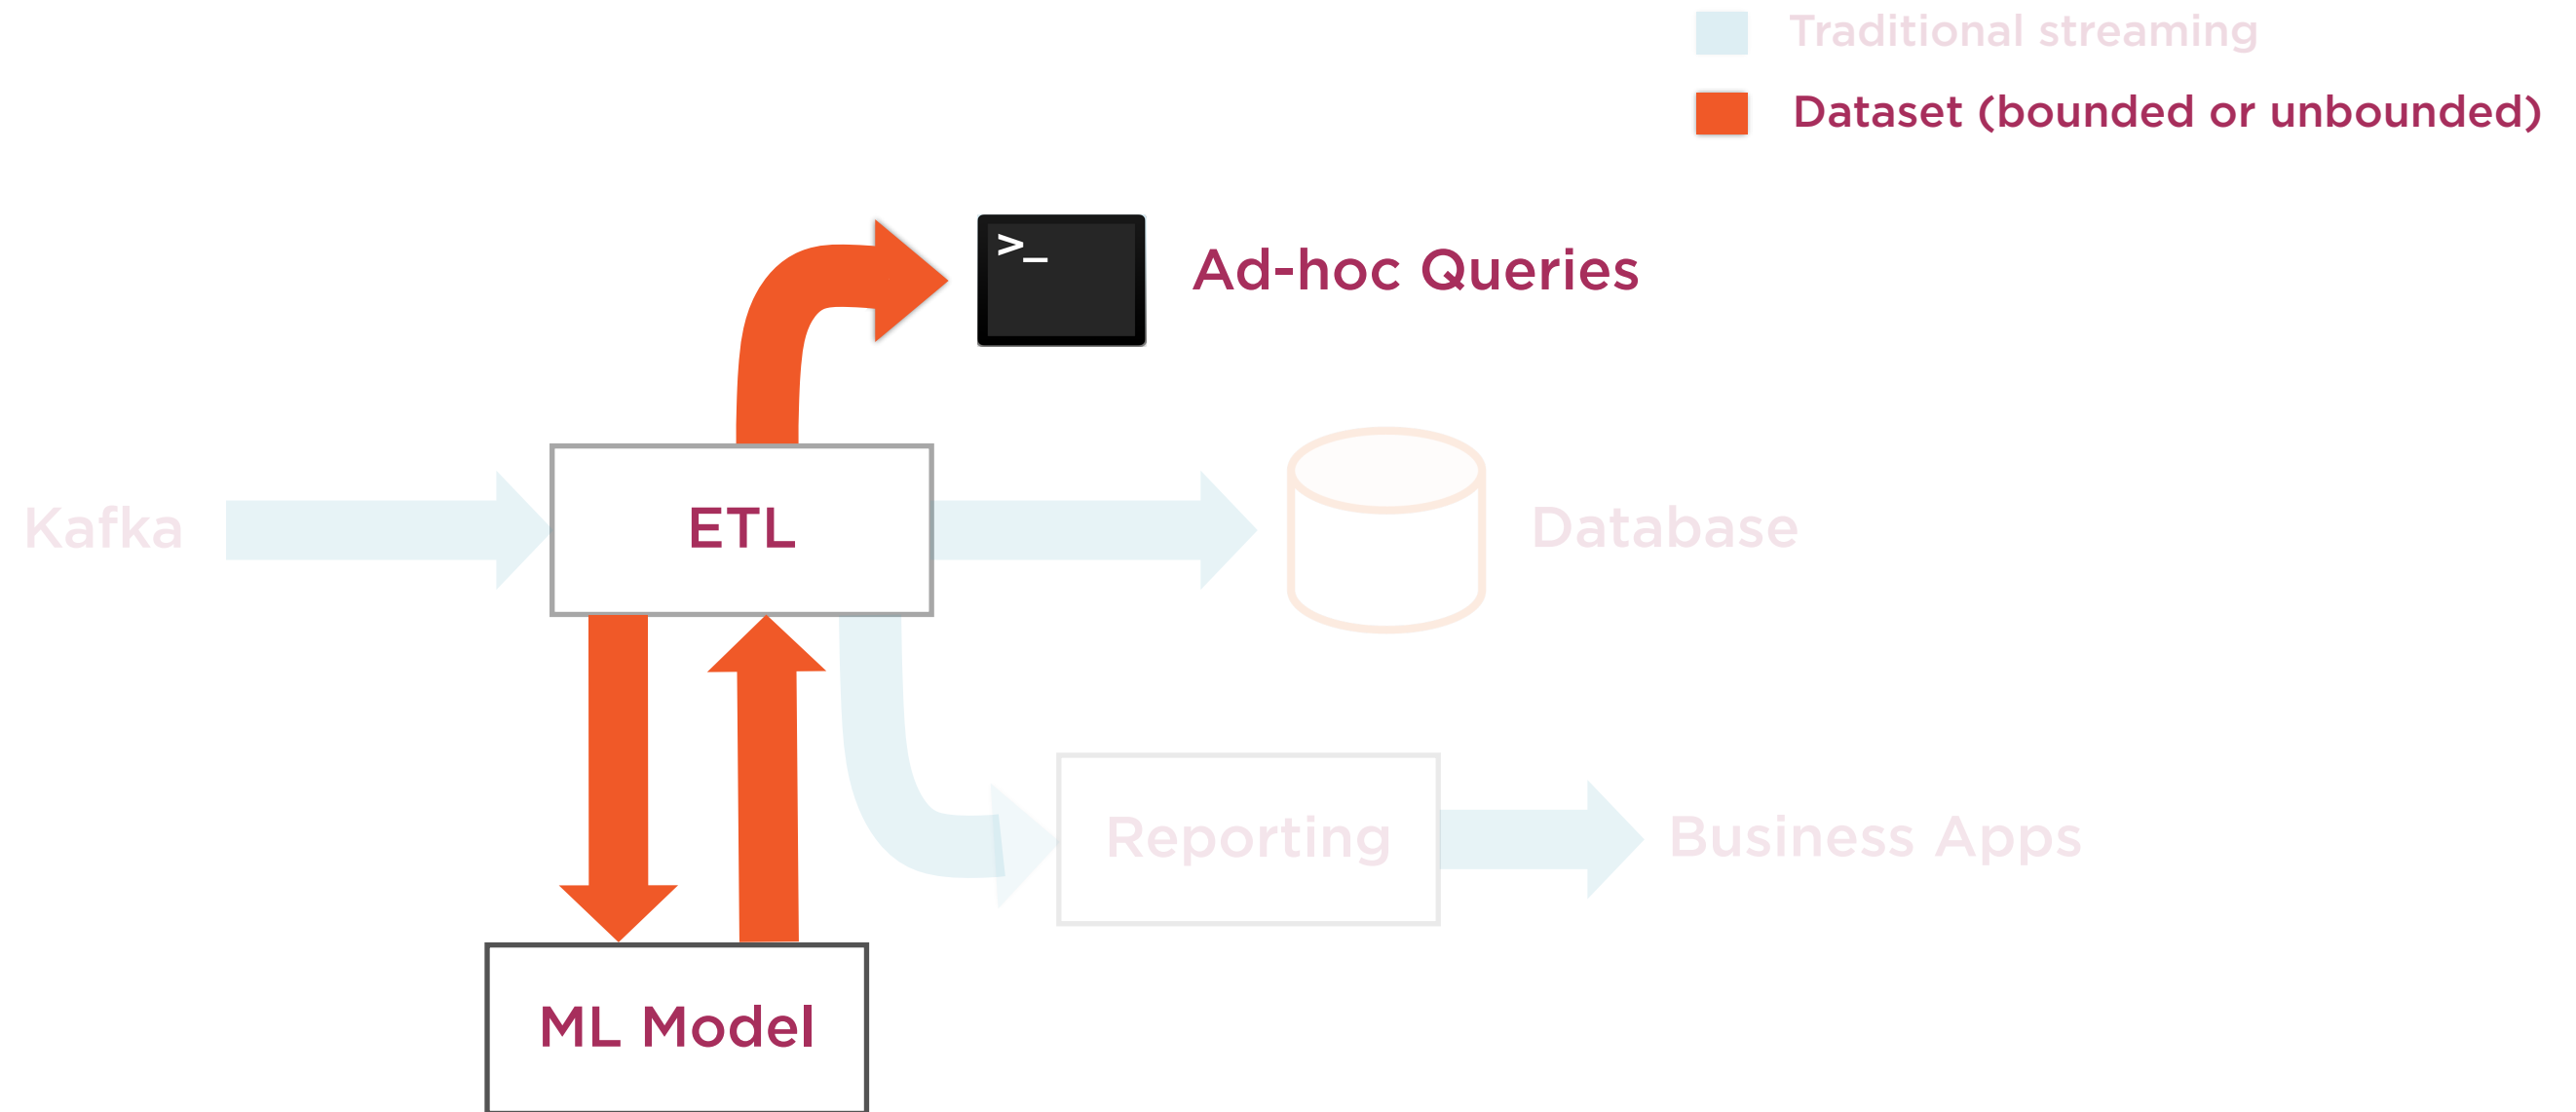
User implements **batch** computation using
DataFrame/Dataset API

Automatic Support for Continuous Apps



Spark automatically **incrementalizes** the batch computation

Automatic Support for Continuous Apps



**i.e. Spark automatically converts the job
from batch to streaming**

Micro-batch Processing



Default stream processing mode in Spark

Data streams processed as a series of batch jobs

End-to-end latencies as low as 100ms

Exactly-once fault-tolerance guarantees

Micro-batch Processing

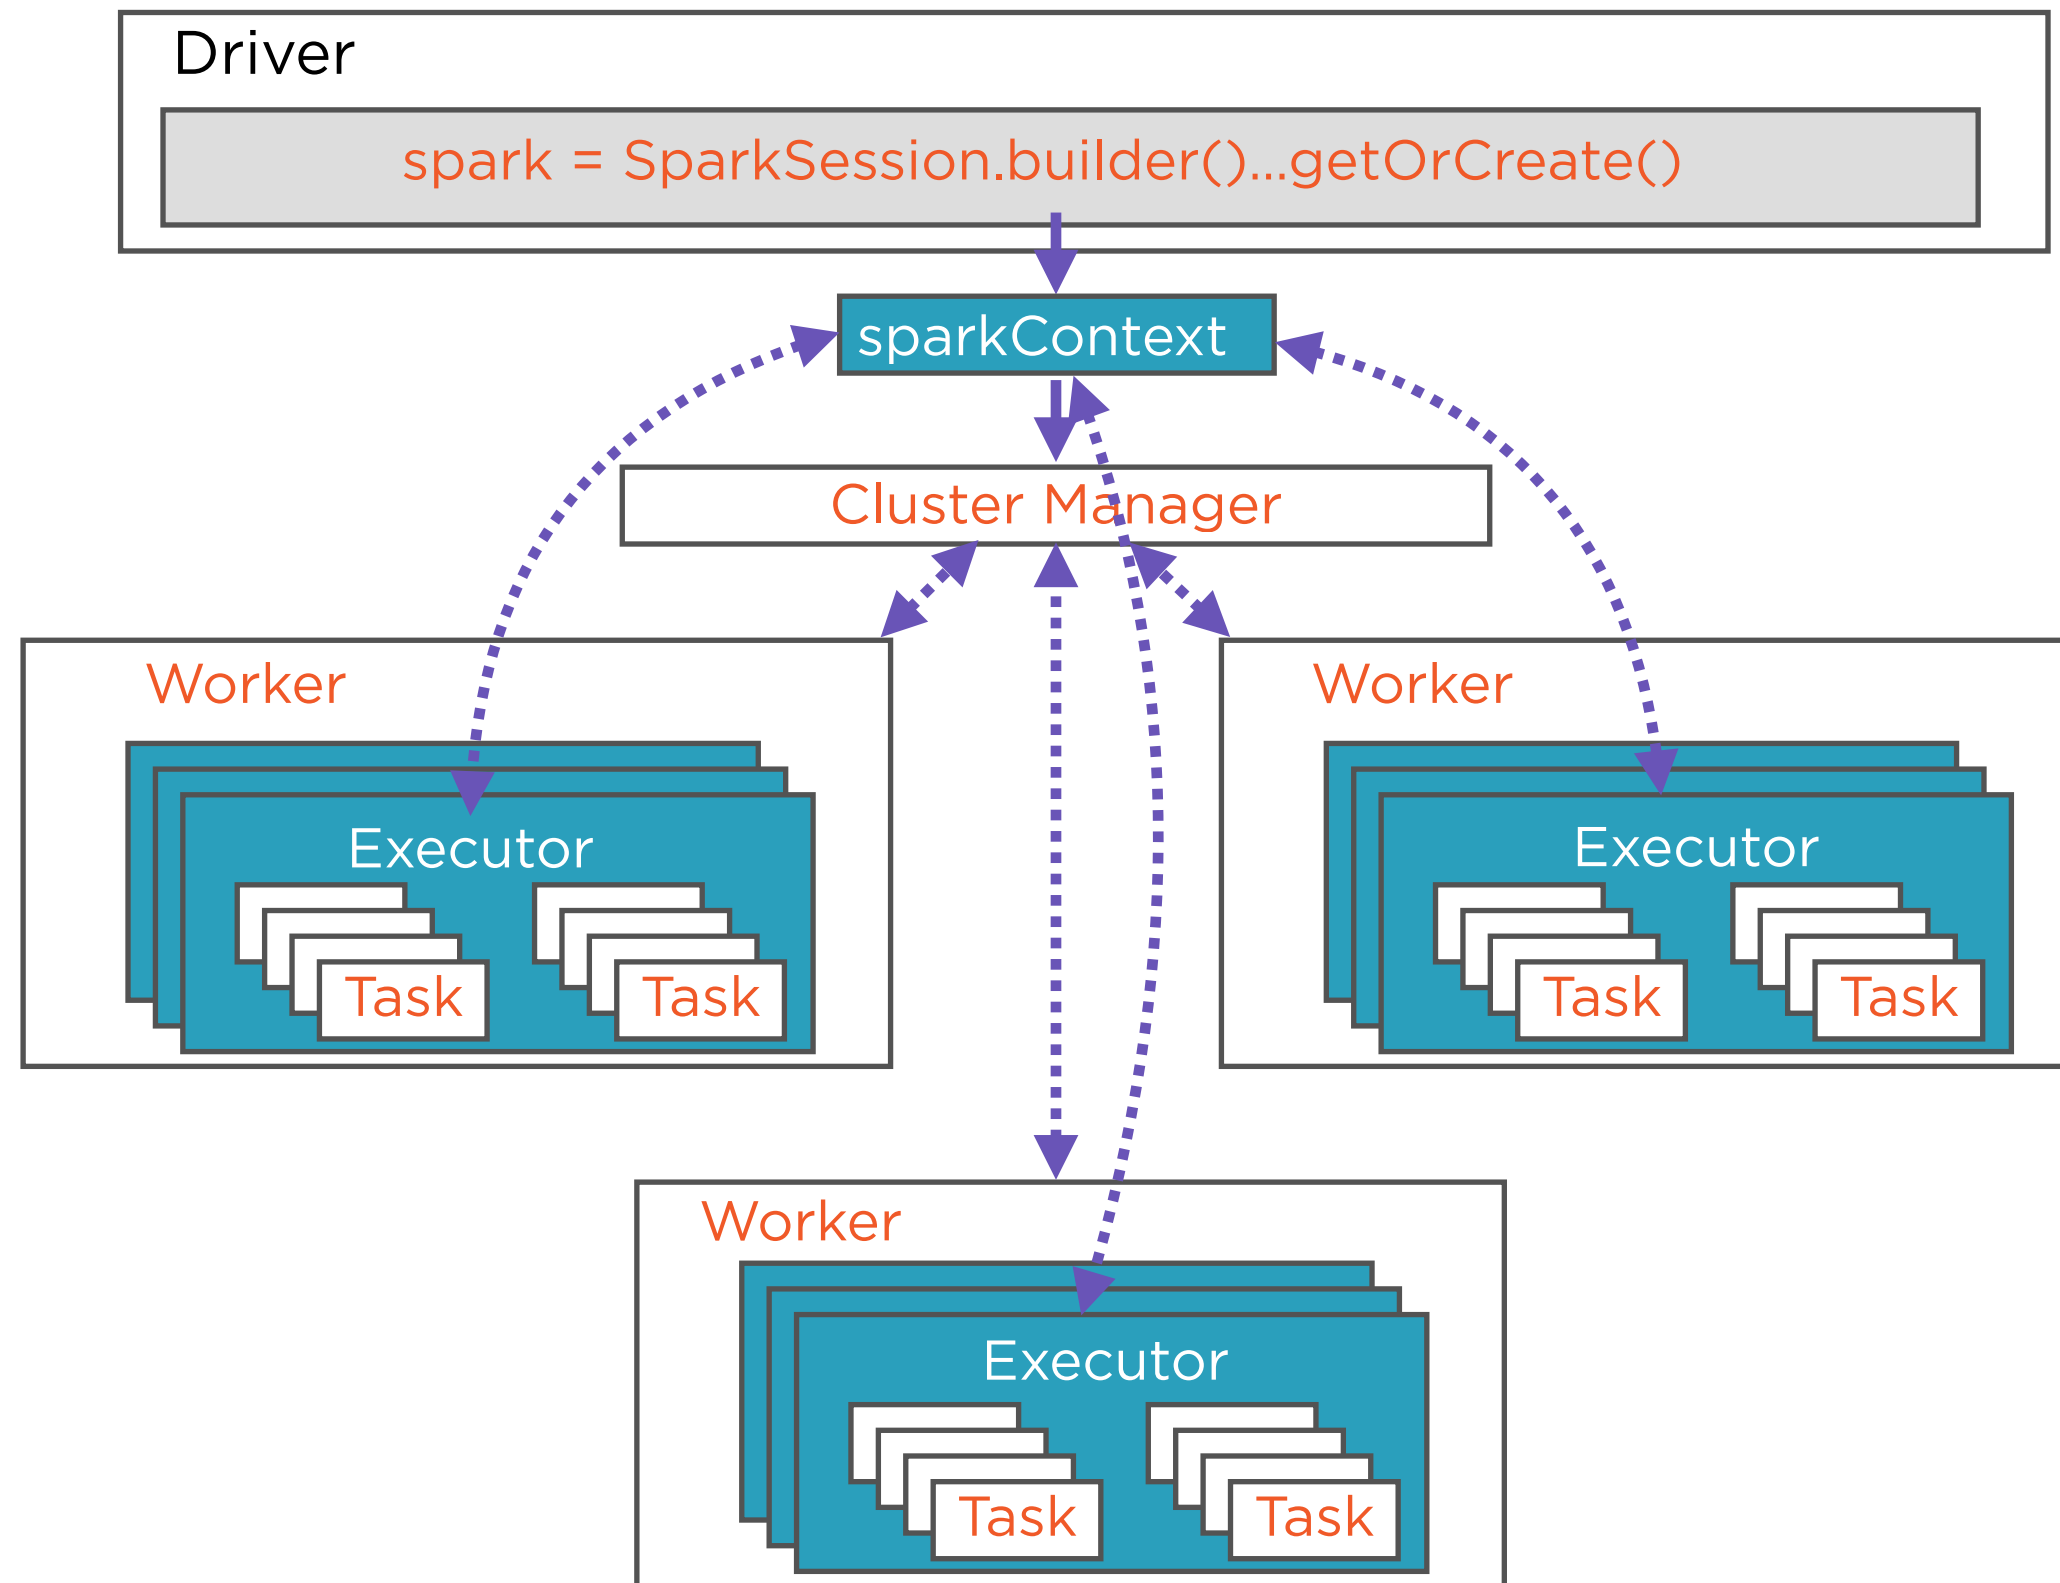


**Since Spark 2.3 there exists a new
continuous processing mode**

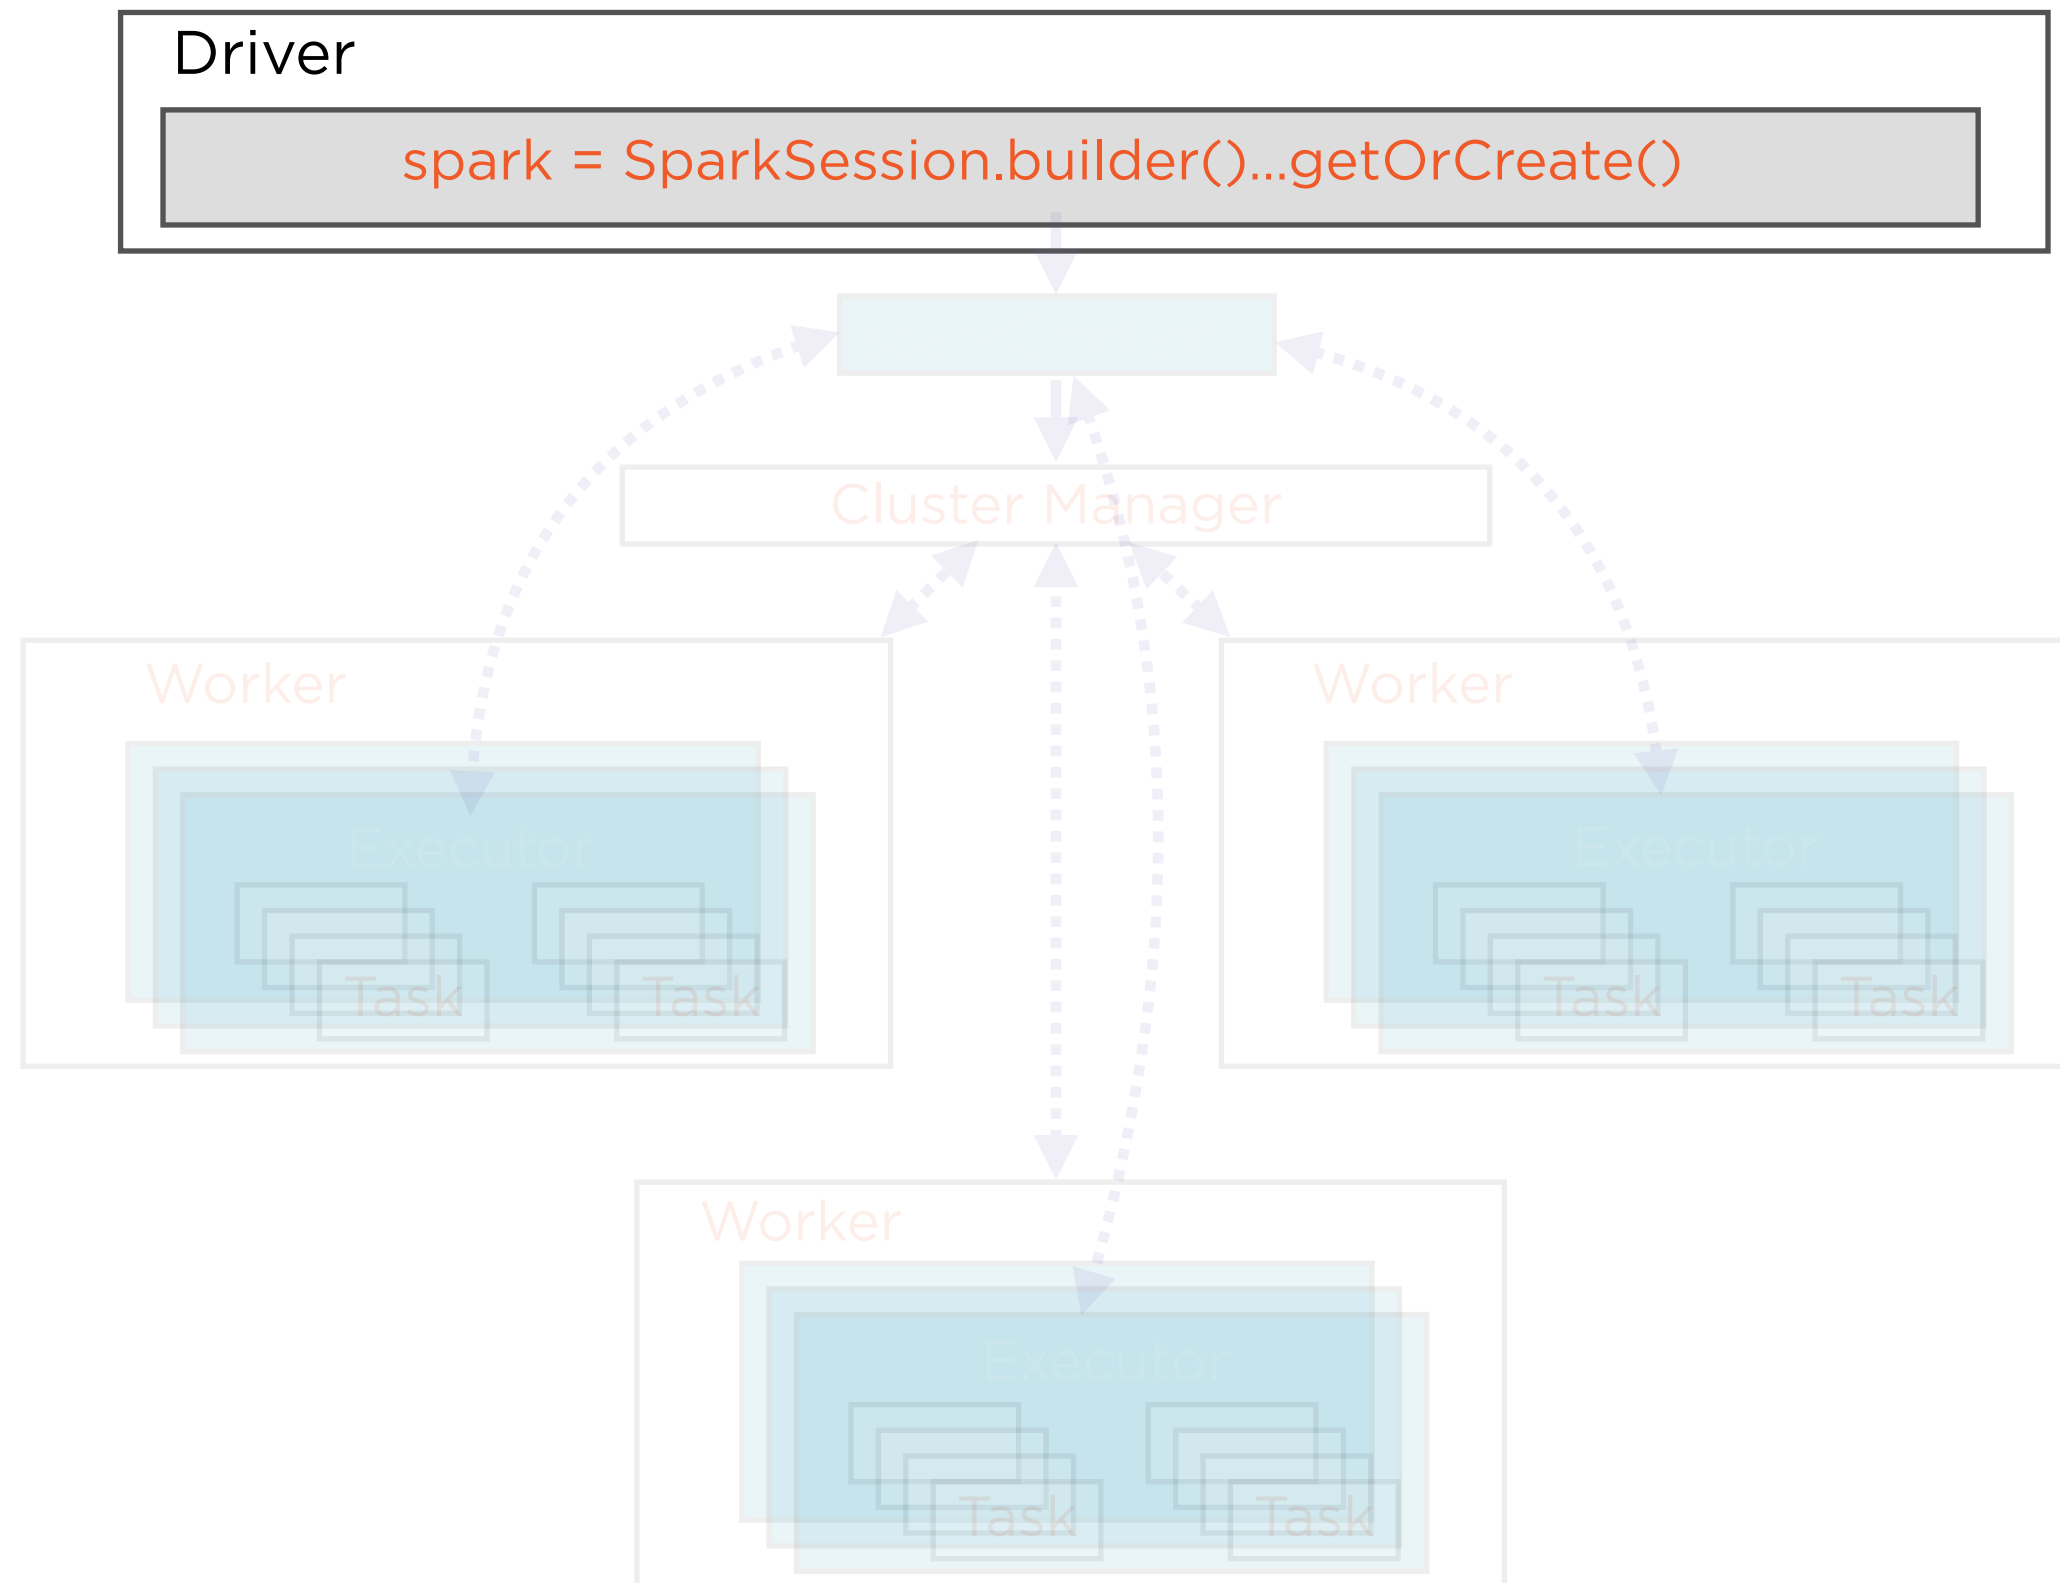
Currently in the experimental stage

Spark Architecture

Spark 2.x Architecture



Spark 2.x Architecture



Driver



Separate process (JVM)

The master node in a Spark application

Launches tasks

Hosts SparkContext

Driver



Several groups of services run inside the driver

- SparkEnv
- DAGScheduler
- Task Scheduler
- SparkUI
- ...

Spark Application



Uses SparkContext as entry point

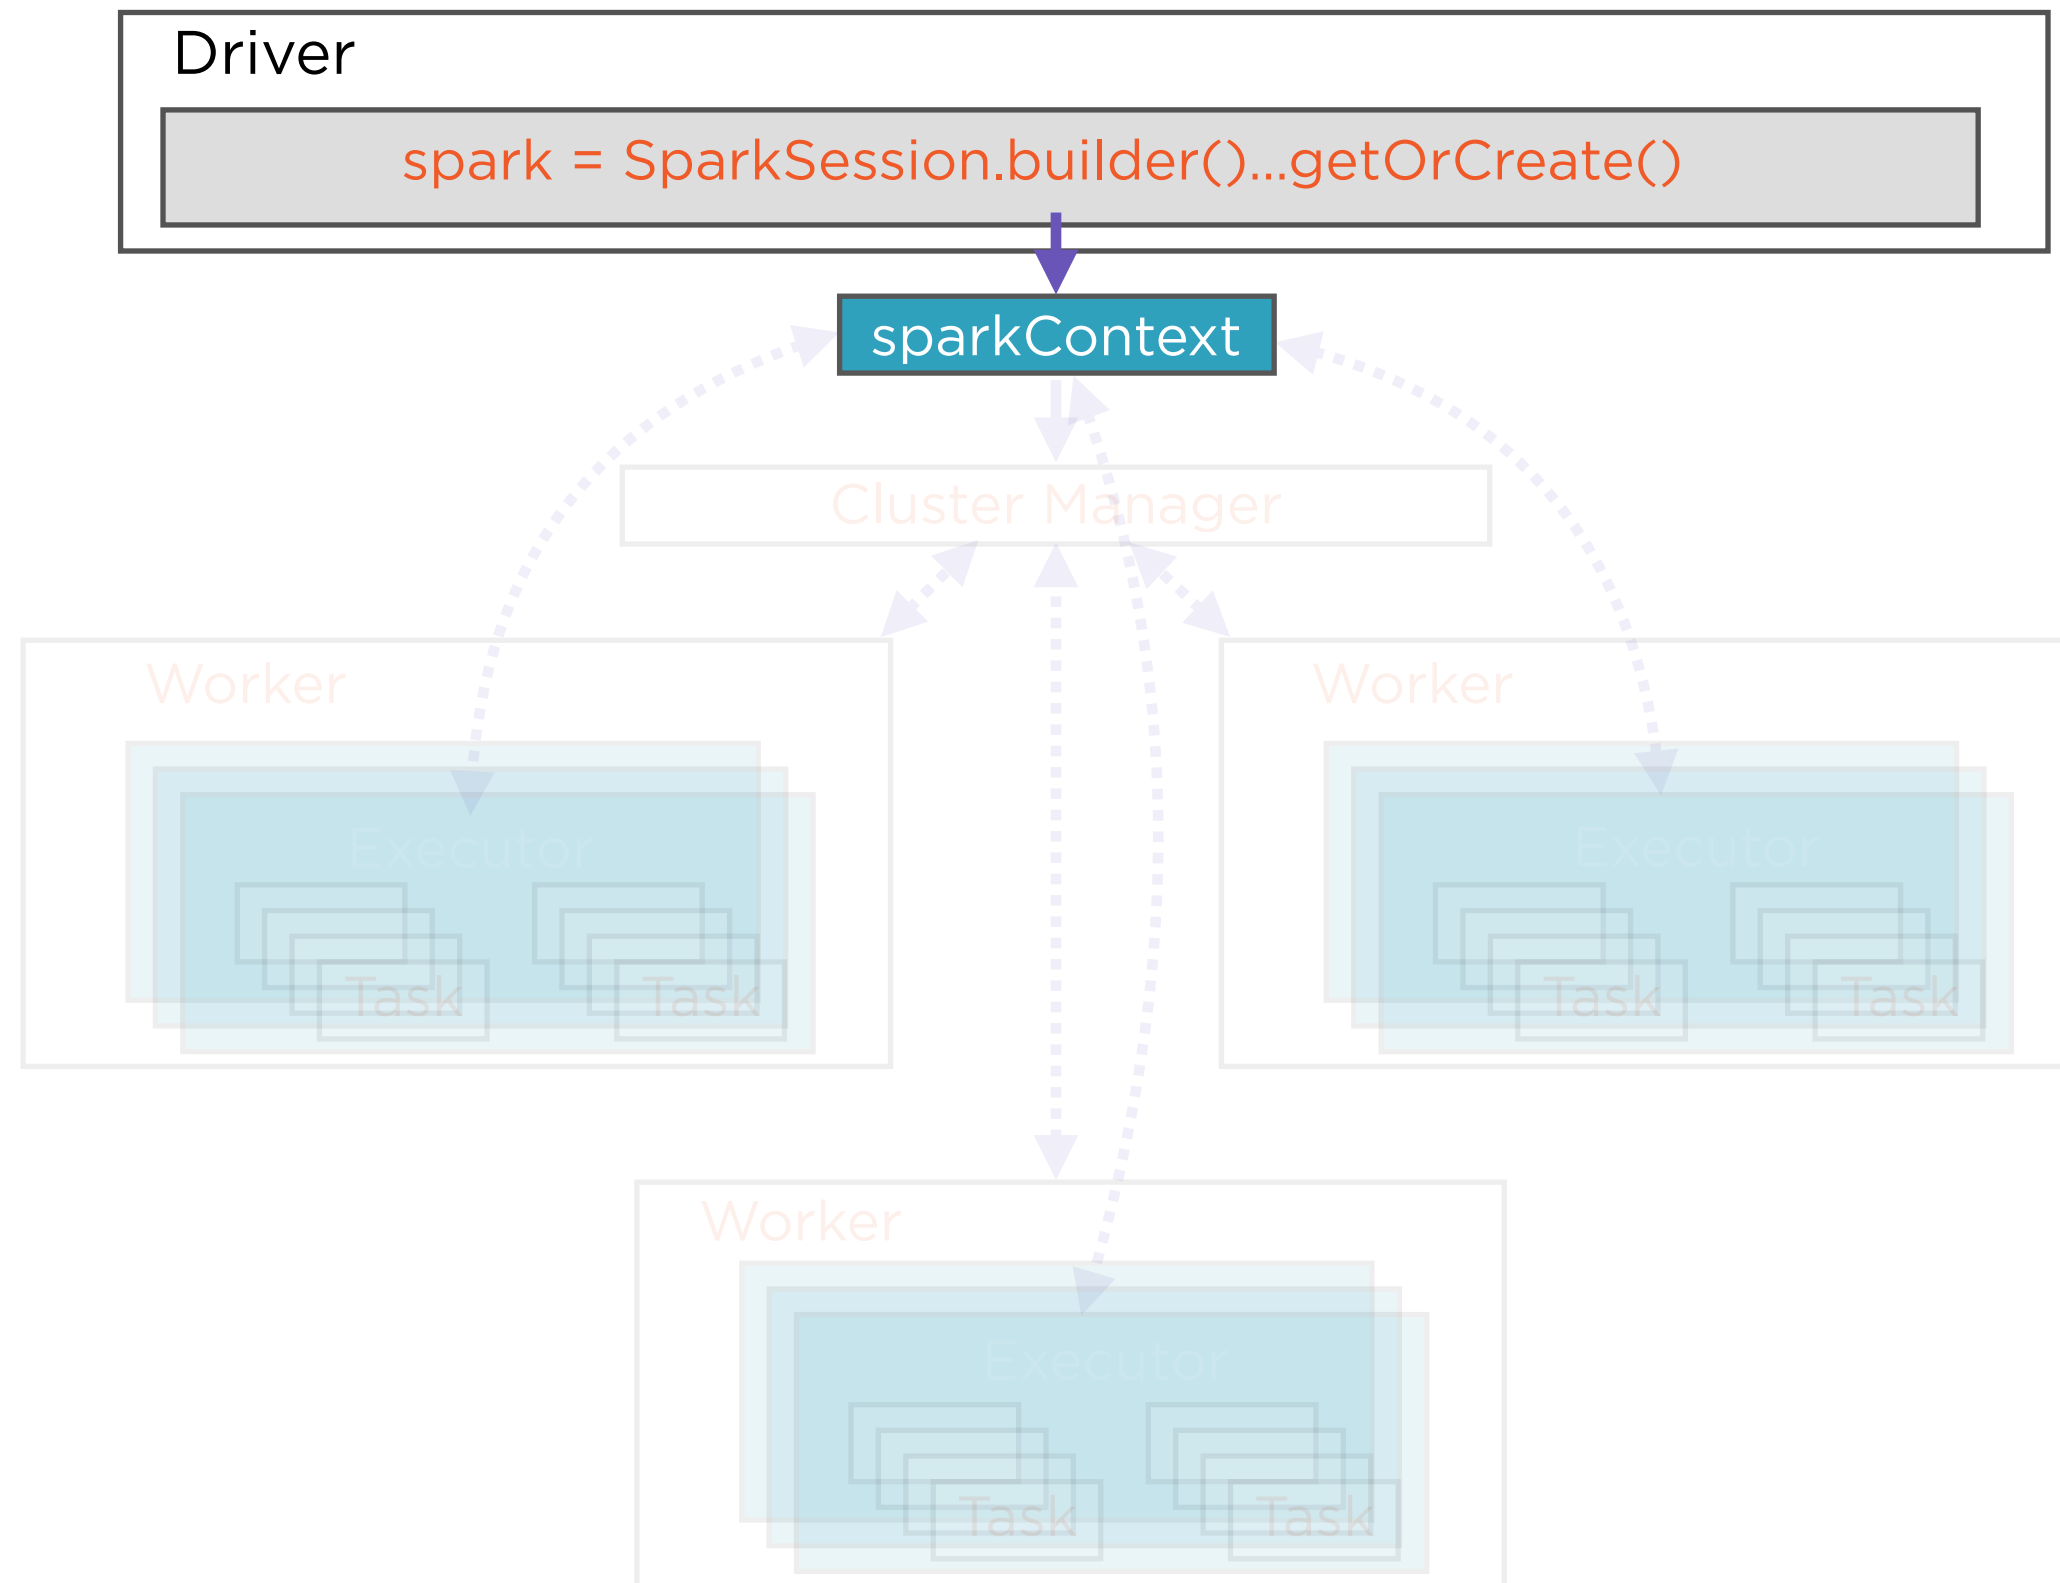
Creates RDD Directed Acyclic Graph

DataFrames run on top of RDDs

Internally, Spark creates *Stages*
(physical execution plan)

Each stage is split into operations on
RDD partitions called *Tasks*

Spark 2.x Architecture



SparkContext (Spark 1.x)



Familiar code entry point to Spark

```
sc = new sparkContext(...)
```

Create RDDs, accumulators...

Run jobs

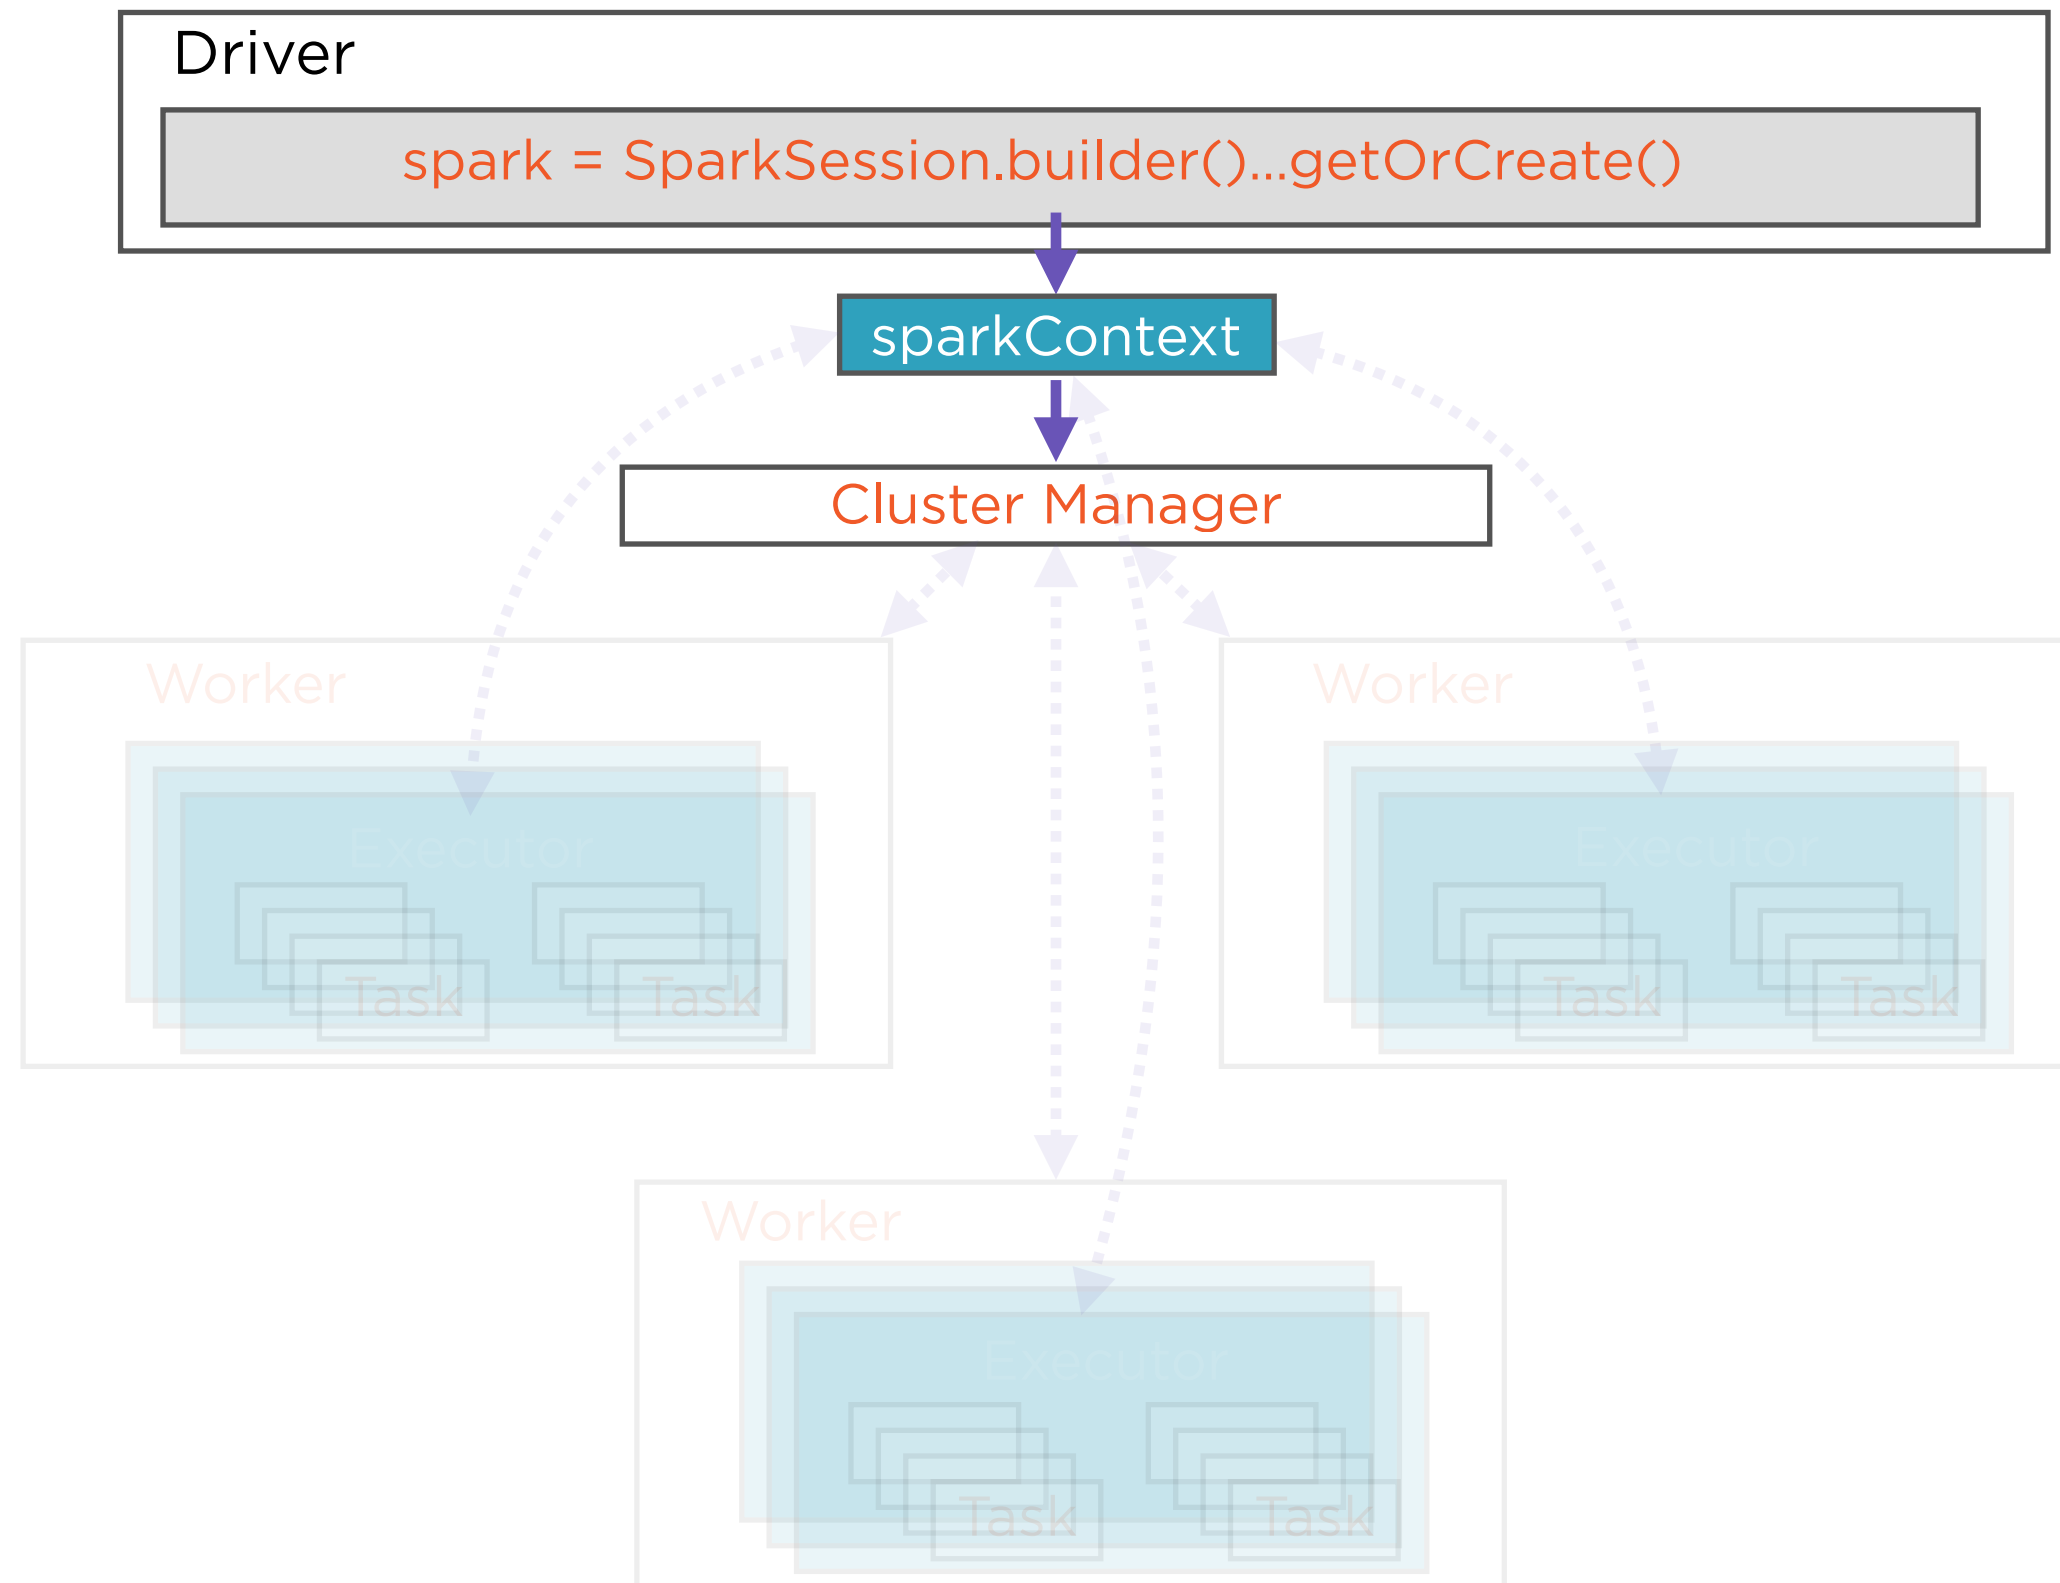
SparkSession (Spark 2.x)



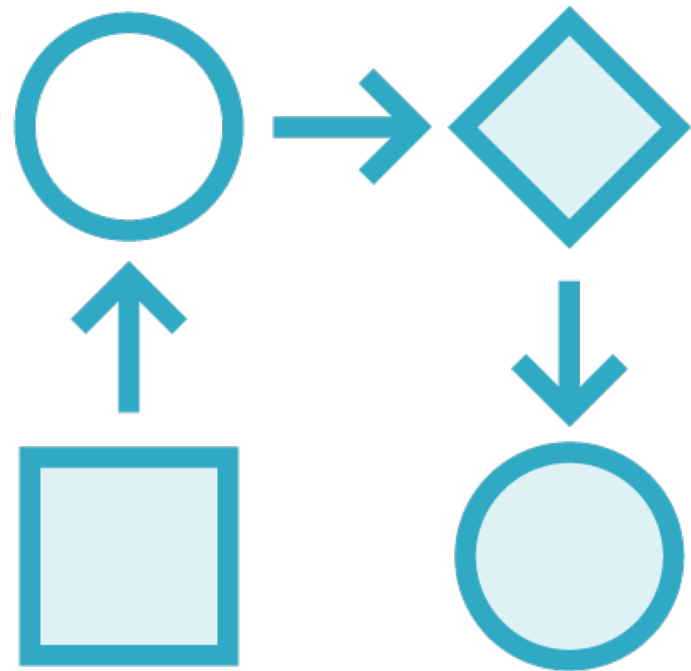
**In Spark 2.x, SparkContext is wrapped
in SparkSession**

**Encapsulates SparkContext,
SQLContext, HiveContext...**

Spark 2.x Architecture



Cluster Manager



Hadoop's YARN

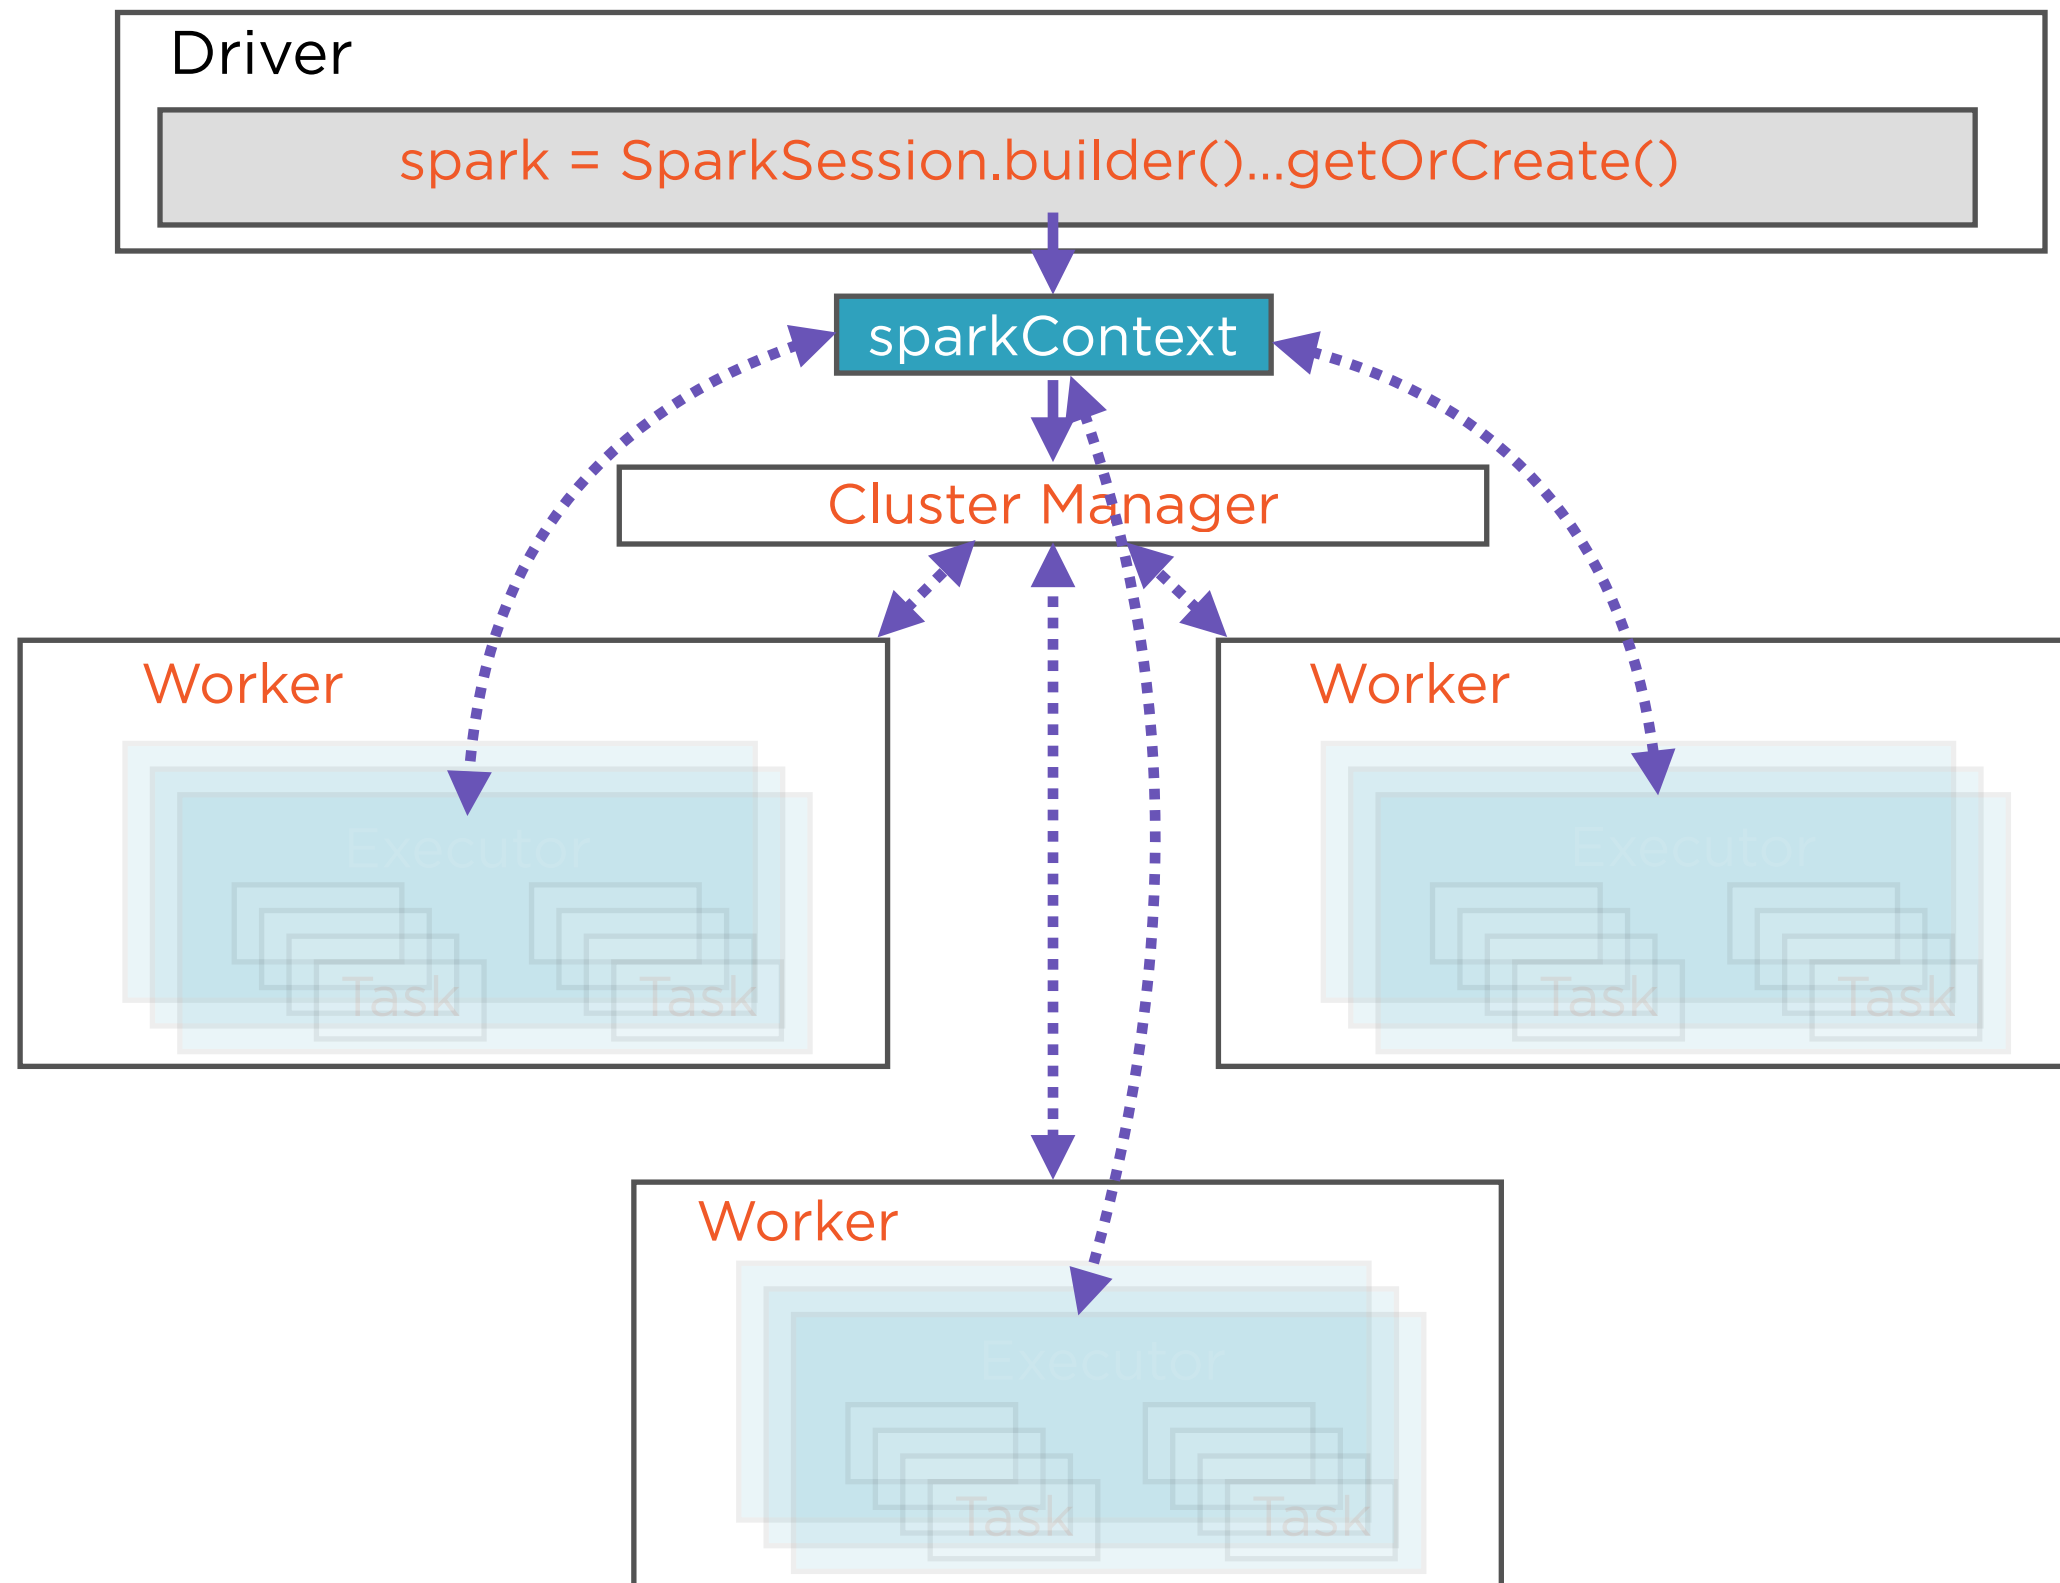
Apache Mesos

Kubernetes

Spark Standalone

Orchestrates execution

Spark 2.x Architecture



Worker



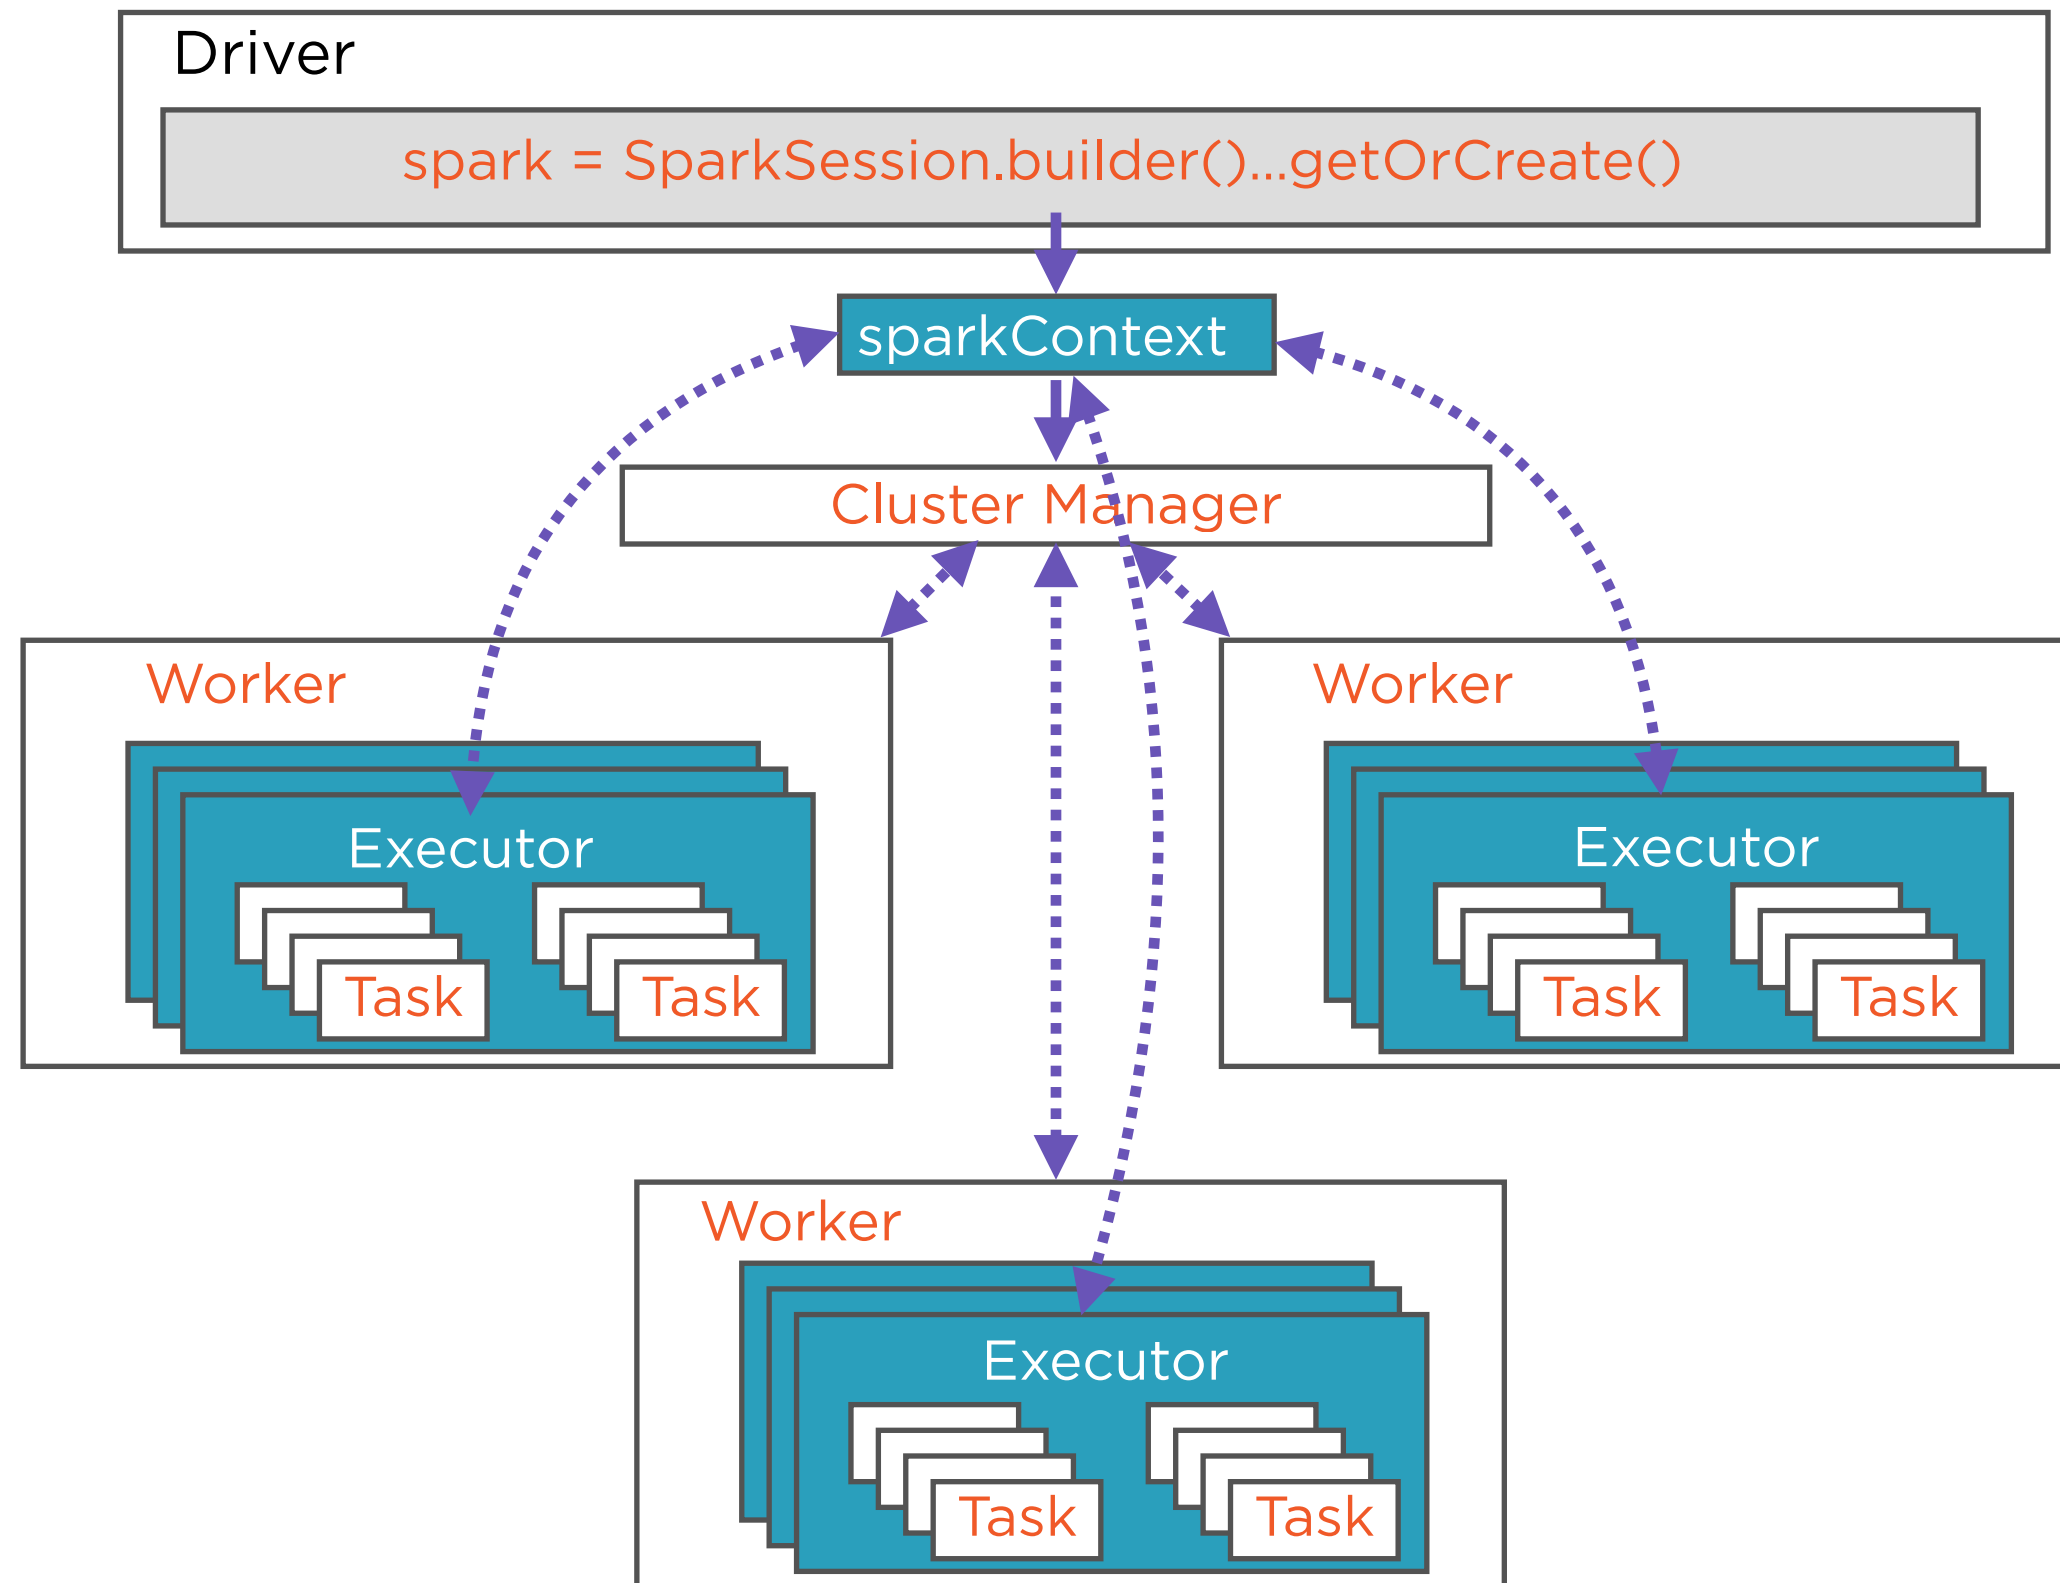
Compute nodes in cluster

Runs the Spark application code

When SparkContext created...

...Each worker starts *executors*

Spark 2.x Architecture



Executors



Distributed agents that execute *tasks*

Tasks are basic units of execution

Tasks belong inside *stages*

Stages are physical units of execution

Demo

**Install and set up Apache Spark on
MacOS**

Demo

**Install and set up Apache Spark on
Windows**

Demo

Streaming word count with a socket source and a console sink

Demo

Aggregations on streaming data with a socket source and console sink

Summary

**Streaming and its place in the
Apache Spark stack**

RDDs and DataFrames

Structured Streaming in Spark 2.x

Batch as a prefix of stream

**ClusterManager, SparkSession, and
Executors**

Up Next:

Executing Streaming Queries
