

Executing Streaming Queries



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Prefix integrity and implications

**Output modes - Append, Complete,
and Update**

Executing streaming queries in Spark

Schema and schema auto-detection

Prefix Integrity

Streaming Data Spark 2.x


Data stream



Every data item that is arriving on the stream is like a new row being appended to the input table

Streaming Data Spark 2.x

Data stream




Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream




Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



--	--	--

--	--	--

--	--	--

Every data item
that is arriving on
the stream is like a
new row being
appended to the
input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



--	--	--

--	--	--

--	--	--


--	--	--

Every data item
that is arriving on
the stream is like a
new row being
appended to the
input table

Data stream as an unbounded input table

Batch is Simply A Prefix of Stream






Data stream



**In other words, the
input table (batch)
is simply a prefix of
the stream**

Batch is Simply A Prefix of Stream

Data stream



**All operations that
can be performed on
data frames can be
performed on the
stream**

Prefix Integrity

Running job on continuous data yields same result as running job on batch data (where the batch is a prefix or snapshot of continuous data).

Prefix Integrity

Running job on continuous data yields **same result** as running job on batch data (where the batch is a prefix or snapshot of continuous data).

Prefix Integrity


Running job on continuous data yields same result as running job on batch data (where the batch is a prefix or snapshot of continuous data).

Structured Streaming treats a live data stream as a table that is being continuously appended

Burden of stream-processing shifts from user to system

Triggers and Output Modes

Stream Processing Model





Data Source

File source

Kafka source

Socket source

Rate source



Data Sink

File sink

Kafka sink

Foreach sink

Console sink

Memory sink



Data Sink

Structured streaming reads data from source

- Processes incrementally
- Updates result
- Discards source

Only minimal intermediate state maintained

Trigger

Events that determine when transformations on accumulated input data need to be re-performed. Each trigger event emits new data into the Result Table.

Trigger

Events that determine when transformations on accumulated input data need to be re-performed. Each trigger event emits new data into the Result Table.

Trigger

Events that determine when transformations on accumulated input data need to be re-performed. Each trigger event emits new data into the Result Table.

Trigger

Events that determine when transformations on accumulated input data need to be re-performed. Each trigger event emits new data into the Result Table.

Result Table

Executing a query on input data generates the Result Table. Rows in the Result Table are written out to an external data sink.

Types of Triggers

Default

Fixed interval micro-batch

One-time micro-batch

**Continuous with fixed
checkpoint interval**

Micro-batch Processing Mode

Default

Fixed interval micro-batch

One-time micro-batch

Continuous with fixed
checkpoint interval

Continuous Processing Mode


Default

Fixed interval micro-batch

One-time micro-batch

**Continuous with fixed
checkpoint interval**

Default




Used when no trigger setting specified

Query executed in micro-batch mode

Each new micro-batch generated when previous one completes processing


Fixed Interval Micro-batch



Micro-batch kicked off at user-specified intervals

If no data available no processing

Fixed Interval Micro-batch




If previous micro-batch completes **within** the interval:

- engine waits till interval is over

If previous micro-batch takes **longer** than specified interval:

- next micro-batch starts as soon as data arrives

One-time Micro-batch



Execute only one micro-batch to process all available data


Once processed query will stop

Used when cluster periodically spun up to process data since last period


May result in significant cost savings

Result Table


**Input
Table**



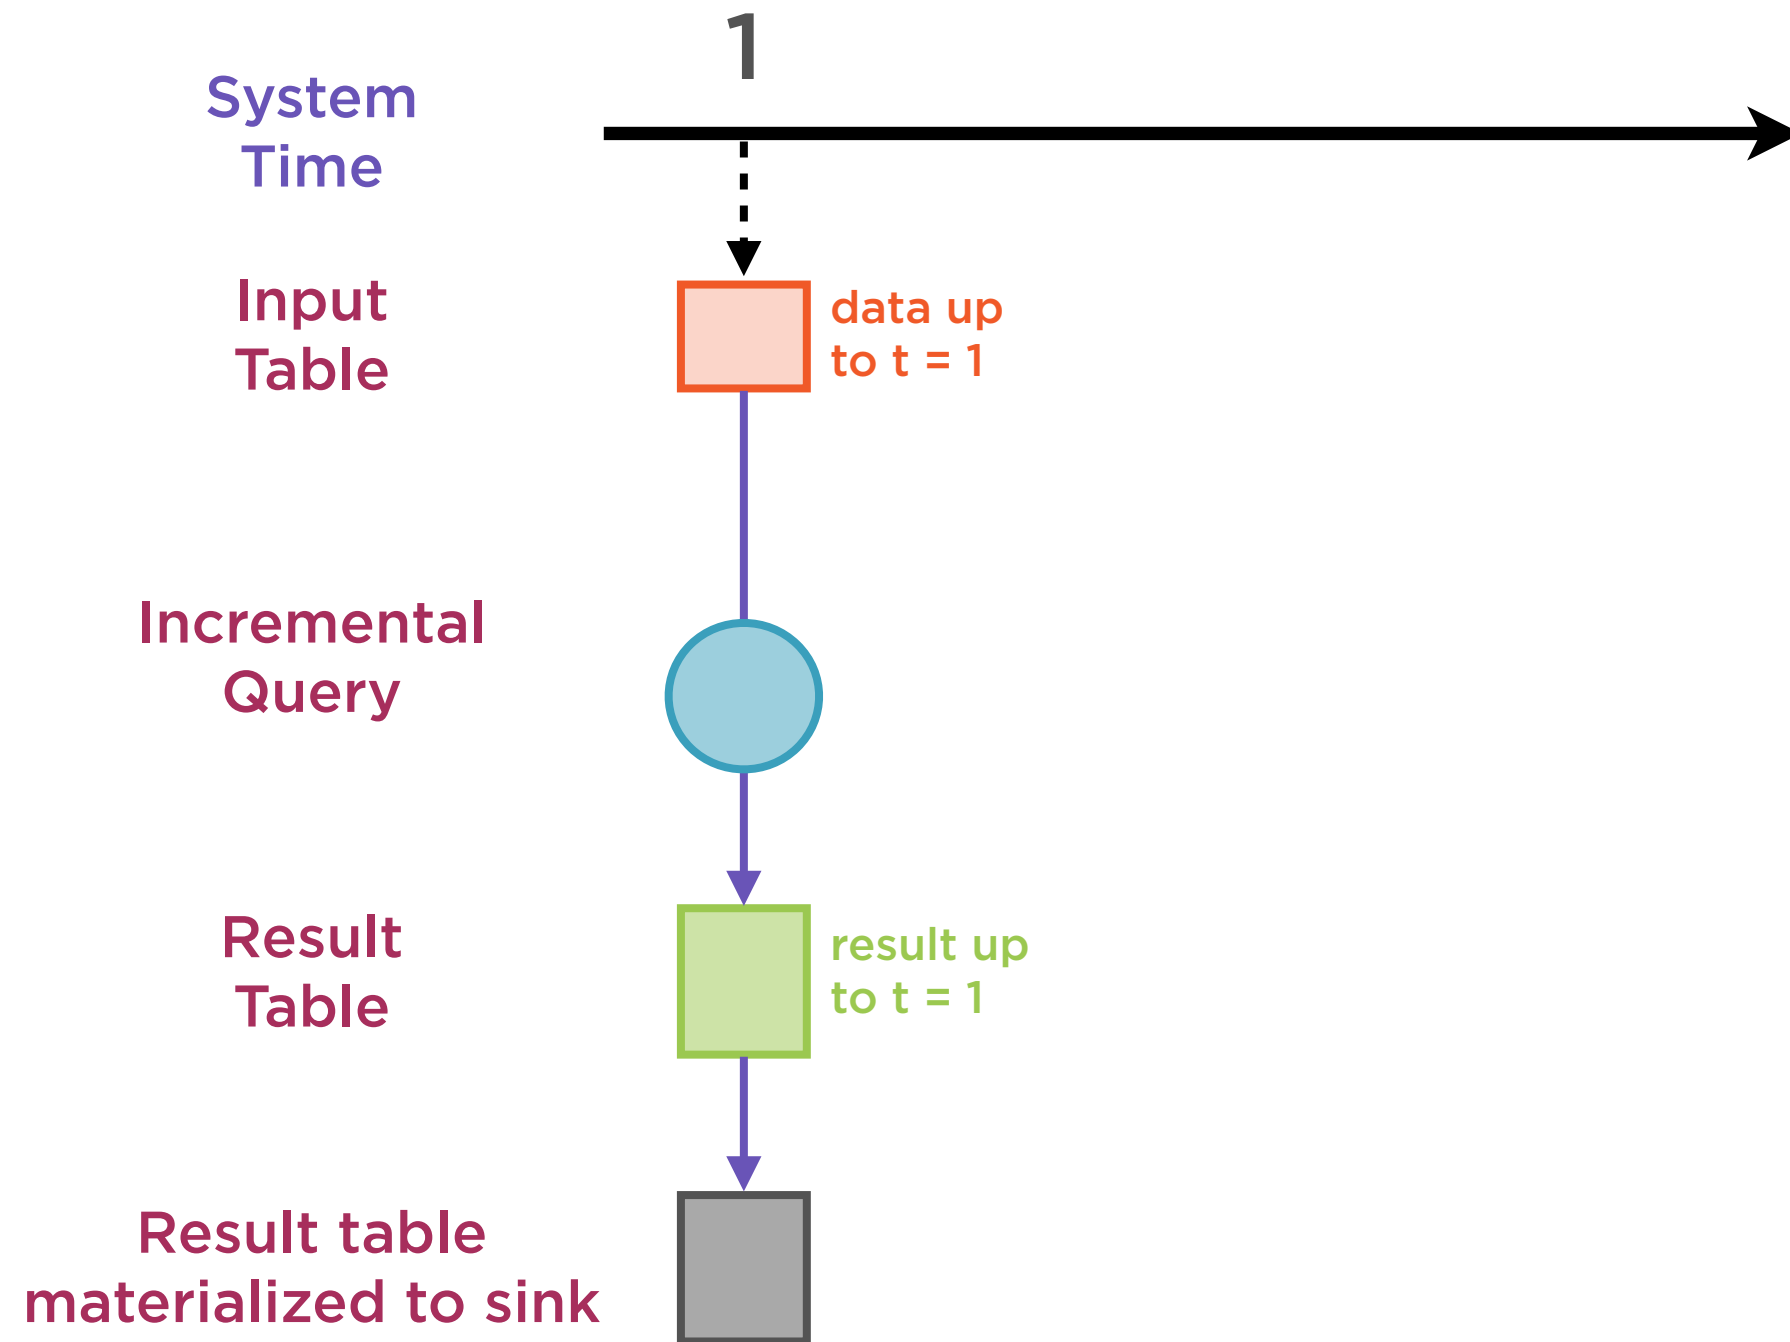
**User
Query**



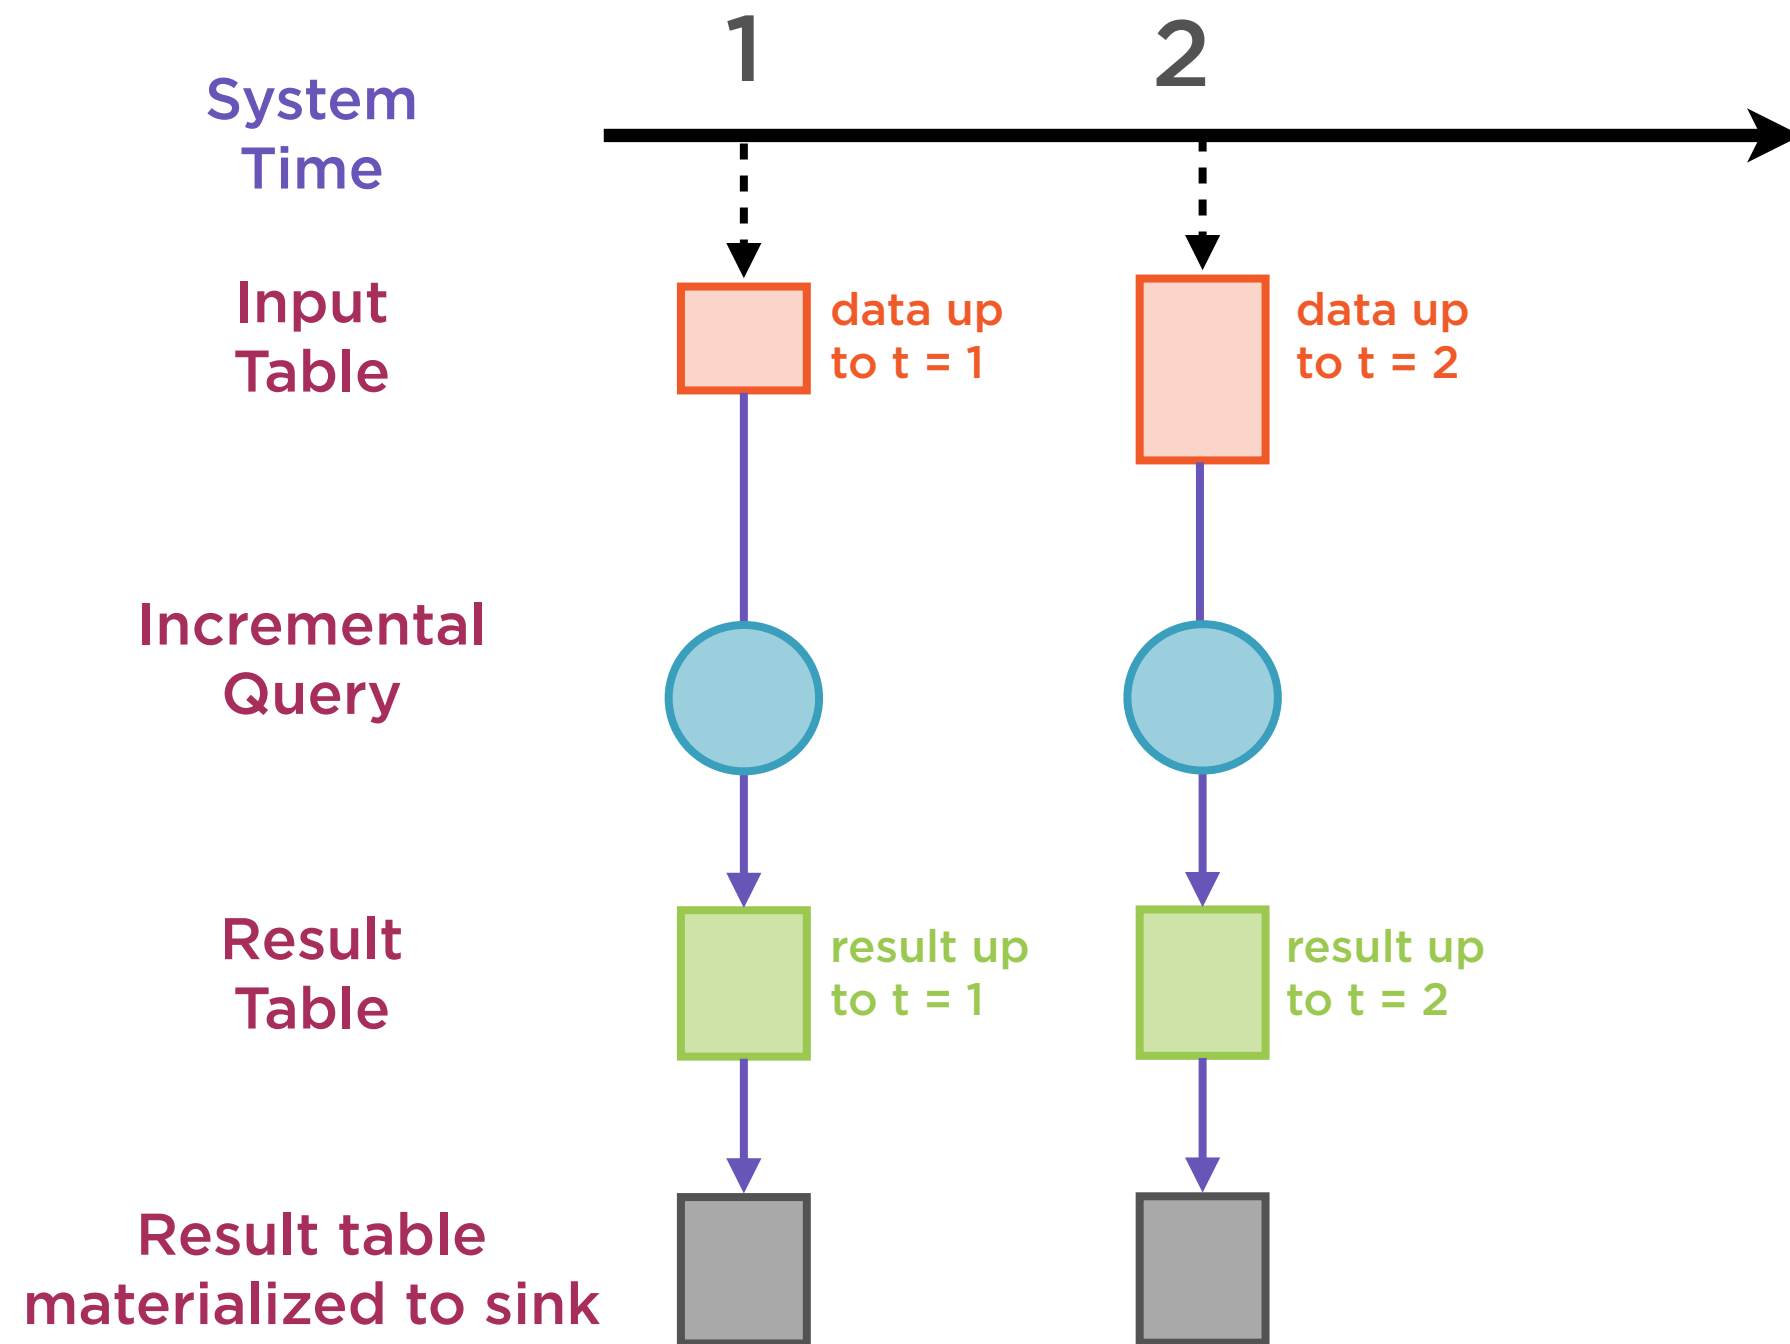
**Result
Table**




Result Table



Result Table




Result Table



When writing to the sink the entire
Result Table is not materialized

What is written out depends on
the **mode**


Output Modes



Determines what Result Table rows get sent to storage

- Update mode
- Append mode
- Complete mode

Output Modes




Update mode - only Result Table rows updated since last trigger

Even previous results will be updated in case of aggregations



Append mode - only Result Table rows appended since last trigger


Previous (existing) output rows cannot change



Complete mode - entire updated Result Table is sent across

Storage connector must decide how to use all that data

Output Modes



Update mode

Selections, projections, and aggregations



Append mode

Selections, projections, aggregations not supported



Complete mode

Selections, projections, aggregations, ordering

Demo

Projections and filtering in append mode using DataFrames

Demo

Projections and filtering in append mode using Spark SQL

Demo

**Aggregations and grouping in
complete mode using DataFrames**

Demo

**Aggregations and grouping in
complete mode using Spark SQL**

Demo

Aggregations and grouping in update mode using DataFrames

Demo

Aggregations and grouping in update mode using Spark SQL

Schema Inference

Schema Inference



Streaming DataFrames can be untyped

- Schemas will not be checked at compile-time

However some operations need schema information at compile-time

- map, flatMap

Untyped Streaming DataFrames can be converted to typed in such cases

Schema Inference



Structured Streaming from file sources requires schema to be specified

- Done by default to ensure consistent schema even in case of failures

For ad-hoc cases, can turn on schema inference

- `spark.sql.streaming.schemaInference = true`

Schema Inference



Perform schema inference on a small subset of the input stream

Process as batch data

Once schema known, use with streaming sources

Demo

**Adhoc schema inference using batch
data**

Summary

Prefix integrity and implications

Output modes - Append, Complete, and Update

Executing streaming queries in Spark

Schema and schema auto-detection

Up Next:

Understanding Scheduling and Checkpointing
