

Configuring Processing Models



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Stream processing models

Micro-batch execution and continuous processing

Considerations of latency, scaling, and recovery

At-least-once guarantees

Running Spark in continuous processing mode

Stream Processing Models

Stream Processing Models



Stream Processing Models



Stream processing does not necessarily mean
continuous real-time processing

Micro-batch Processing




Run transformations on smaller accumulations of data

Collect say less than one minute of data


Process this micro-batch in near real-time

Micro-batch Processing




A stream of integers

Micro-batch Processing



Grouped into batches


Micro-batch Processing



If the batches are small enough...

Close to real-time processing

Micro-batch Processing



Exactly once semantics

Replay micro-batches

Latency-throughput trade-off based on batch sizes

Batch Processing for Streams



Latency, freshness of data are not considerations

Complex analytical operations

Joins on relational data

- Data might be in a data warehouse, need not be in an RDBMS

Continuous Stream Processing for Streams



Latency and freshness of data are **most important** considerations

Rate of arrival is high

- Latency in seconds/milliseconds only possible with continuous processing

Micro-batch Processing for Streams



Latency and freshness of data are important

but

Real-time processing is overkill


Rate of arrival is low/moderate

- Latency in seconds/milliseconds less important
- Acceptable latency possible with micro-batches


Micro-batch processing is
the **default** mode in Spark

Spark also supports a new
**experimental, continuous
processing** mode

Micro-batch and Continuous Processing



Micro-batch and Continuous Processing




Micro-batch Processing in Spark

Structured Streaming treats a live data stream as a table that is being **continuously** appended

Streaming Data Spark 2.x

Data stream



--	--	--

--	--	--

--	--	--

--	--	--

Every data item
that is arriving on
the stream is like a
new row being
appended to the
input table

Data stream as an unbounded input table

Micro-batch Processing in Spark



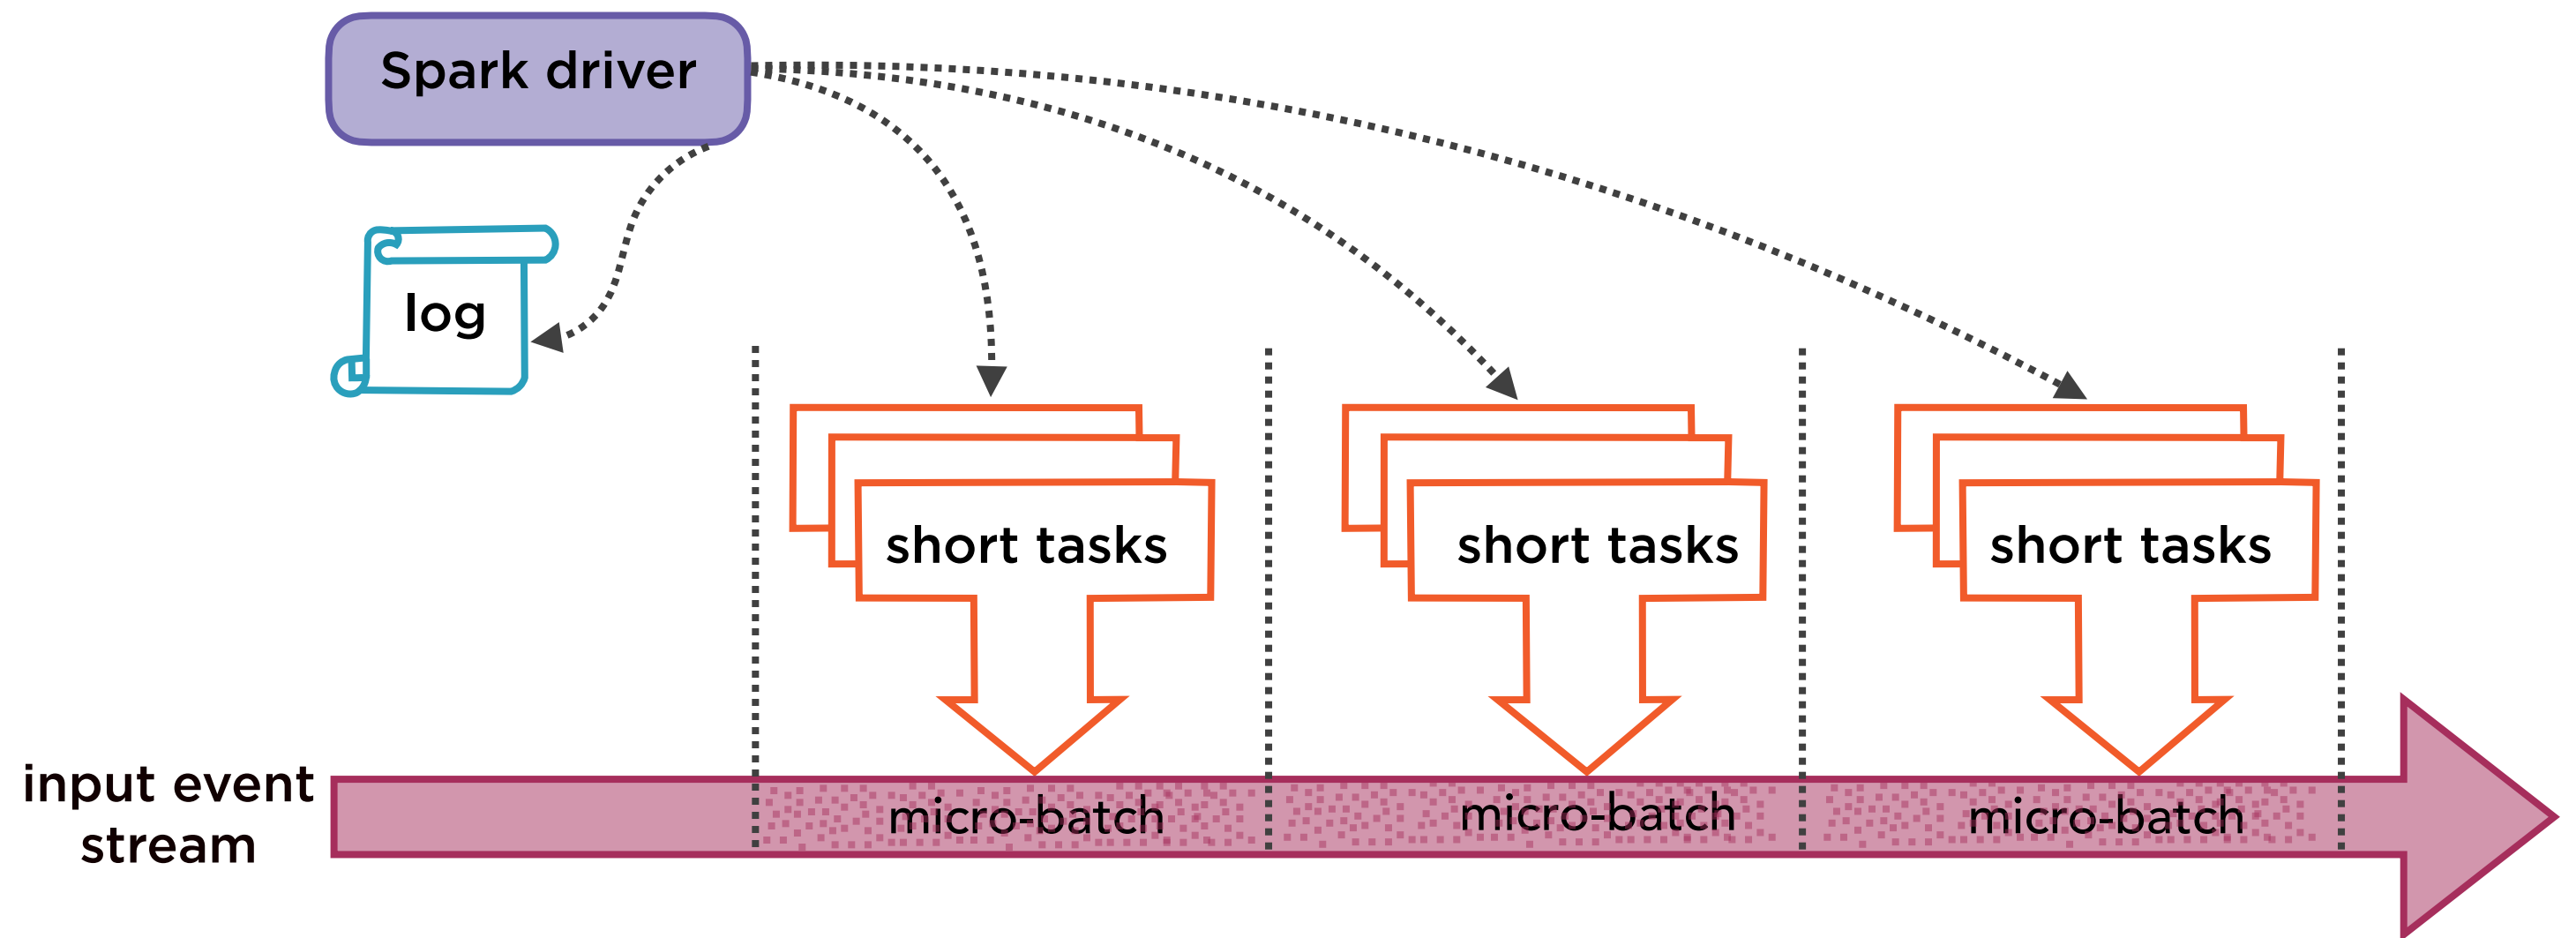
Default stream processing mode in Spark

Data streams processed as a series of batch jobs

End-to-end latencies as low as 100ms

Exactly-once fault-tolerance guarantees

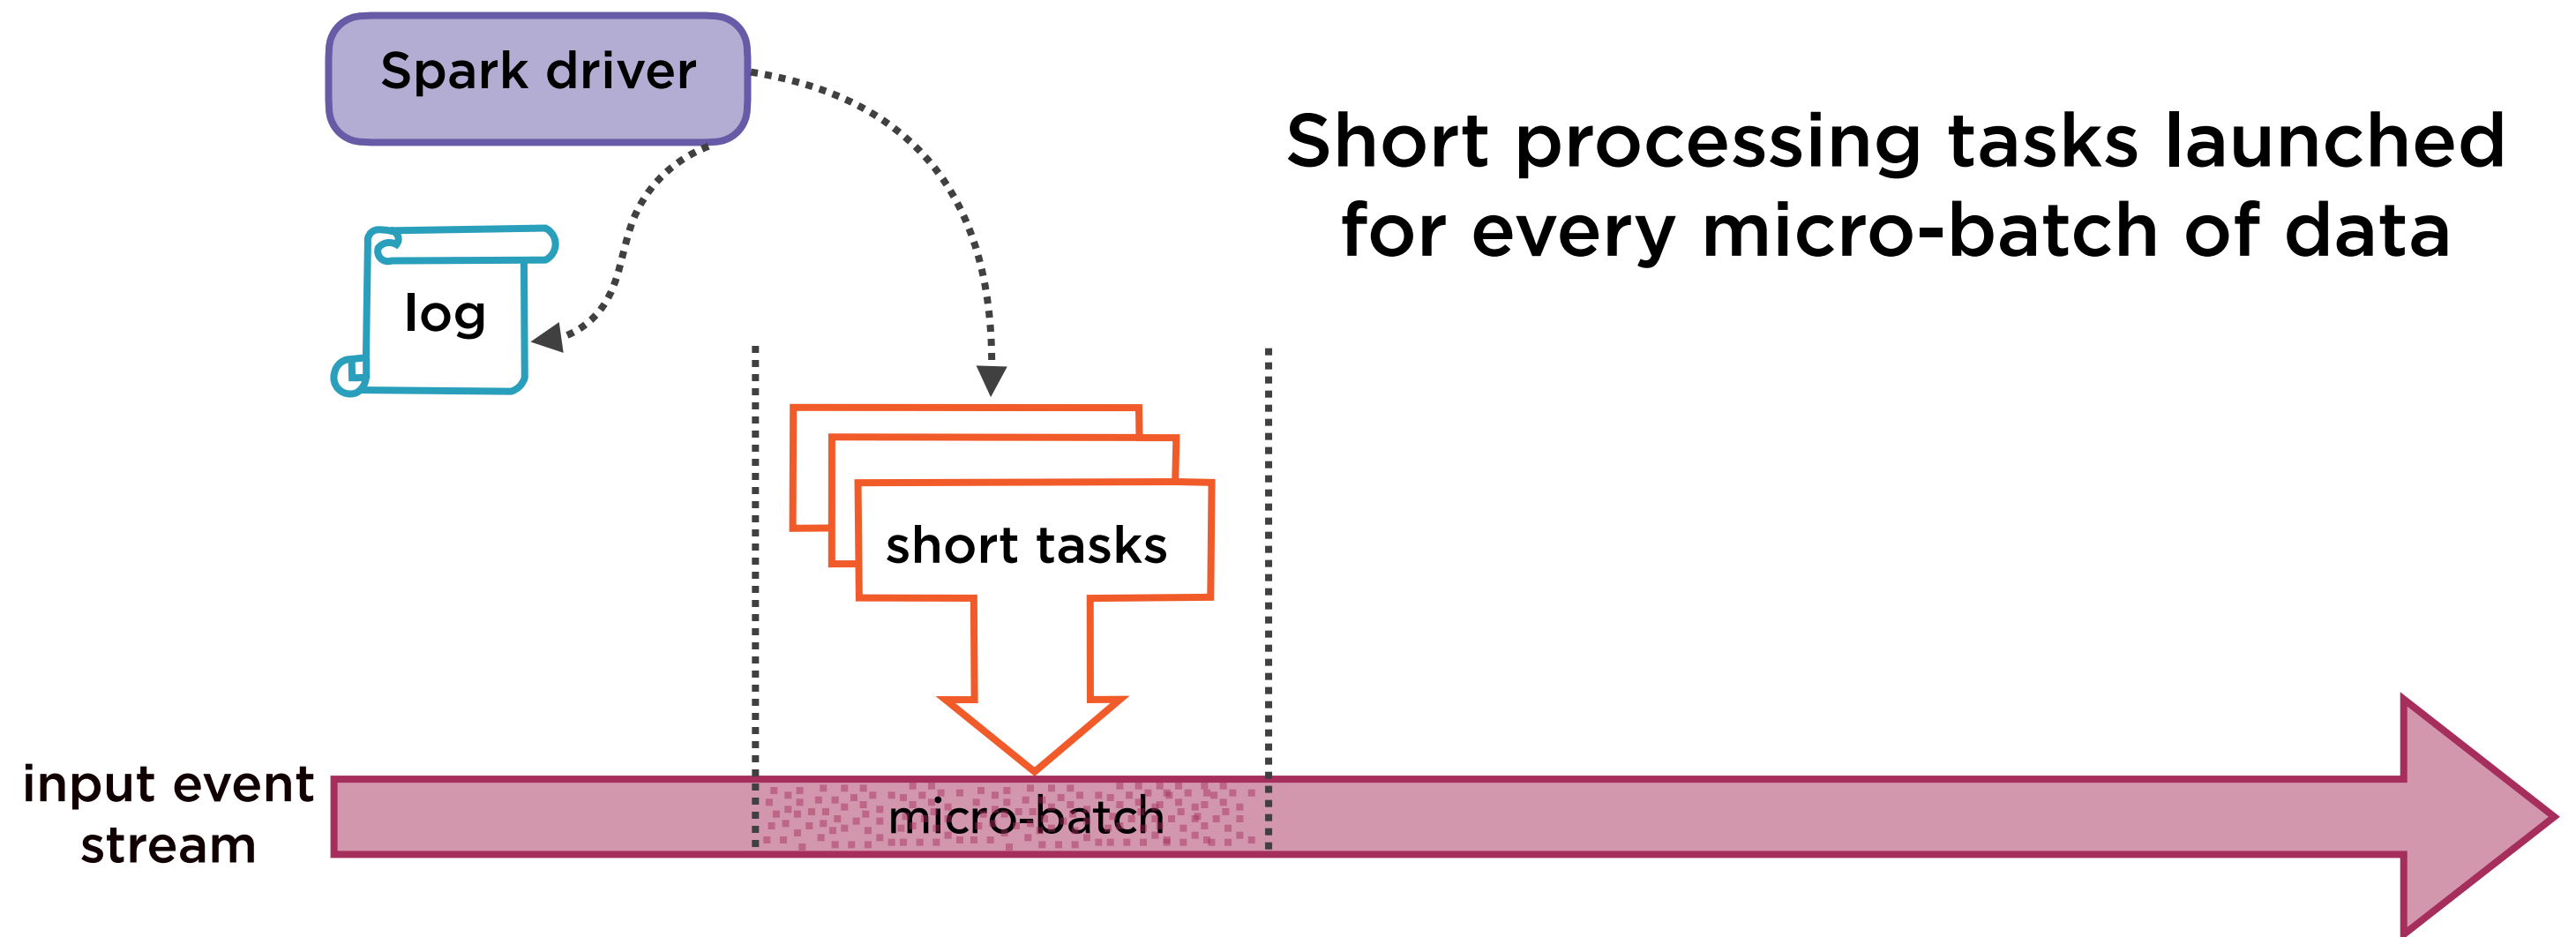
Micro-batch Processing in Spark



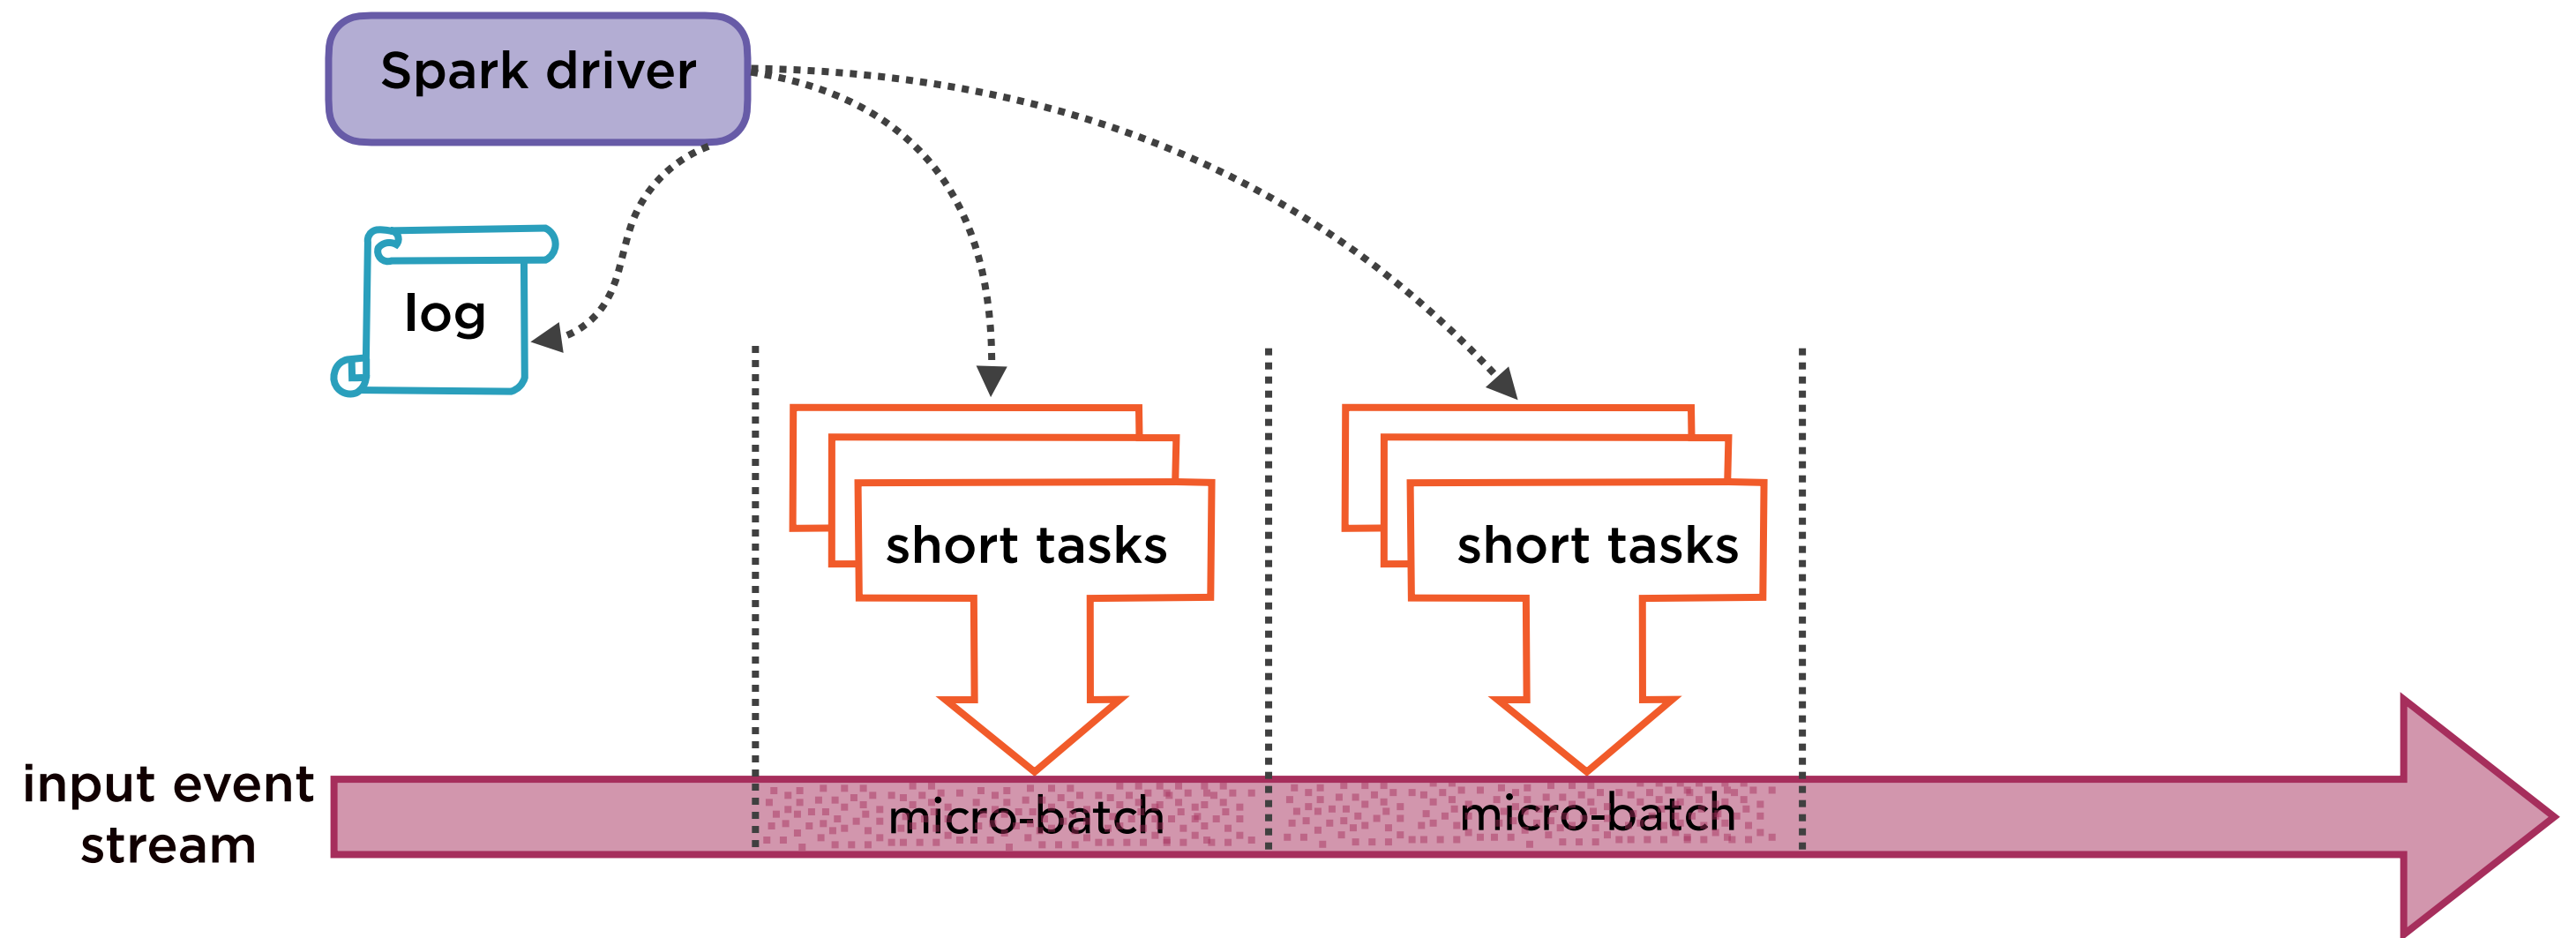
Micro-batch Processing in Spark



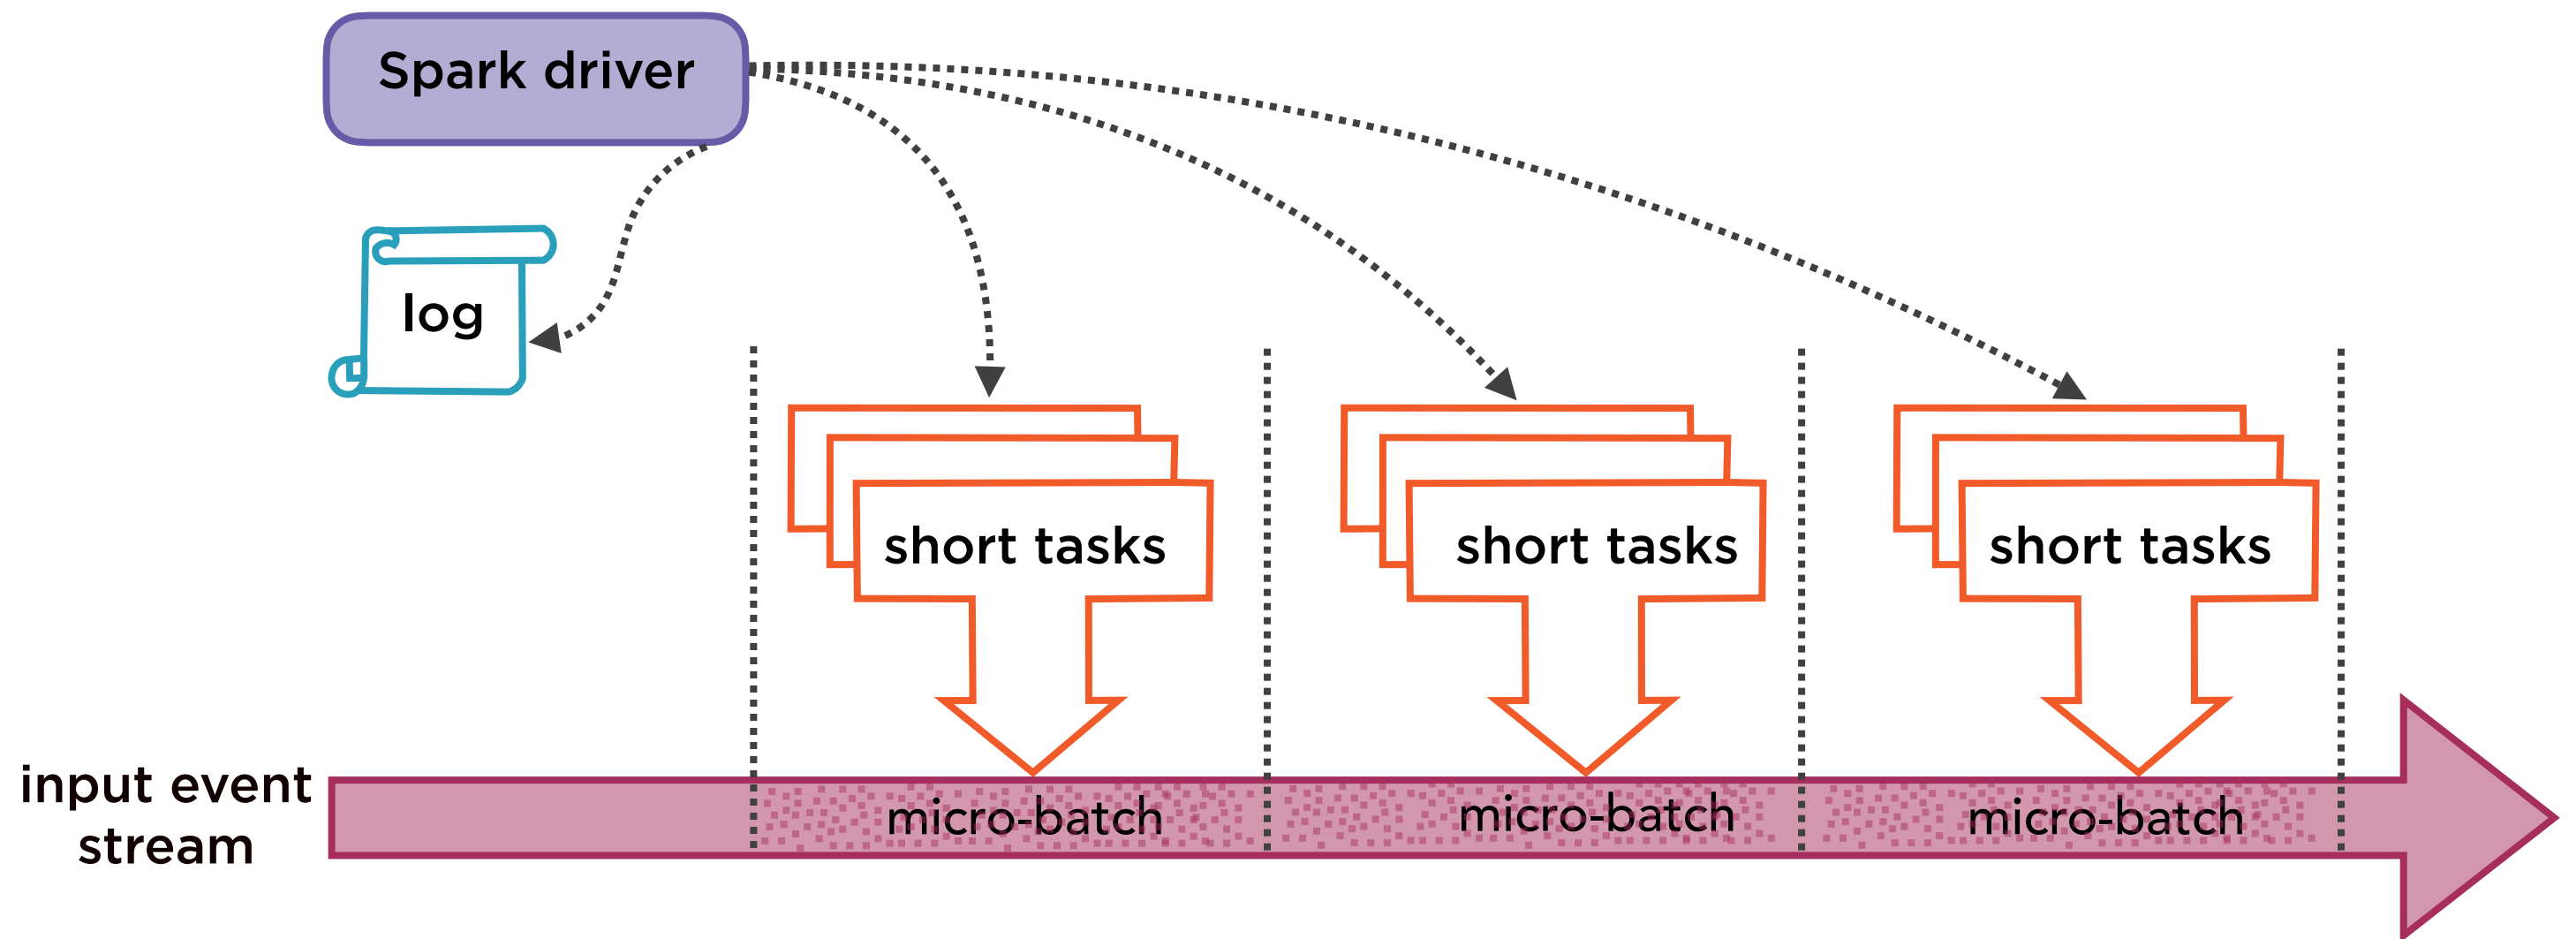
Micro-batch Processing in Spark




Micro-batch Processing in Spark




Micro-batch Processing in Spark



Higher End-to-end Latencies




Higher End-to-end Latencies




**Incoming streaming data
at the source**

Higher End-to-end Latencies




**Collected into a micro-batch which is
processed and written out to the sink**

Higher End-to-end Latencies




**Collected into a micro-batch which is
processed and written out to the sink**

Higher End-to-end Latencies




**Interval before another micro-batch is
created and processed**

Higher End-to-end Latencies



Higher End-to-end Latencies



End-to-end latency for processing is usually in the order of seconds

Continuous Processing in Spark

Continuous Processing in Spark




Experimental stream processing mode
in Spark

Data streams processed using long-
running tasks

End-to-end latencies a few milliseconds

At least-once fault-tolerance guarantees


Continuous Processing in Spark



Continuous Processing in Spark




Continuous Processing in Spark




Long-running tasks to continuously process the input stream of events

Continuous Processing in Spark



**Processed data written to write-ahead
logs every epoch**

Lower End-to-end Latencies




Lower End-to-end Latencies



**Incoming streaming data
at the source**

Lower End-to-end Latencies




**Processed continuously
using long running jobs**

Lower End-to-end Latencies




**Epoch markers used to write progress of
query out to checkpoint locations**


Lower End-to-end Latencies



Lower End-to-end Latencies



Lower End-to-end Latencies



**End-to-end latencies for
processing in milliseconds**

Micro-batches vs. Continuous Processing

Micro-batches

Latency in hundreds of milliseconds

Exactly-once guarantee

Spark Streaming launches many periodic tasks, each short-lived

Basis of Structured Streaming since Spark 2.0

All operations supported

Continuous Processing

Latency in few milliseconds

Only at-least-once guarantee

Spark Streaming launches a few, long-running tasks

Introduced as experimental feature in Spark 2.3

Restrictions apply - e.g. aggregation functions currently unsupported

Demo

Continuous processing using a rate source

Summary

Stream processing models

**Micro-batch execution and
continuous processing**

**Considerations of latency, scaling,
and recovery**

At-least-once guarantees

**Running Spark in continuous
processing mode**

Up Next:

Understanding Query Planning
